

DEMOS AI LAB

Pitch Deck

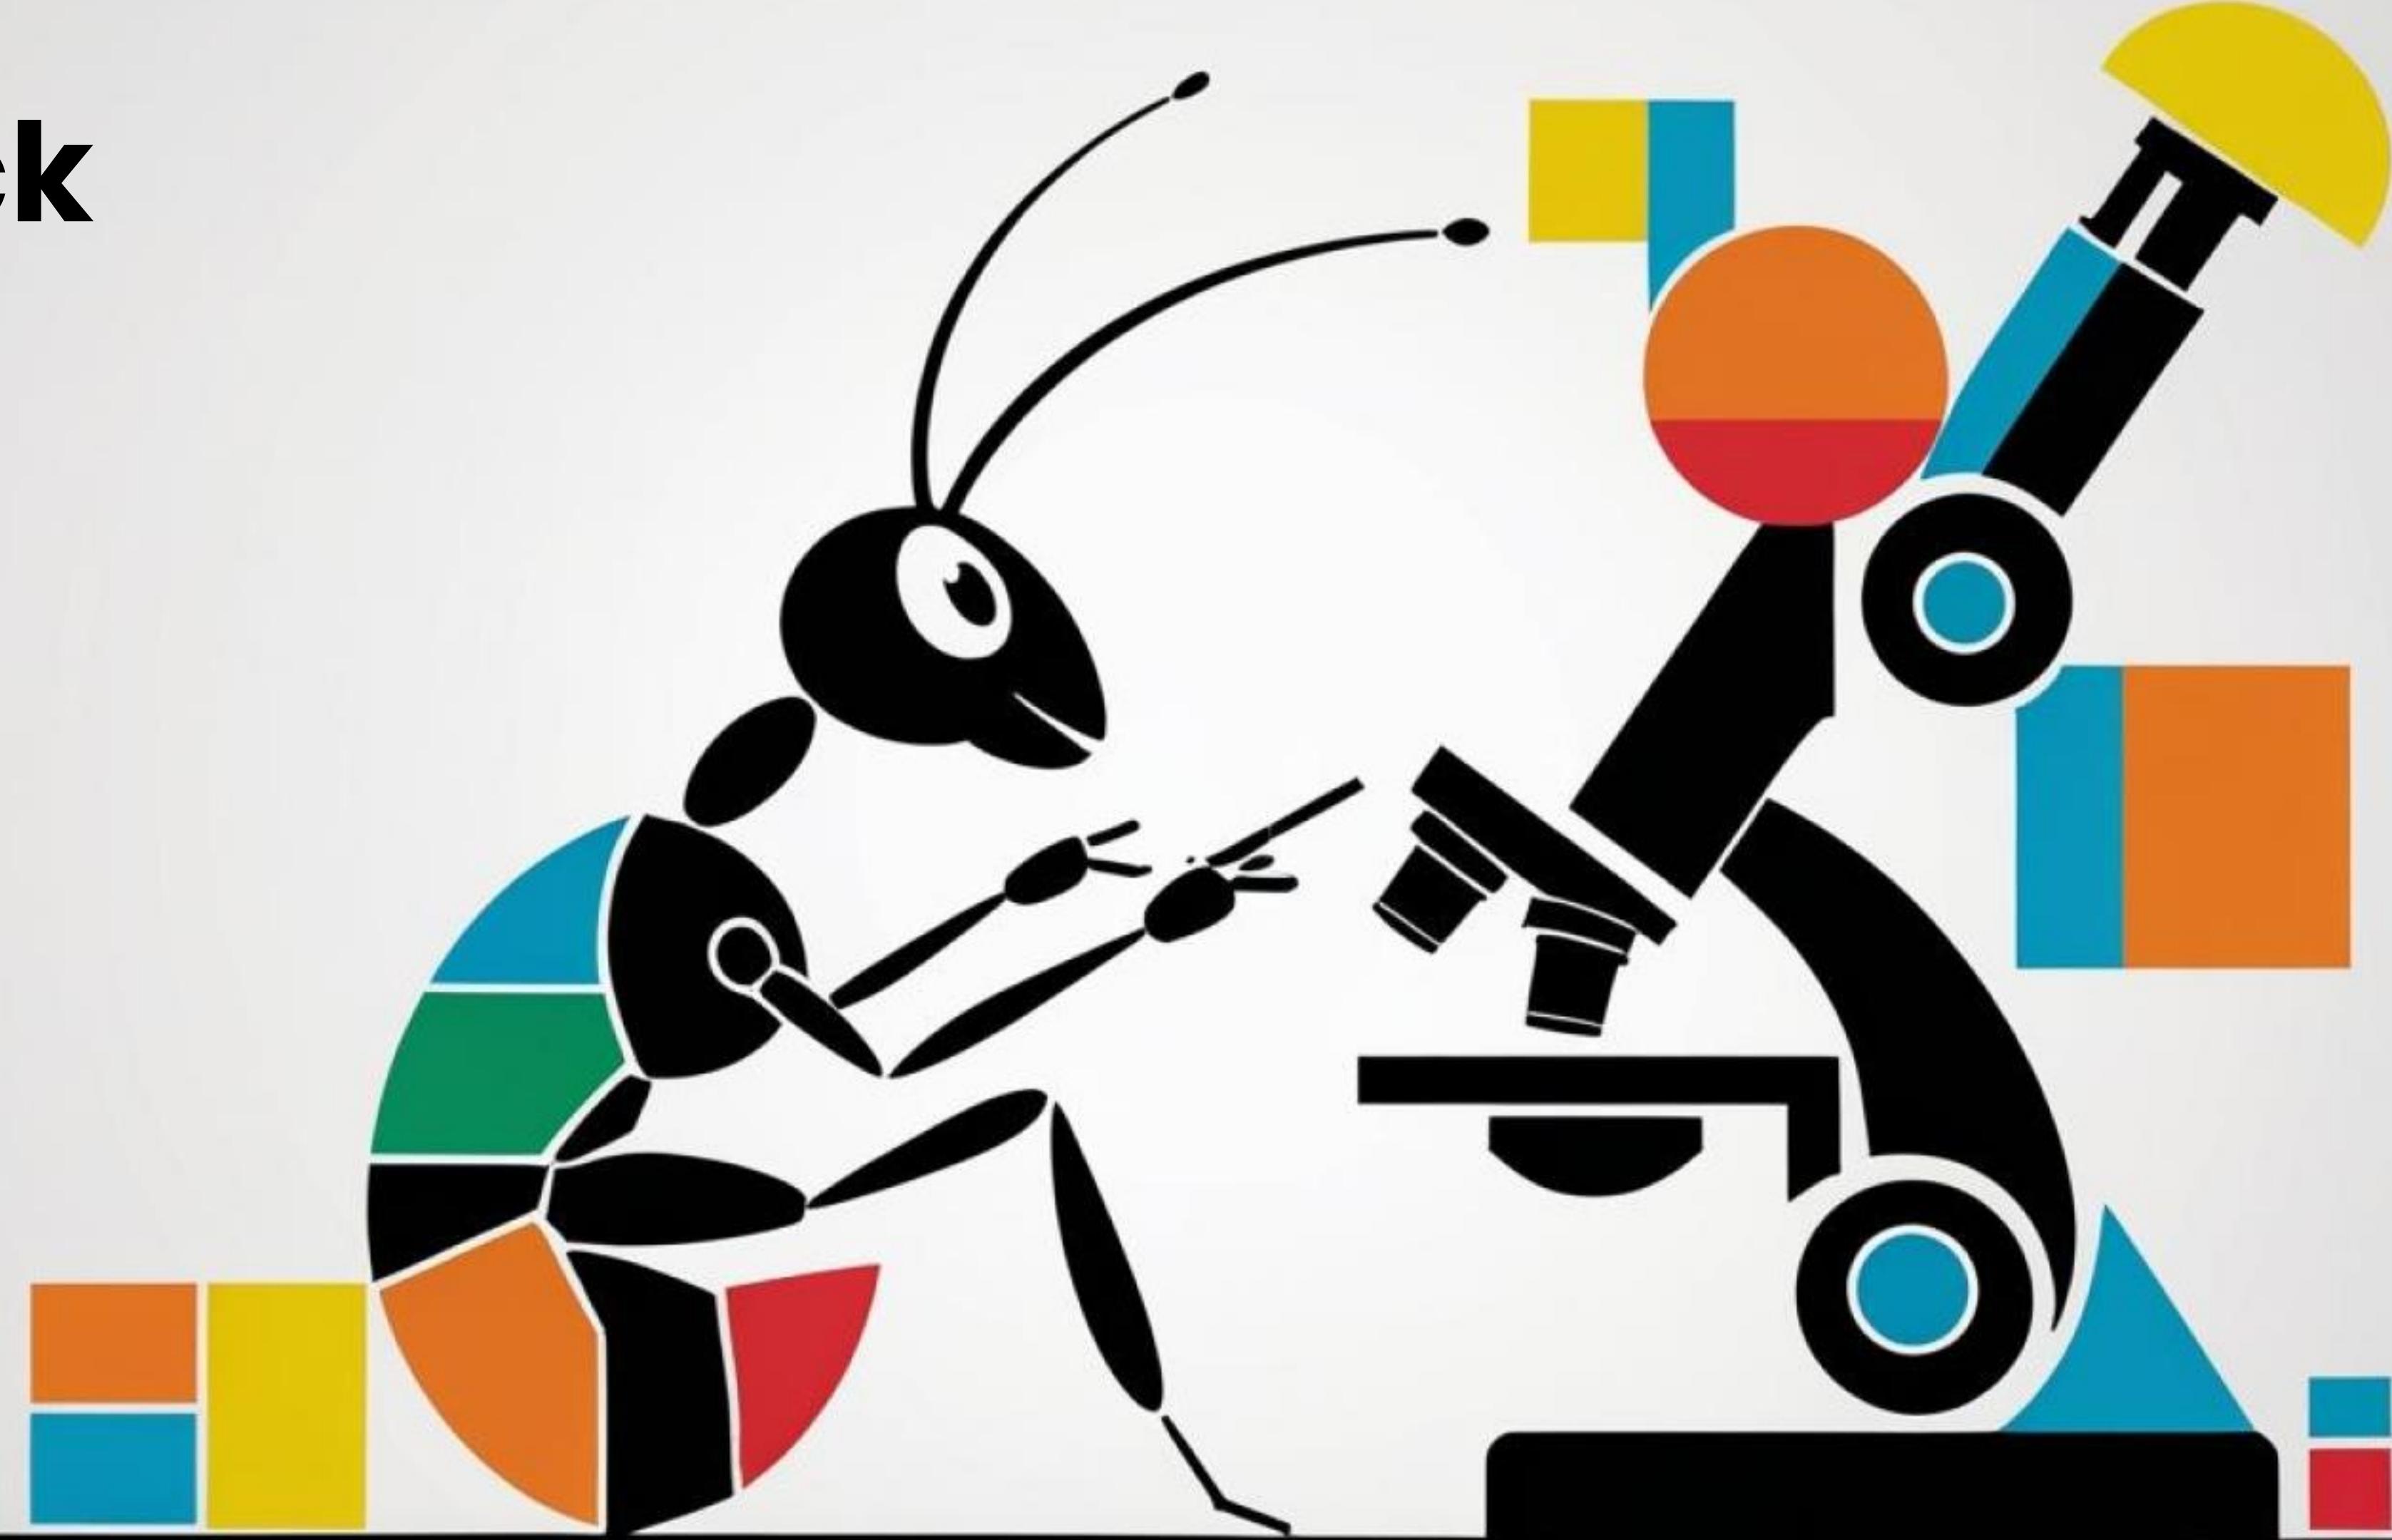


Table Of **Contents**

• The Problem	6
The 6.5 Billion Person Opportunity	8
The Real Cost Of Cloud Dependence	9
The Real Cost Of Language Apartheid	10
The Great AI Language Gap	11
The Strategic Revenue Opportunity	12
Recent MIT Study Reveals Big AI Is Failing Its Promise	13
• The Solution	14
Language-Based Edge AI	15
Edge AI	16
What Are Small-Parameter Models?	17
Types of Small Models for Edge Deployment	18
Natural Language Processing: Building The Language Foundation	19
Natural Language Processing: The Hidden Core Of The AI Economy:	20
DEMONS AI Will Build Native Language Infrastructure	21
Why This Changes Everything	22
• The Market	23
The Largest TAM In AI History	24
Edge Is Where The AI Economy Lives	26
The \$140B Battleground Where AI Meets The User	27

Table Of **Contents**

Edge AI Examples	29
AI Demand By Region	30
Natural Language Processing	38
Global Training Data Infrastructure	39
Market Overview	40
AI Training & Data Market	42
US Government AI SPENDING	43
Department Of Defense AI Spending (Direct + Embedded)	44
US Government Infrastructure Investments	45
Defense AI Budget (FY 2026)	46
GOVT Contract Awards & Market Outlook	47
US Government Contractor Example: Palantir	48
US Government Market Validation: Auterion	49
• The Products	50
Product 1: Data Labeling Ecosystem	51
Product 2: Edge AI Ecosystem Platform	54
Product 2: Implementation of Edge AI Agent	58
Product 3: Omnestra - The Agent APP Store	59
• The Business Model	61
Demos AI Lab Structure	62

Table Of **Contents**

Edgentic Intelligence Division	63
Rosetta ML Division	64
The Profitability Case – Why Smaller Models Win	65
Edge AI	66
GPT-5 (Orion) – The \$500 Million+ Challenge	67
The \$100 Fine-Tuning Revolution	68
Databricks Dolly – AI for Less Than \$30	69
The Great Divide	70
Revolutionary Unit Economics	71
Eight Revenue Streams Per Language	72
Financial Revenue Projections	73
Financial Cost Projections	75
Financial Breakdown	76
Financial EBITDA Projections	78
• The Roadmap	79
Global Language Roadmap: 3 Phases	80
Product Roadmap: 4 Phases	81
48-month Product Roadmap To Market Dominance	82
• The Competitive Advantage	86
Early Traction & Momentum	87
Our Unbreakable Moats	88

Table Of **Contents**

Competitive Matrix	89
• The Team	90
Founders With \$3B+ Operational Experience	91
Board Of Directors – Unmatched Access	92
World-class Advisory Board	93
• The Ask	94
The Investment Opportunity	95
Use Of Funds & Path To Unicorn	96
• The Vision	97
Puerto Rico: The Perfect Launchpad For Hermigas	98
Join The AI Revolution	100
AI For Every Language, Every Human	102
Closing Statement	103
• The Appendix: Branding	104
Appendix: Branding	105

The Problem

AI For **Every Language, Every Human**

Language Is The Horizon Of Consciousness Itself

"Language is not merely a tool; it is the very horizon of consciousness itself."

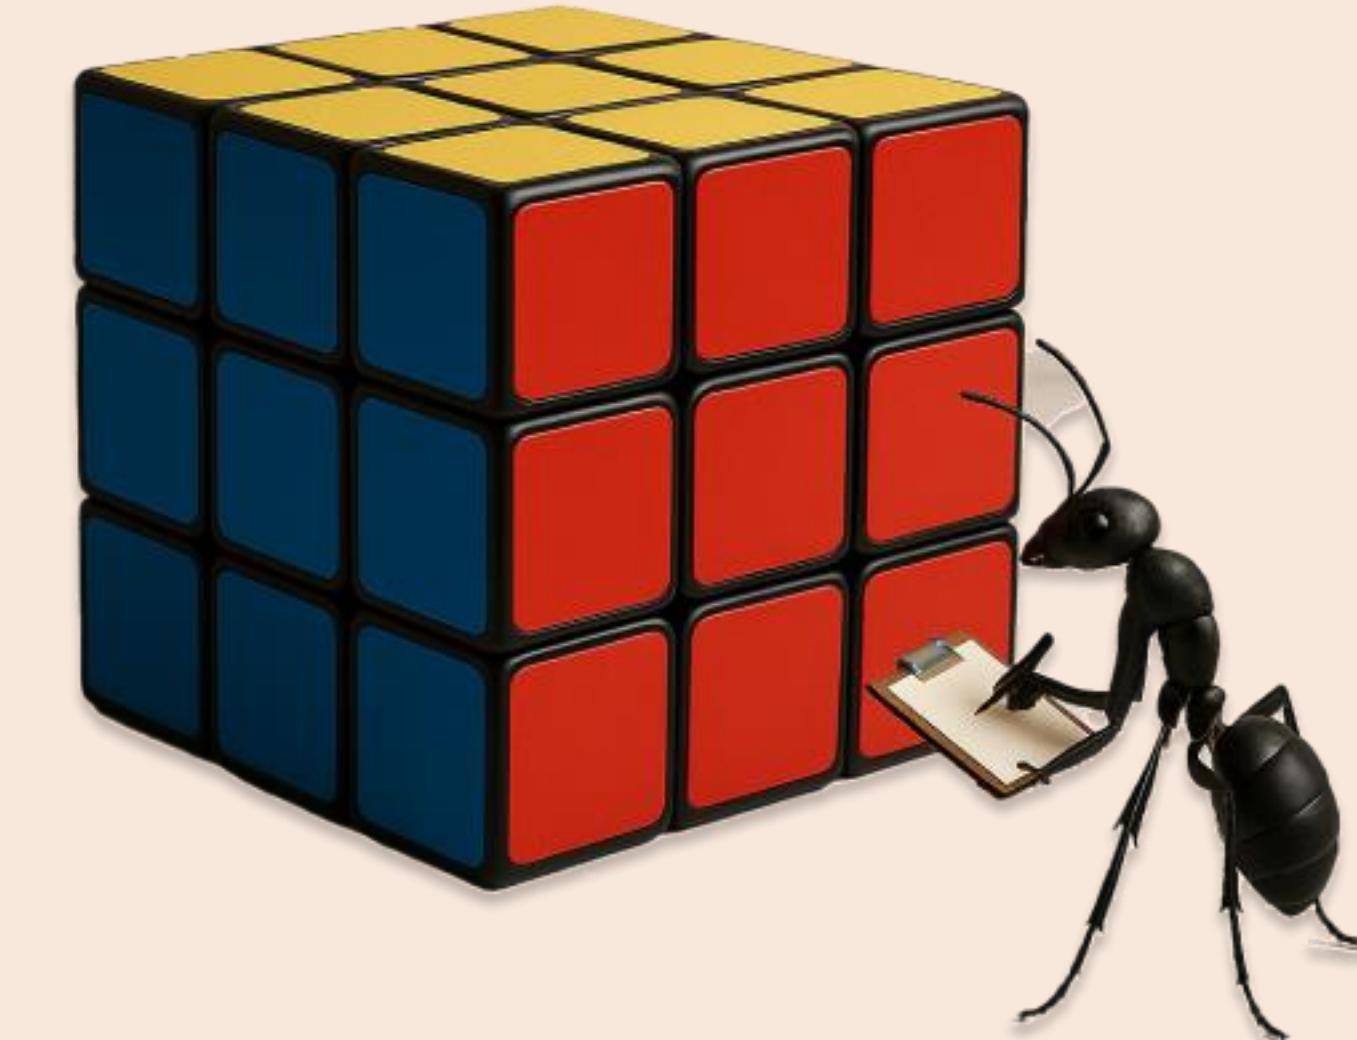
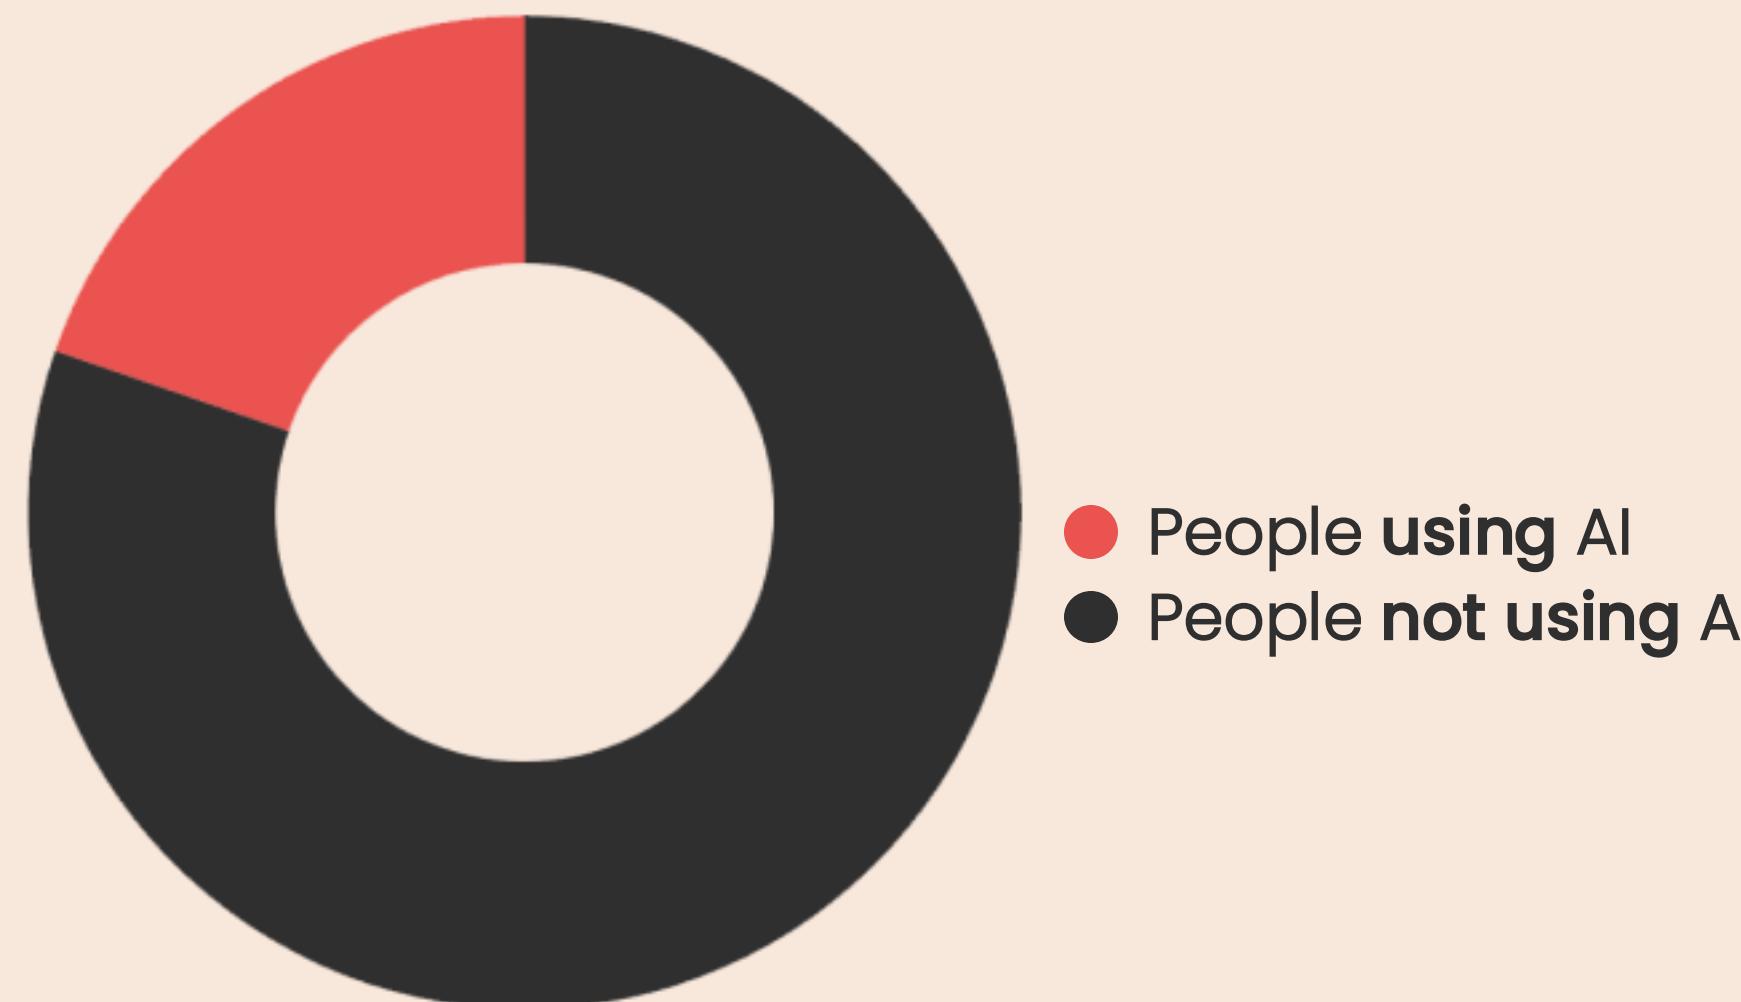
- Jacques Derrida

What This Means: AI trained only in English can never truly serve Spanish speakers. Each language deserves its own AI ecosystem that preserves its unique ways of seeing, thinking, and being.

The 6.5 Billion Person Opportunity

6.5 Billion People Are Locked Out Of The AI Revolution

- AI today is both English-centric and cloud-dependent
- 7,000+ languages exist – yet **less than 1%** are properly served by AI
- **Cloud AI can't reach most of the world:** unstable Internet, unreliable power, or unaffordable \$20/month subscription
- As a result, **only ~2 billion people** have practical access to advanced AI tools
- This isn't a tech problem—**it's a \$2 trillion market opportunity** (Grandview Research, Artificial Intelligence Report 2025)



The **Real Cost** Of Cloud Dependence

Cloud AI's Physical Limitations Create Massive Market Gaps

- **Infrastructure Requirements:** Stable power, fiber internet, expensive subscriptions
- **Resource Intensity:** One data center can easily consume 300,000 gallons of water daily
- **Market Reality:** Communities providing training data can't access or afford the AI
- **Opportunity:** Edge Based Language AI serves these markets at 10% of infrastructure cost—finally making them accessible and profitable

Cloud AI's infrastructure cost

Edge AI's infrastructure cost



The **Real Cost** Of Language Apartheid

Current AI Language Bias

Definition: **Language Apartheid** is the digital divide created when AI systems serve only English-speaking populations—excluding billions who speak other languages

Key Characteristics:

- **6.5 Billion People:** Excluded from advanced AI tools because their languages are not supported
- **\$\$ Trillions in Economic Potential:** Lost across non-English markets.
- **Data Bias:** Approx. 90% of training data used in major models is in English, erasing cultural and linguistic diversity
- **Sovereignty Risk:** Nations without language infrastructure depend on foreign cloud systems for their own knowledge processing
- **Technological Inequality:** AI built for a **few languages** embeds **bias and misinterpretation into global systems**

Why It Matters: When AI can't understand a language, it can't serve its people, industries, or governments. Language Apartheid isn't just a technical failure—it's an economic, cultural, and geopolitical one.

DEMOS AI LAB Context: DEMOS ends Language Apartheid by creating **language-native AI ecosystems**—built locally, owned locally, and operated on-device. Each language gains its own data infrastructure, models, and edge-based intelligence—restoring **sovereignty, inclusion, and dignity** in the age of AI

The Great AI Language Gap

English Centric AI

Current AI speaks to 1.4 billion people. It ignores the other 6.5 billion.

- According to Statista (2024): Roughly 1.4 billion people speak English – native or otherwise. The world's population? Over 8 billion (UN, 2024).
- Do the math: 80% of humanity doesn't primarily use English. Yet nearly every commercial AI system is designed, trained, and optimized in English first.

It gets worse:

- Ethnologue (2024) – the world's definitive linguistic database counts 7,164 living languages on earth. How many show up in modern AI models? Fewer than 100 (Bender et al., arXiv 2020).

Let that sink in:

- 80% of people underserved
- 98% of languages invisible
- Leaving Trillions on the table

The bottom line: The world's largest tech opportunity isn't in the next LLM breakthrough. It's in building AI that actually speaks to everyone else.

The **Strategic Revenue** Opportunity

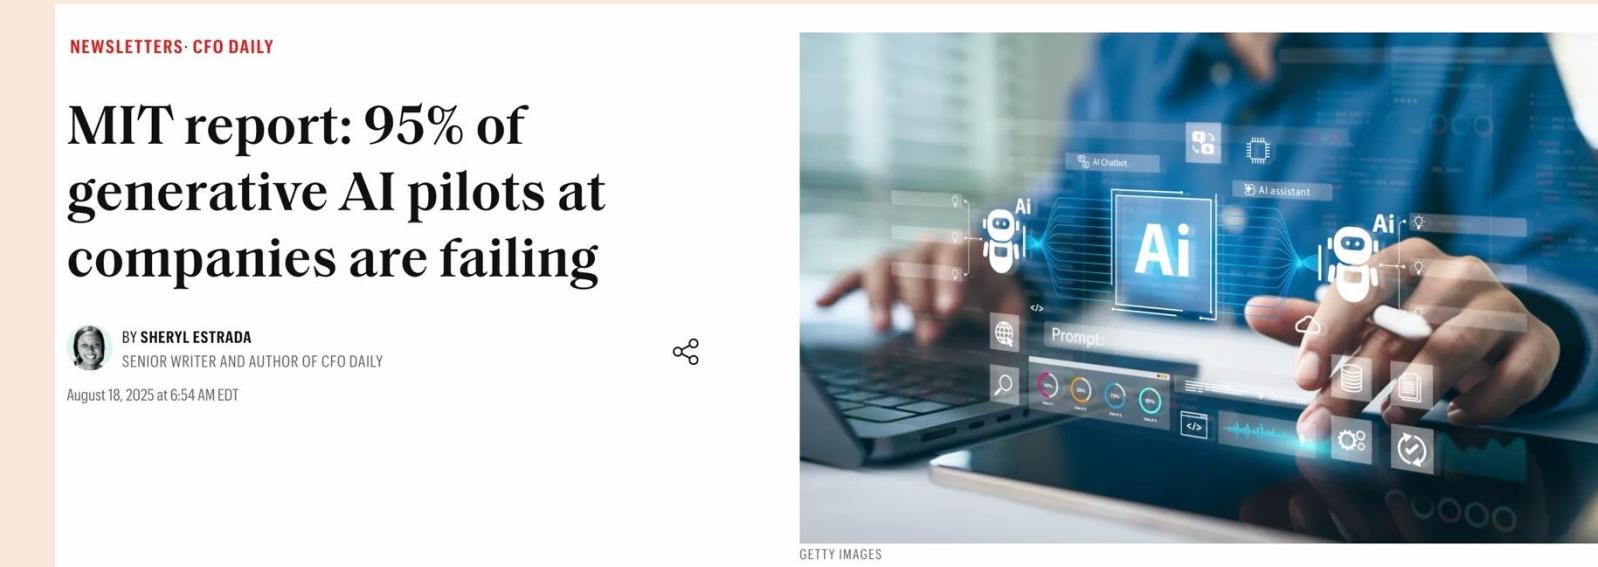
Language Edge Based AI Becomes **Strategic Necessity And Revenue Driver**

- 7,000+ languages exist, <1% properly served by AI
- Cloud AI physically can't reach most—no stable internet, power, or \$20–25/month to spare
- Current AI serves only those with infrastructure = **2B people maximum**
- This isn't a tech problem—**it's a \$ trillion market opportunity**



Recent MIT Study Reveals Big AI Is Failing Its Promise

Wall Street's Fear Realized: Zero ROI from Big AI



The image consists of two parts. On the left is a screenshot of a news article from 'NEWSLETTERS - CFO DAILY'. The headline reads 'MIT report: 95% of generative AI pilots at companies are failing'. Below the headline is a photo of Sheryle Estrada, the author. The text below the photo indicates the article was published 'BY SHERYLE ESTRADA SENIOR WRITER AND AUTHOR OF CFO DAILY August 18, 2025 at 6:54 AM EDT'. On the right is a photograph showing a person's hands interacting with a laptop screen. The screen displays a futuristic, semi-transparent user interface with various icons and the word 'AI' prominently featured. The source of the image is cited as 'GETTY IMAGES'.

- 95% of organizations saw zero ROI on GenAI investments despite \$30–40B in enterprise spending, according to MIT (2025)
- **Cloud AI's economics are collapsing** – massive energy, water, and data center costs exceed any measurable return
- Centralized models can't scale globally, breaking down in multilingual, low-connectivity, and sovereign environments
- Wall Street and governments are losing confidence in the centralized AI model, creating an opening for Edge-native innovation

The Solution

Language-Based Edge AI

The Solution

To serve the 6.5 billion people excluded from today's cloud AI:

We must combine two breakthroughs:

- Edge Computing: runs AI directly on local devices — phones, laptops, drones, IoT sensors
— no internet or data centers required
- Natural Language Processing (NLP): enables AI to understand and respond in every native language, dialect, and cultural context

Why It Matters:

- The fusion of these technologies creates Language-Based Edge AI — intelligent, offline systems designed for each language ecosystem
- Each model operates locally, learns from its environment, and contributes to a shared knowledge network coordinated by the Ant Colony Architecture — a decentralized system where thousands of edge models continuously exchange insights with a central "Queen" node that refines and redistributes collective intelligence back to the network

Edge AI

The Solution

Definition: Edge AI is the practice of running **artificial intelligence models directly on local devices** — from smartphones and laptops to drones, IoT sensors, and microcontrollers — without requiring cloud access. Unlike traditional edge computing, which only processes raw data locally, Edge AI performs **model inference and decision-making** on-device using **Small Parameter Models (SPMs)** optimized for low power and memory environments.

Key Characteristics:

- **Local Inference:** AI models execute directly where data is generated — on the device itself.
- **Small Parameter Models (SPMs):** Lightweight models (<30B parameters) trained for high accuracy and efficiency, often using FP8 or FP4 quantization.
- **No Cloud Dependency:** Operates entirely offline — reliable even during outages or in disconnected regions.
- **Low Latency:** Millisecond responses with no round-trip to remote servers
- **Data Sovereignty:** Sensitive data never leaves the device or network
- **Energy & Cost Efficiency:** 80–90% lower power and compute requirements than cloud AI

Why It Matters: Edge AI democratizes access to intelligence — enabling advanced AI capabilities in **low-connectivity environments** such as rural clinics, defense zones, and developing regions. It protects **privacy**, enforces **sovereignty**, and brings **AI accessibility to the 6.5 billion people** without stable cloud infrastructure

DEMOS AI LAB Context: Our **Ant Colony Architecture** connects thousands of lightweight edge models—each running locally on phones, drones, and IoT devices—to a coordinating “home cloud” we call the **Queen**. These edge agents learn from their environments and send summarized insights to the Queen, which refines and retrains the shared model. Updated intelligence is then redistributed back to the network, ensuring every device improves continuously. This living feedback loop creates a **self-learning, sovereign ecosystem** that operates independently of centralized cloud infrastructure while becoming smarter with every interaction

What Are Small-Parameter Models?

Definition: Small-parameter models are AI models with 1–30 billion parameters designed for efficient deployment on edge devices. Unlike massive cloud models (100B+ parameters), these compact models deliver high performance while running locally on smartphones, laptops, and embedded systems.

The Power of Specialization: Small models trained on curated, domain-specific datasets often outperform massive general models on targeted tasks—medical Q&A in Spanish, legal document analysis, real-time translation, code generation.

Paramter Ranges & Deployment:

- **1–3B parameters:** Microcontrollers and low-end smartphones
- **3–7B parameters:** Modern smartphones and tablets
- **7–30B parameters:** Laptops, desktops, and edge servers

Key Advantages:

- **On-device deployment:** Works offline, no cloud dependency
- **Fast inference:** Millisecond response times
- **Cost-effective:** No per-query fees, one-time deployment cost

Why It Matters:

- **Democratizes AI**—works on hardware people already own
- **Enables privacy and sovereignty**—data never leaves the device
- **Serves 6.5 billion people** with unreliable internet or limited cloud access

DEMOS Context: Edgentic Intelligence Division creates small-parameter models, while Rosetta ML Division trains them on premium multilingual datasets—delivering sovereign AI that runs anywhere, for anyone.

Types of Small Models for Edge Deployment

Small Language Models (SLMs): Text processing, conversation, translation, code generation (1–30B parameters) | Example: Native Spanish medical Q&A, legal contract review

Small Vision Models: Image recognition, object detection, autonomous navigation | Example: Drone obstacle detection, crop disease identification

Small Audio Models: Speech recognition, sound classification, voice commands | Example: Real-time speech-to-text, acoustic monitoring

Small Multimodal Models: Combined vision + language understanding | Example: Visual question answering, document analysis with explanations

Time-Series Models: Sensor data analysis, predictive analytics, anomaly detection | Example: Equipment failure prediction, energy consumption forecasting

Reinforcement Learning Models: Robotics control, autonomous decision-making, optimization | Example: Drone navigation, warehouse automation

TinyML Sensor Models: Ultra-compact (KB–MB) for specific sensor tasks | Example: Heart monitoring on wearables, vibration analysis

Why It Matters: Real-world applications require multiple model types working together – vision, language, audio, sensors, and control systems – deployed as coordinated swarms for complex tasks.

DEMOS Context: Edgentic Intelligence Division builds small models across all these types, creating a complete ecosystem of specialized agents for our Ant Colony Architecture.

Natural Language Processing: Building The Language Foundation The Solution

Definition: Natural Language Processing (NLP) enables machines to understand and generate human language across regions, dialects, and cultures.

Key Characteristics:

- Learns language structure, meaning, and intent from text and speech
- Adapts to regional context, idioms, and tone
- Expands comprehension beyond English to thousands of local languages
- Requires language-specific data and fine-tuning to ensure accuracy

Why It Matters:

- AI cannot serve the world until it speaks its languages
- Building NLP infrastructure for every language ensures inclusion, trust, and sovereignty—making intelligence truly global

DEMOS AI LAB Context:

- Unsupervised learning can train models on massive text corpora, giving them a general sense of structure and meaning.
- But this approach misses cultural and linguistic precision—especially in underrepresented languages.
- DEMOS starts by building language-specific data infrastructure, beginning with local annotation and labeling led by native linguists
- These smaller, high-quality datasets become the foundation for fine-tuning our Language-Based Edge AI models, ensuring accuracy, authenticity, and sovereignty in every language we serve

Natural Language Processing

The Hidden Core Of The AI Economy:

Training Data Infrastructure

Component	What It Involves	Examples/players
Annotation & Labeling	Human-in-the-loop Labeling Of Text, Images, Video, Audio, And Sensor Data – The Foundational Step That Transforms Raw Data Into Structured Training Material.	Scale AI, Labelbox, Remotasks
Data Curation & Qa	Deduplication, Cleaning, Bias Detection, Metadata Tagging, And Quality Scoring To Ensure Datasets Are Accurate, Representative, And High-value.	Scale AI, Labelbox
Feedback And Loops	Human Ranking And Comparison Of Model Outputs To Teach Alignment, Preferences, And Safety – Critical For Improving Reasoning And User Experience.	Scale AI, Labelbox
Evaluation Datasets	Specialized, Labeled Datasets Built To Rigorously Test Model Performance, Accuracy, Bias, Safety, And Hallucination Rates.	Scale AI, Labelbox
Continuous Improvement	Post-deployment Labeling And Feedback Collection That Fine-tunes Models On Real-world Use Cases And Evolving Data.	Scale AI, Labelbox

DEMOS AI Will Build **Native Language** Infrastructure

Not Translation—true Language-native AI From The Ground Up

- **Hermigas (Spanish)**: 600M speakers—medical, legal, financial AI in 20+ dialects
- **Formigas (Portuguese)**: 260M speakers—Brazilian and European variants
- **Murahy (Ukrainian) & Mrówki (Polish)**: Strategic European languages
- **Each Language = Complete Ecosystem**: Models, datasets, agents, cultural understanding
- **Built By Native Speakers**: Not Silicon Valley translating, but communities creating

Why This **Changes Everything**

Unlocking 6.5B Consumers Who Become Creators

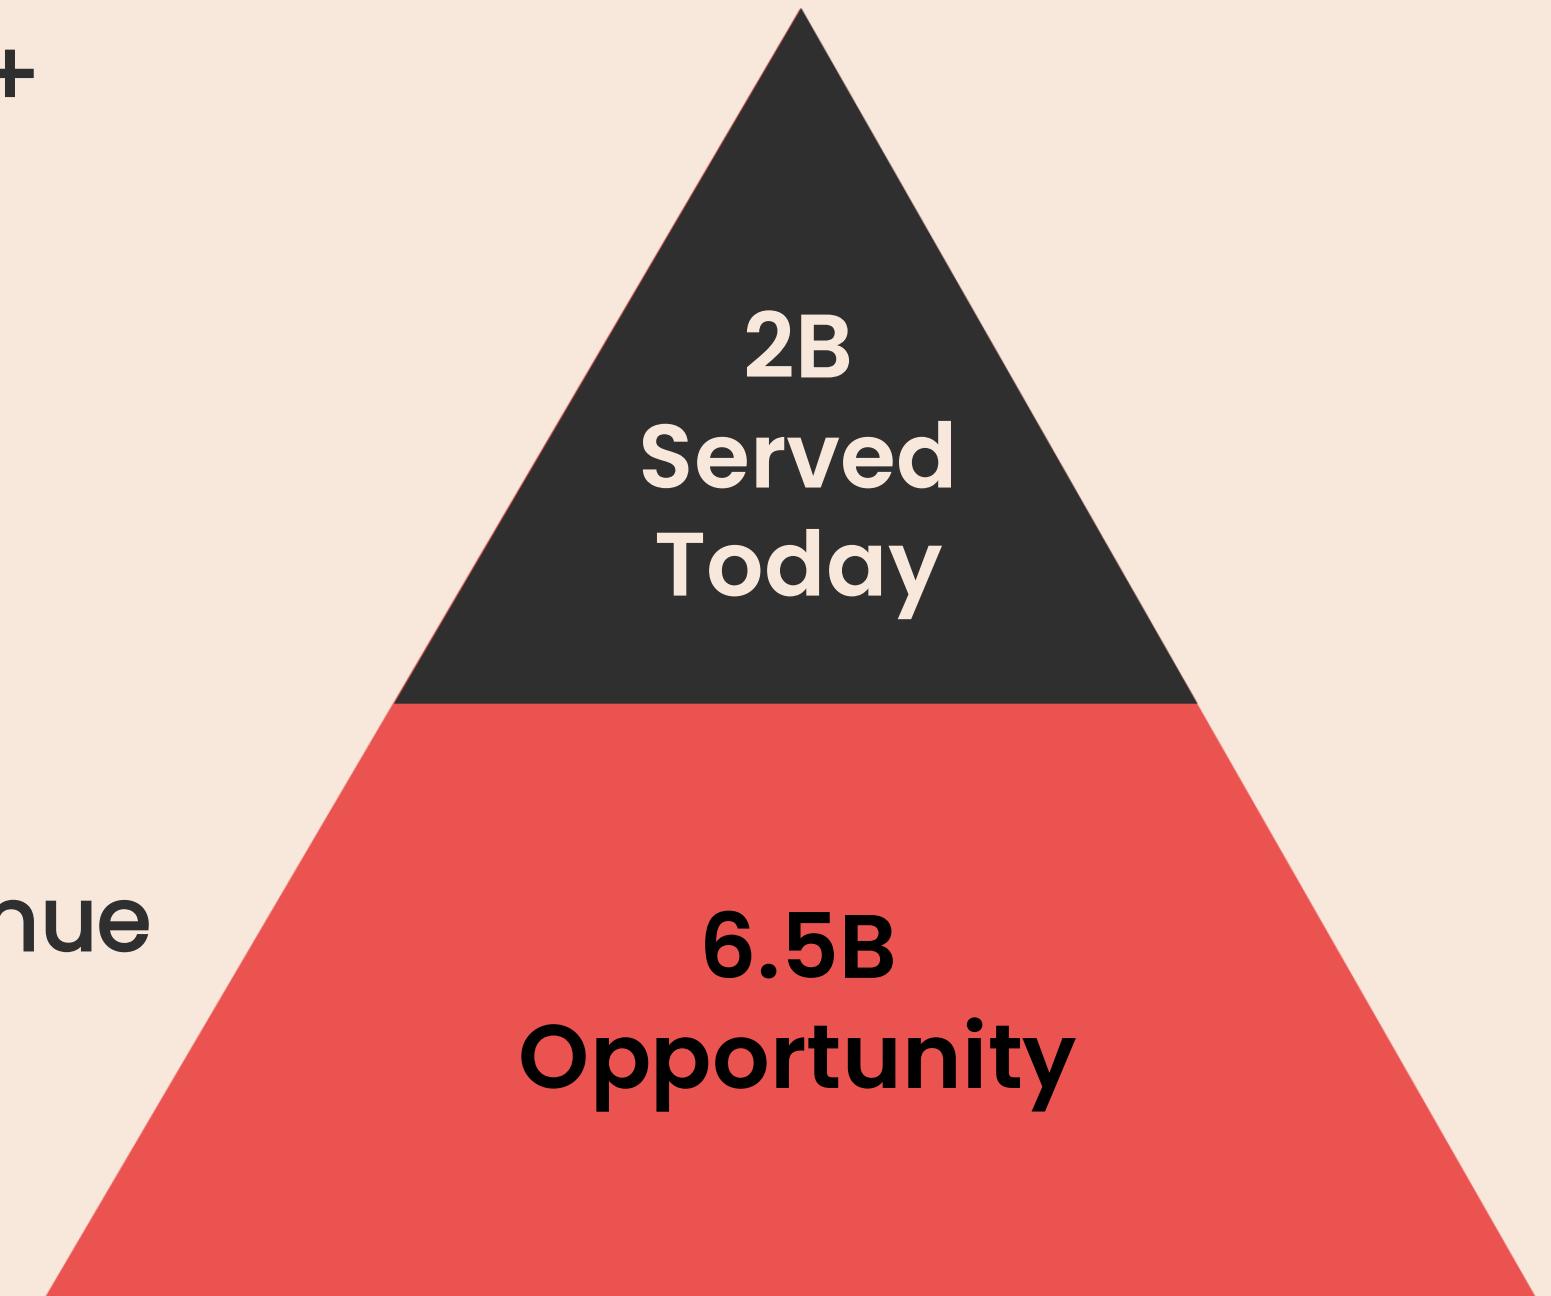
- **Massive New Market:** 6.5B first-time AI users accessible at \$1-25 price points
- **Creator Economy Unleashed:** Local developers build and sell agents, keeping 90% revenue
- **Infrastructure Independent:** Works without internet, during blackouts, in conflicts
- **Sovereignty + Security:** Data never leaves country, defeats adversarial AI
- **Exponential Growth:** Every user can become a creator, every creator attracts more users

The Market

The Largest TAM In AI History

Multi Trillion \$\$ Opportunity: 6.5 Billion Untapped Users

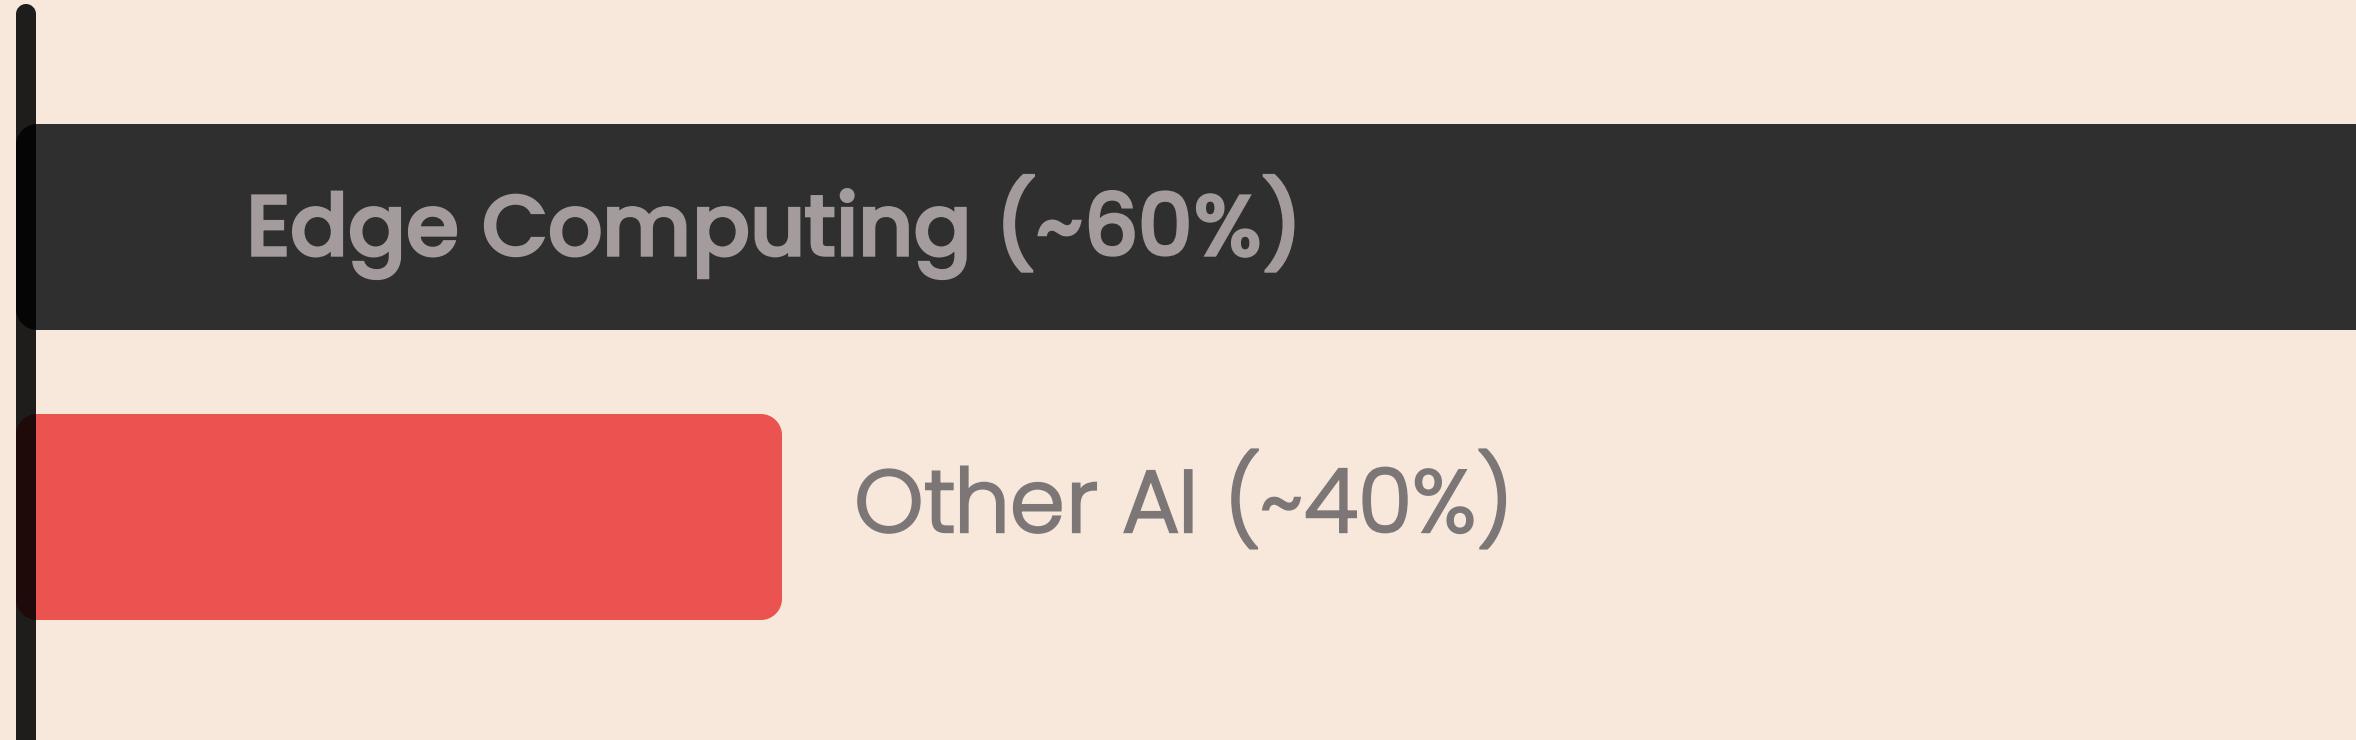
- 6.5B people speaking 7,000+ languages currently locked out of AI
- Immediate addressable market: Spanish (600M) + Portuguese (260M) + Ukrainian/Polish (73M) = 933M people
- AI market growing 40% CAGR—but only for those with infrastructure
- We make the impossible possible: Serving users cloud AI can't reach
- Example: If just 5% of Spanish speakers buy one \$25 agent: \$750M revenue



Edge Is Where The **AI Economy Lives**

By 2028, Nearly 60% Of All Global AI Spending Will Occur On Or Near The Edge -
Where Intelligence Meets The Real World.

- **Global AI Spending:** -632B by 2028 (IDC Worldwide & Gen Ai Spending Aug 19, 2024)
~\$3,497B by 2033 (Grandview Research, Artificial Intelligence Report 2025)
- **Edge Computing Spending:** ~\$380B with 13.8% CAGR by 2028 (IDC, Sep 10, 2024)
- **Other AI (Cloud, Platforms, etc.):** ~\$254 B by 2028
- **Share of Total:** ~60% of AI spending will be on or near the edge, enabling real-time intelligence, autonomy, and local language capabilities.



Edge Computing (~60%)

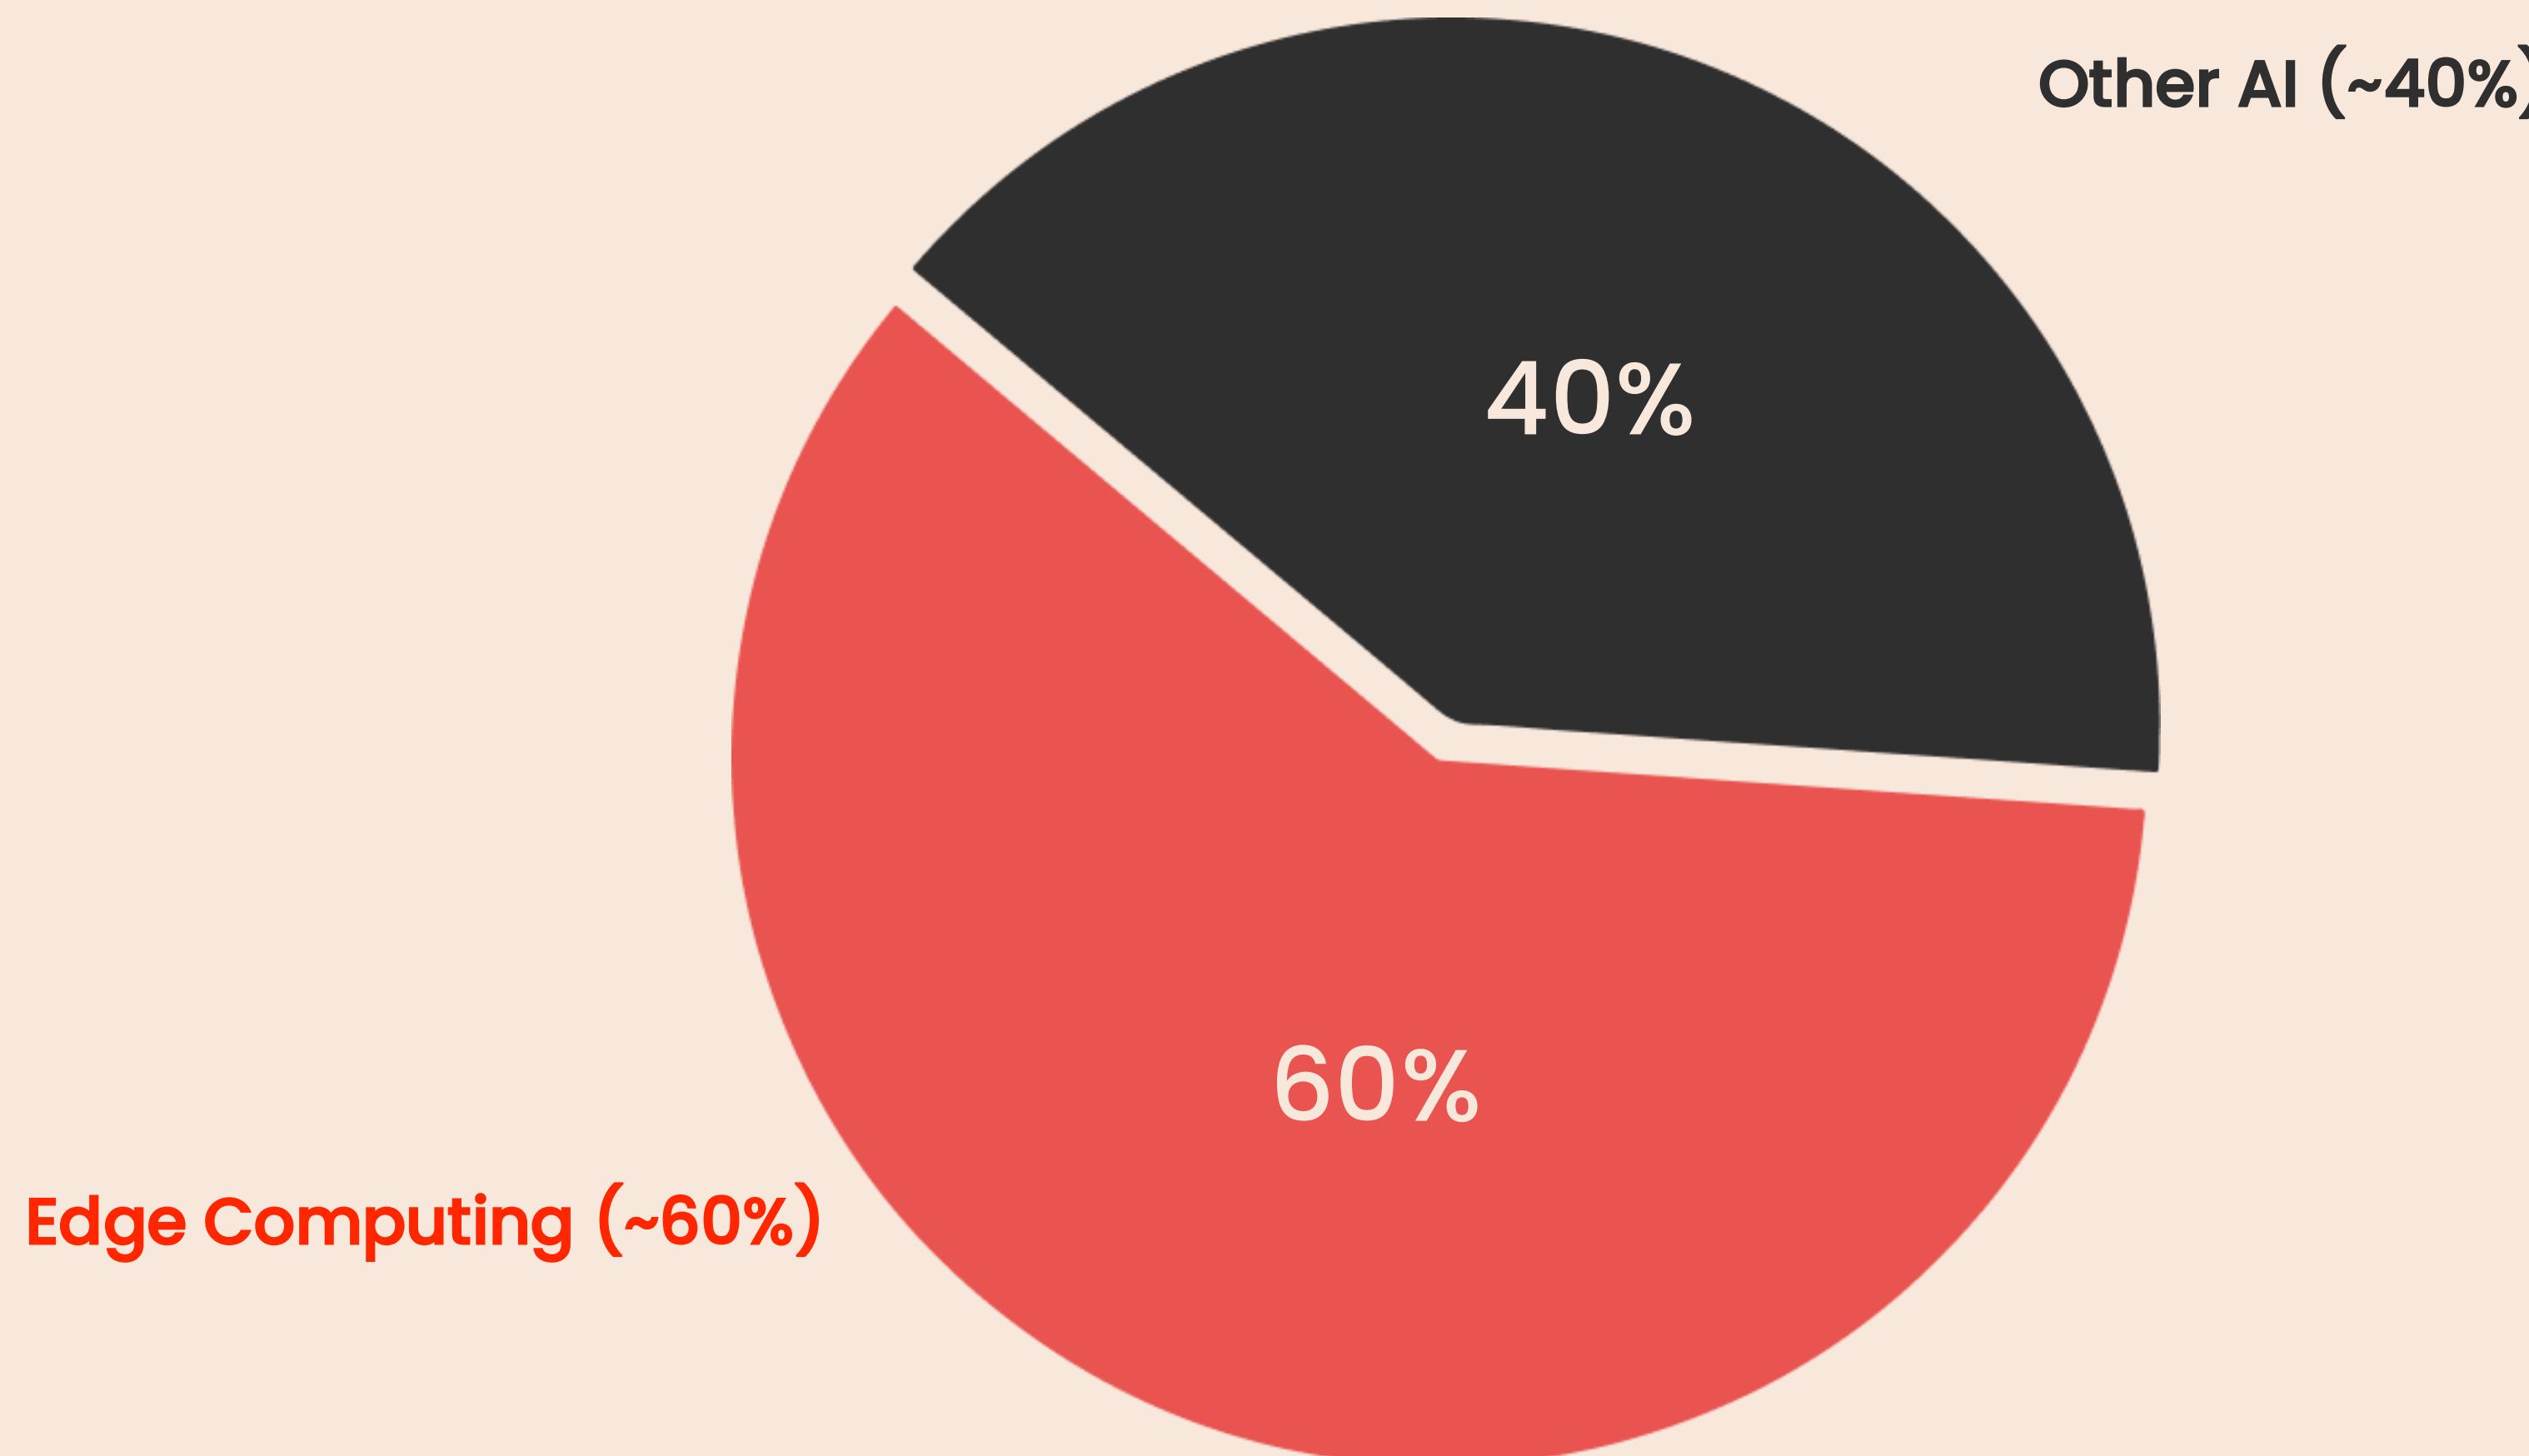
Other AI (~40%)

- IDC, Worldwide AI & GenAI Spending Guide (Aug 19, 2024)
- IDC, Worldwide Edge Spending Forecast (Sep 10, 2024)

Edge Is Where The AI Economy Lives

Global AI Spending 2028: \$632B

By 2028, Nearly 60% Of All Global AI Spending Will Occur On Or Near The Edge -
Where Intelligence Meets The Real World.



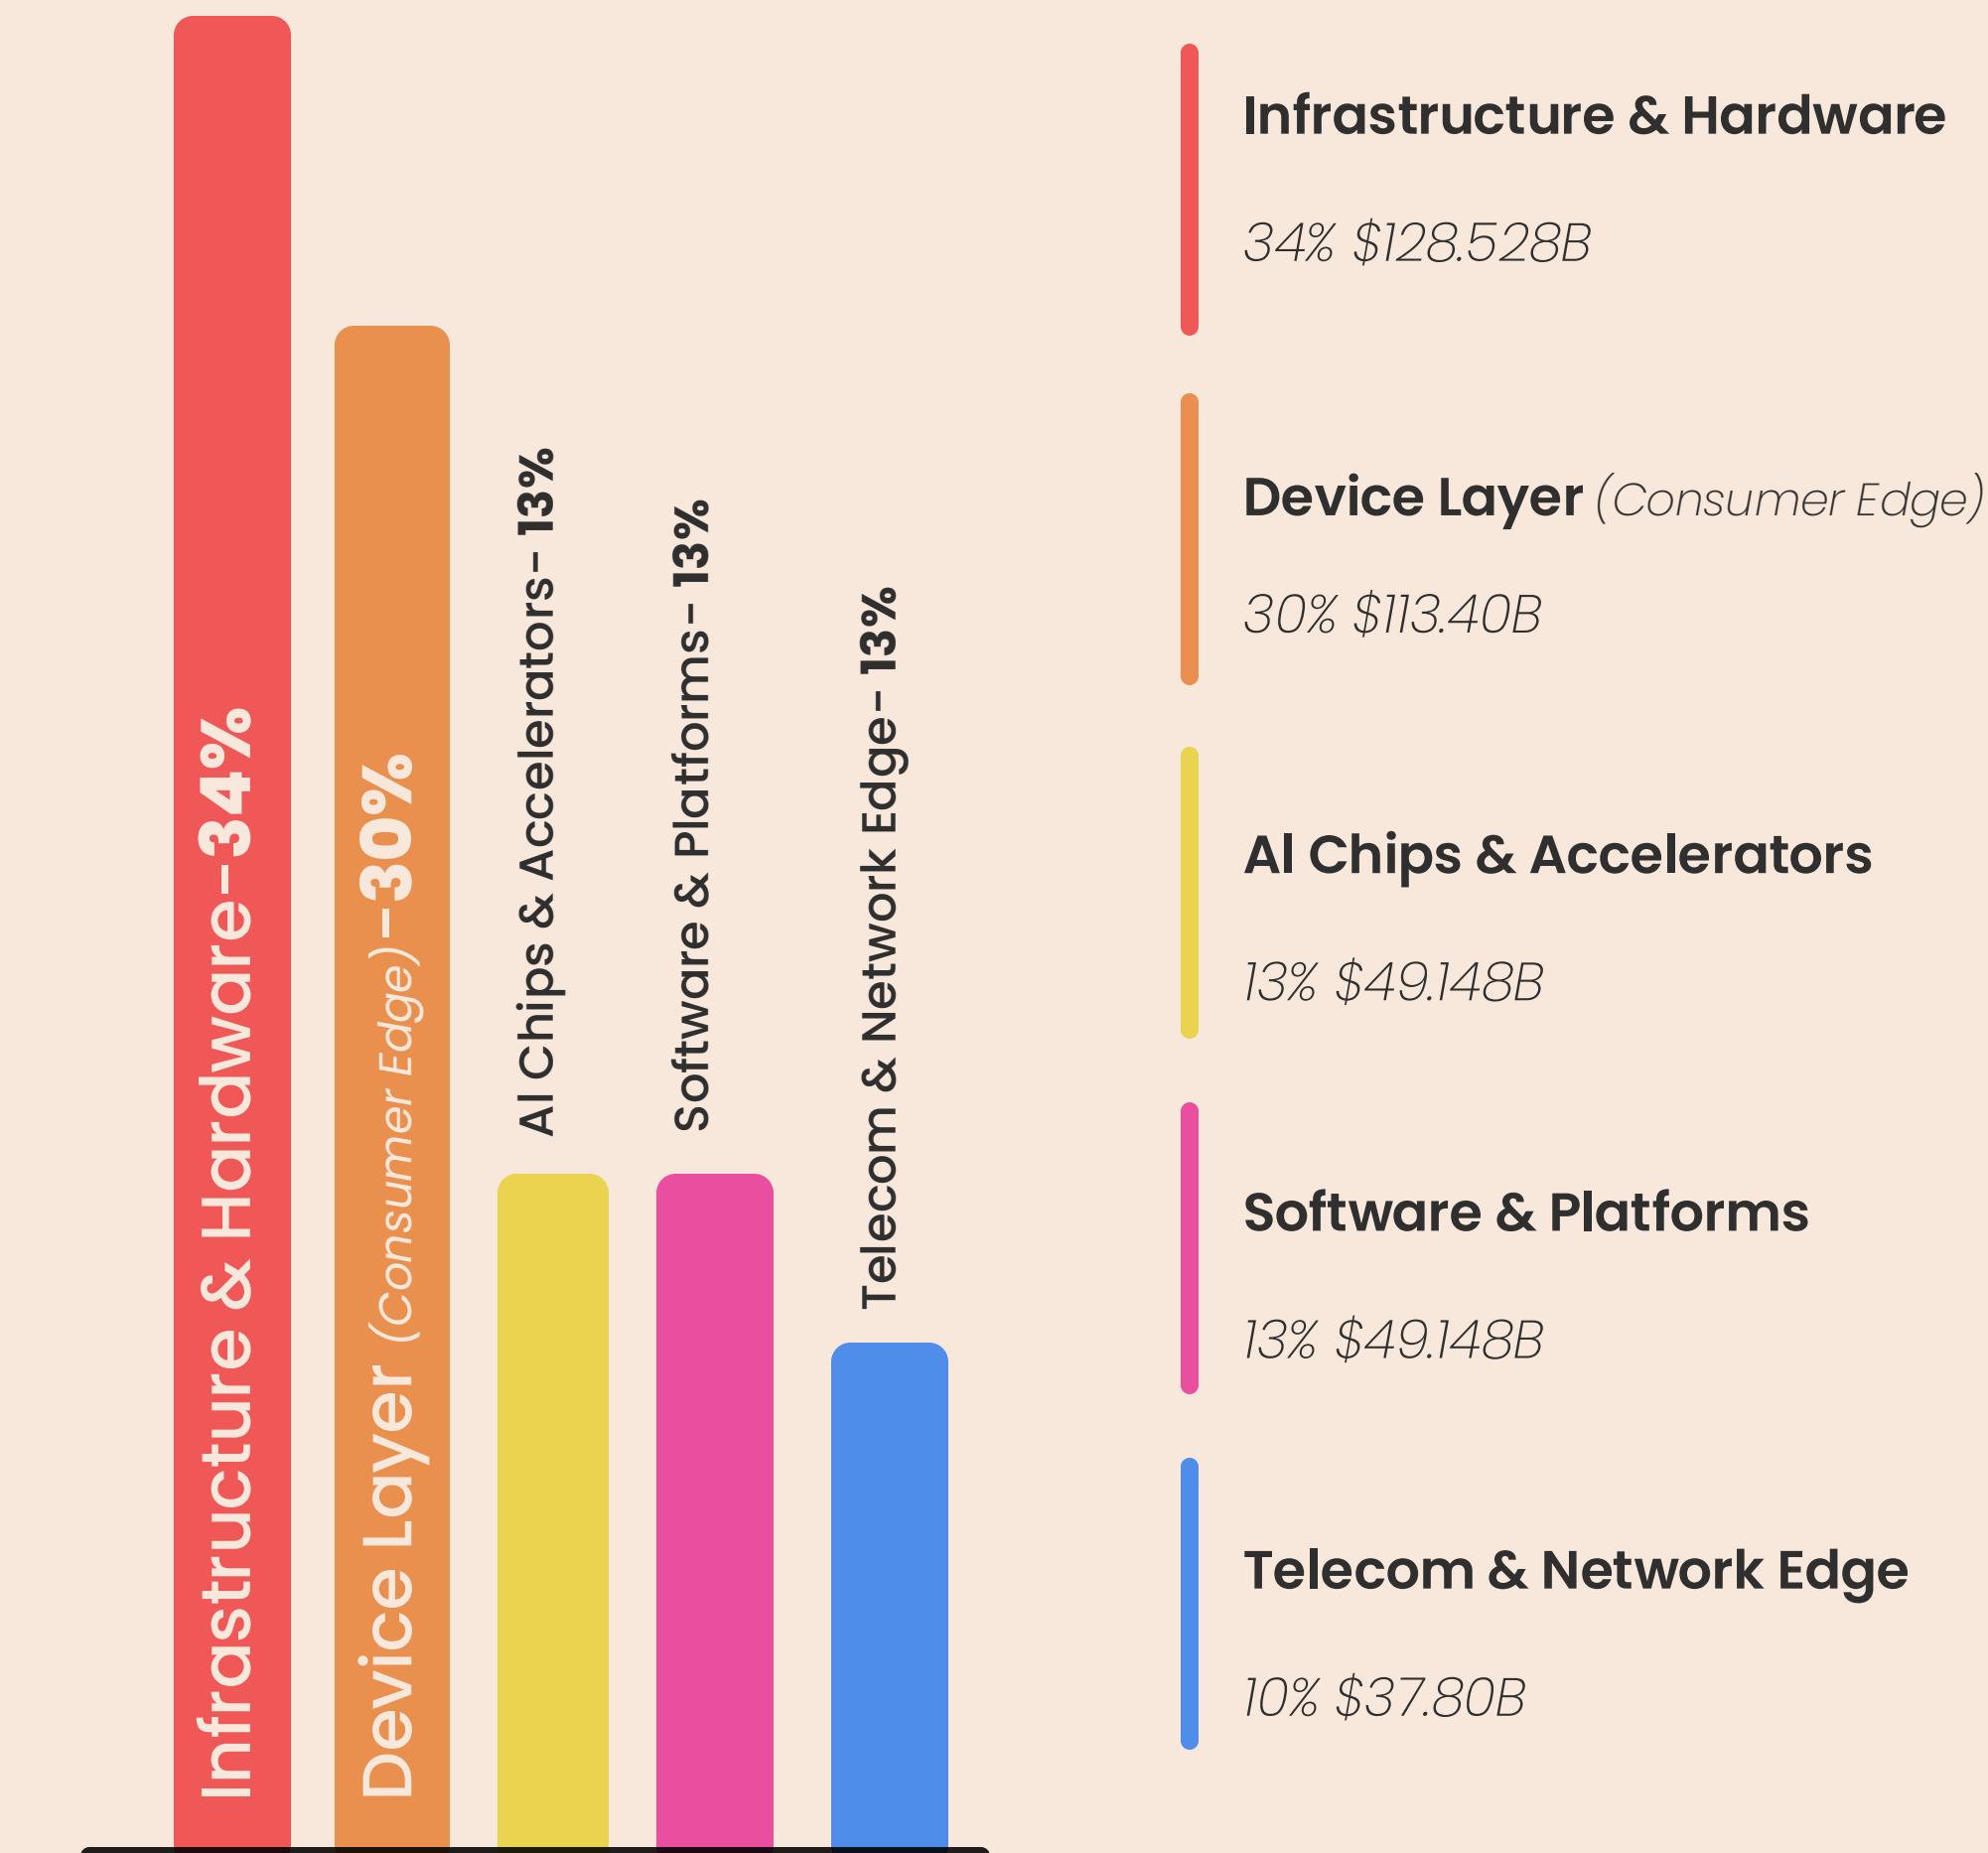
The \$113B Battleground: Where AI Meets The User

Intelligent workloads are shifting to the edge.

Where Intelligence Meets The Real World.

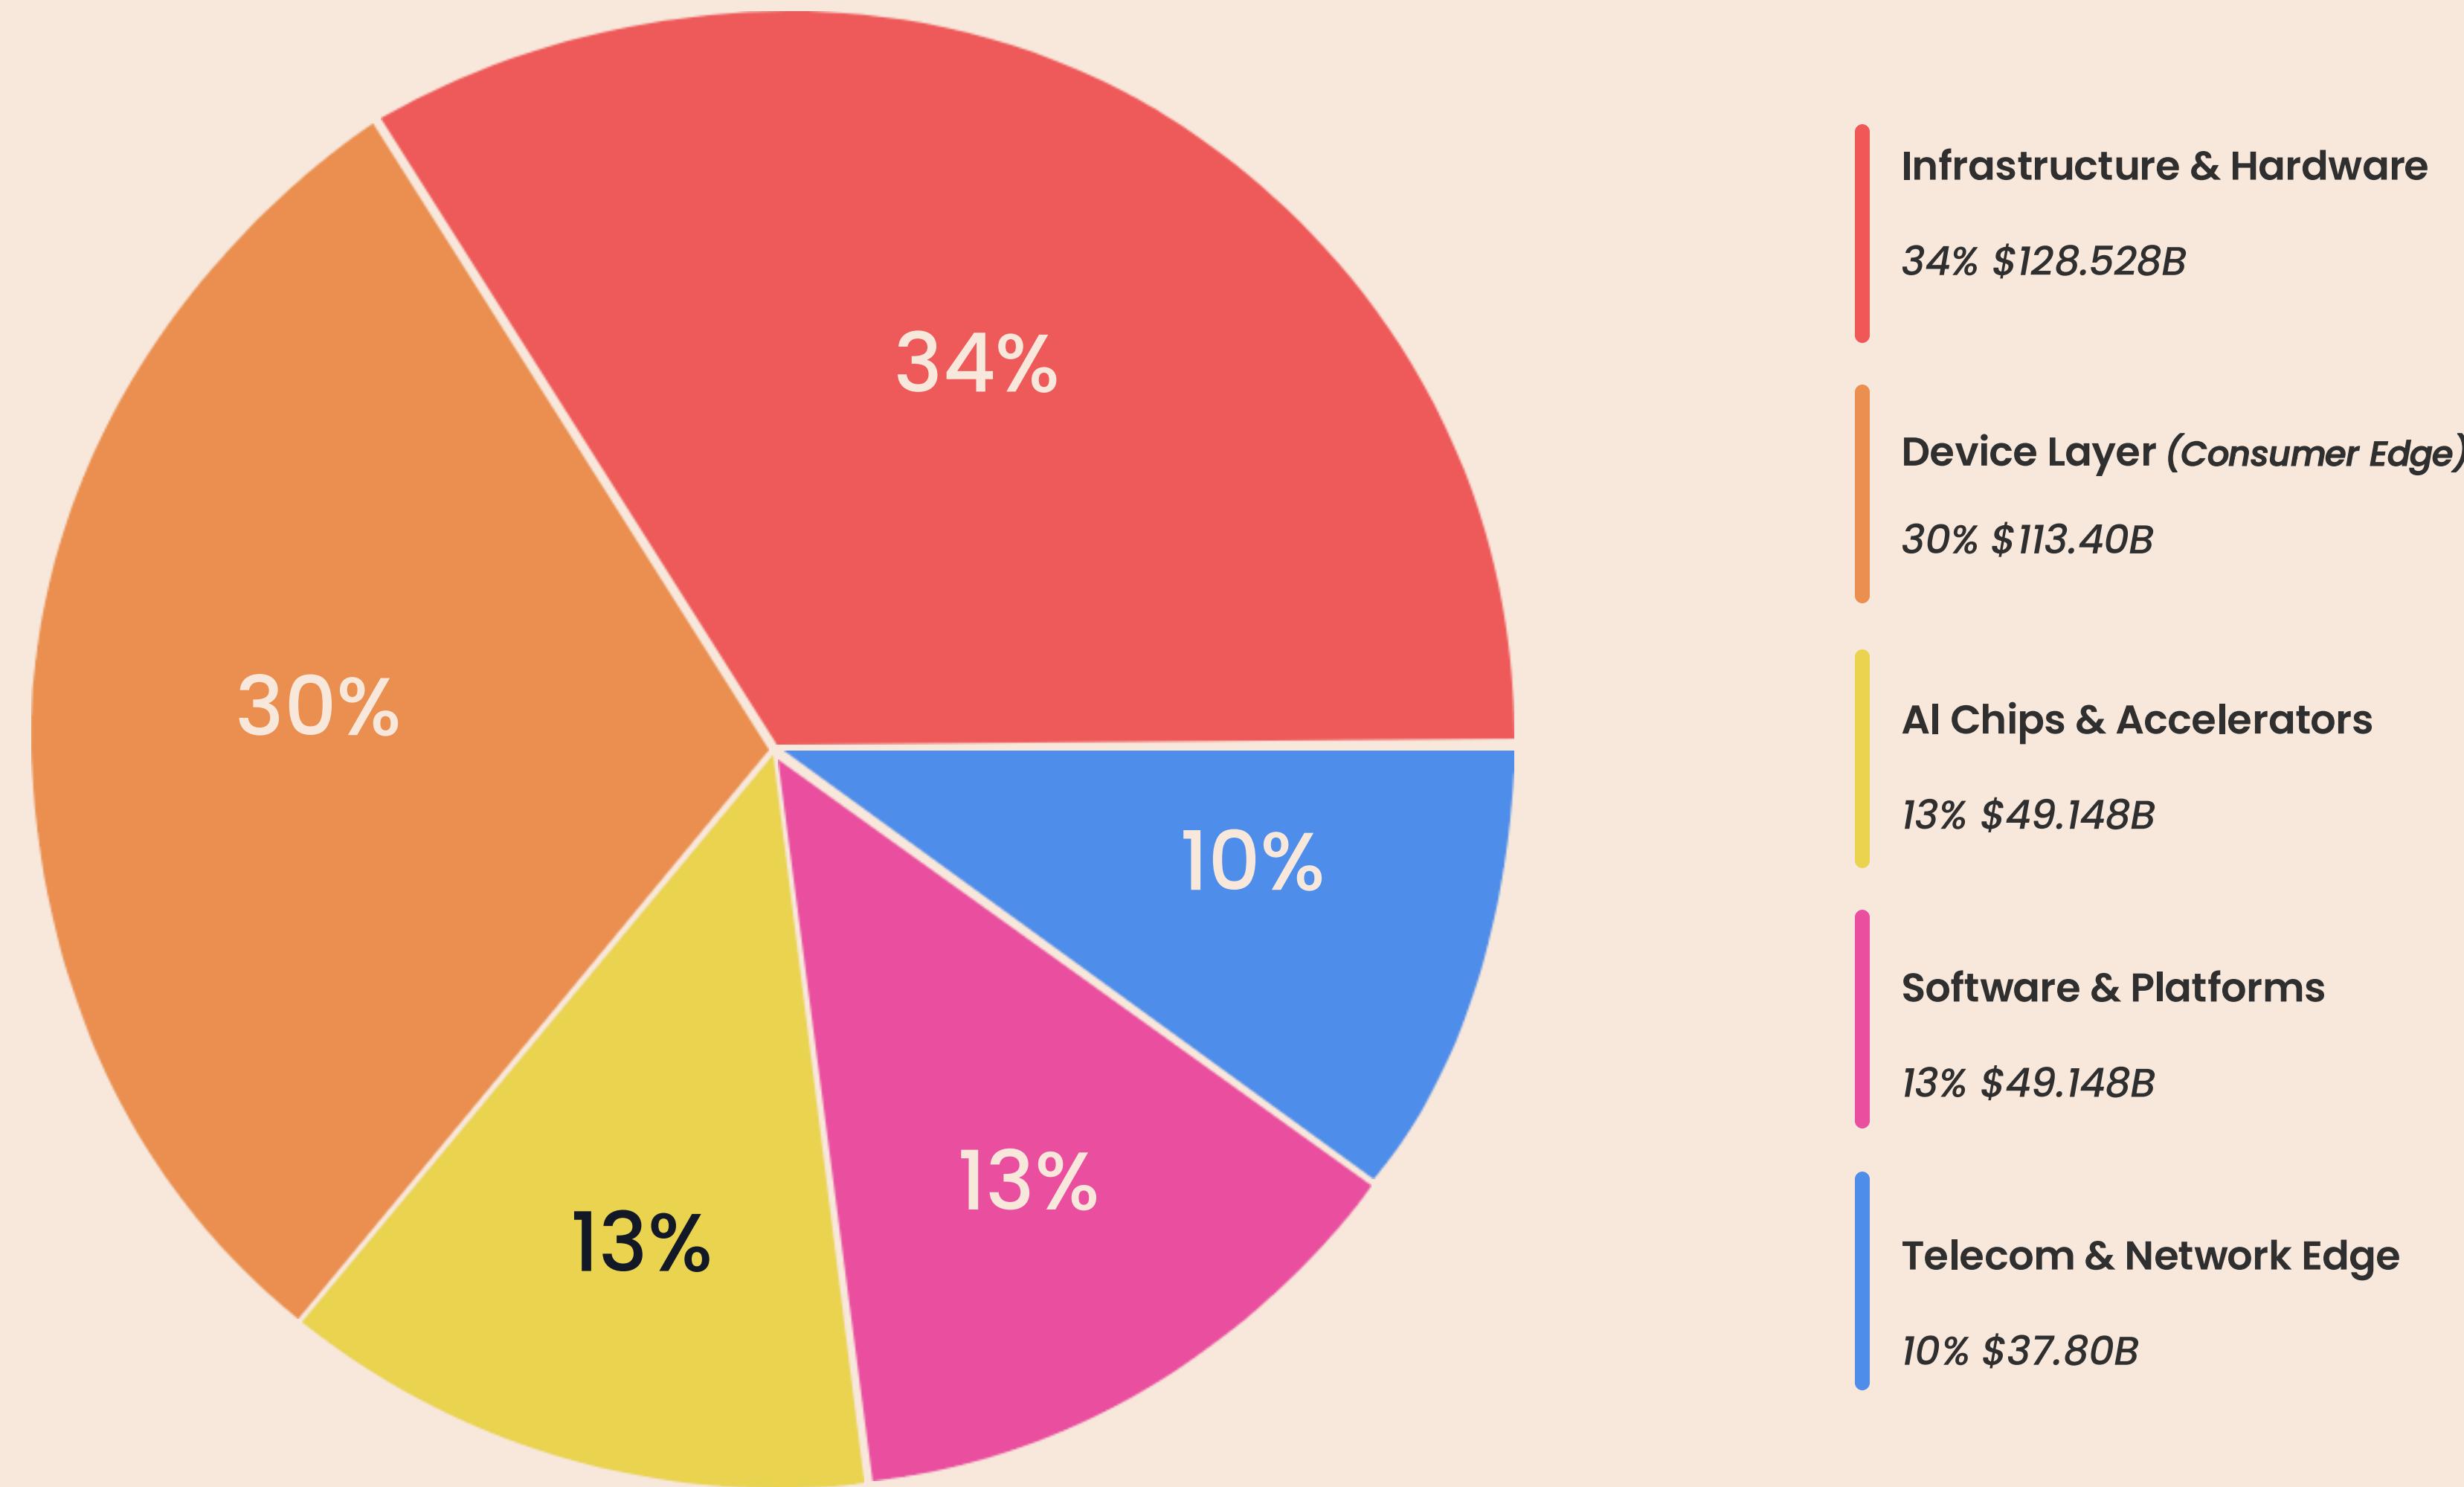
- One-third of all edge spending - \$113 billion - happens inside your pocket, on your desk, and in your home. Smartphones. Laptops. Wearables. AI assistants.
- This is the front line of AI adoption. Where users actually talk to AI. Where copilots, voice agents, and personal assistants must speak the user's language - literally or fail.
- It's the fastest-growing edge segment, exploding alongside GenAI smartphones, AI PCs, and home devices. And it's the most language-sensitive layer of the entire stack.

DEMOS AI LAB sits right at this intersection: on-device intelligence meets multilingual AI. We're powering the next wave of localized, sovereign, user-first experiences - where AI finally speaks everyone's language.



Source: IDC Worldwide Edge Spending Guide (Sept 2024); DEMOS AI LAB analysis
Segment breakdown approximate. Total Edge Computing spend projected at ~\$380B by 2028.

The \$113B Battleground Where AI Meets The User



Edge AI Examples

The world's largest tech companies are pivoting to edge – validating DEMOS' thesis

Apple Intelligence (\$3.4T Valuation)

- Deployed ~3B parameter on-device models across 250M+ devices
- Privacy-first, federated learning architecture
- Zero-cost inference for developers — everything runs locally
- Limitation: Closed ecosystem, premium devices only (\$800+), wealthy markets

Mistral AI (\$14B Valuation)

- Cloud-first company (APIs, Azure, AWS partnerships)
- October 2024 Pivot: Launched Minstral 3B & 8B for edge device
- Built for "local, privacy-first inference" — on-device translation, offline assistants
- Limitation: Edge is secondary; core business remains cloud-based

The Trend: Cloud → Edge Sovereignty

- Google: Gemma edge models | Microsoft: Phi small models | Meta: Llama for edge
- Market: Edge AI reaching ~\$380B by 2028 (13.8% CAGR)

Why It Matters: DEMOS will be the world's first AI research lab exclusively focused on sovereign edge intelligence for all languages, all markets

AI Demand By Region

The Emerging Multilingual Opportunity

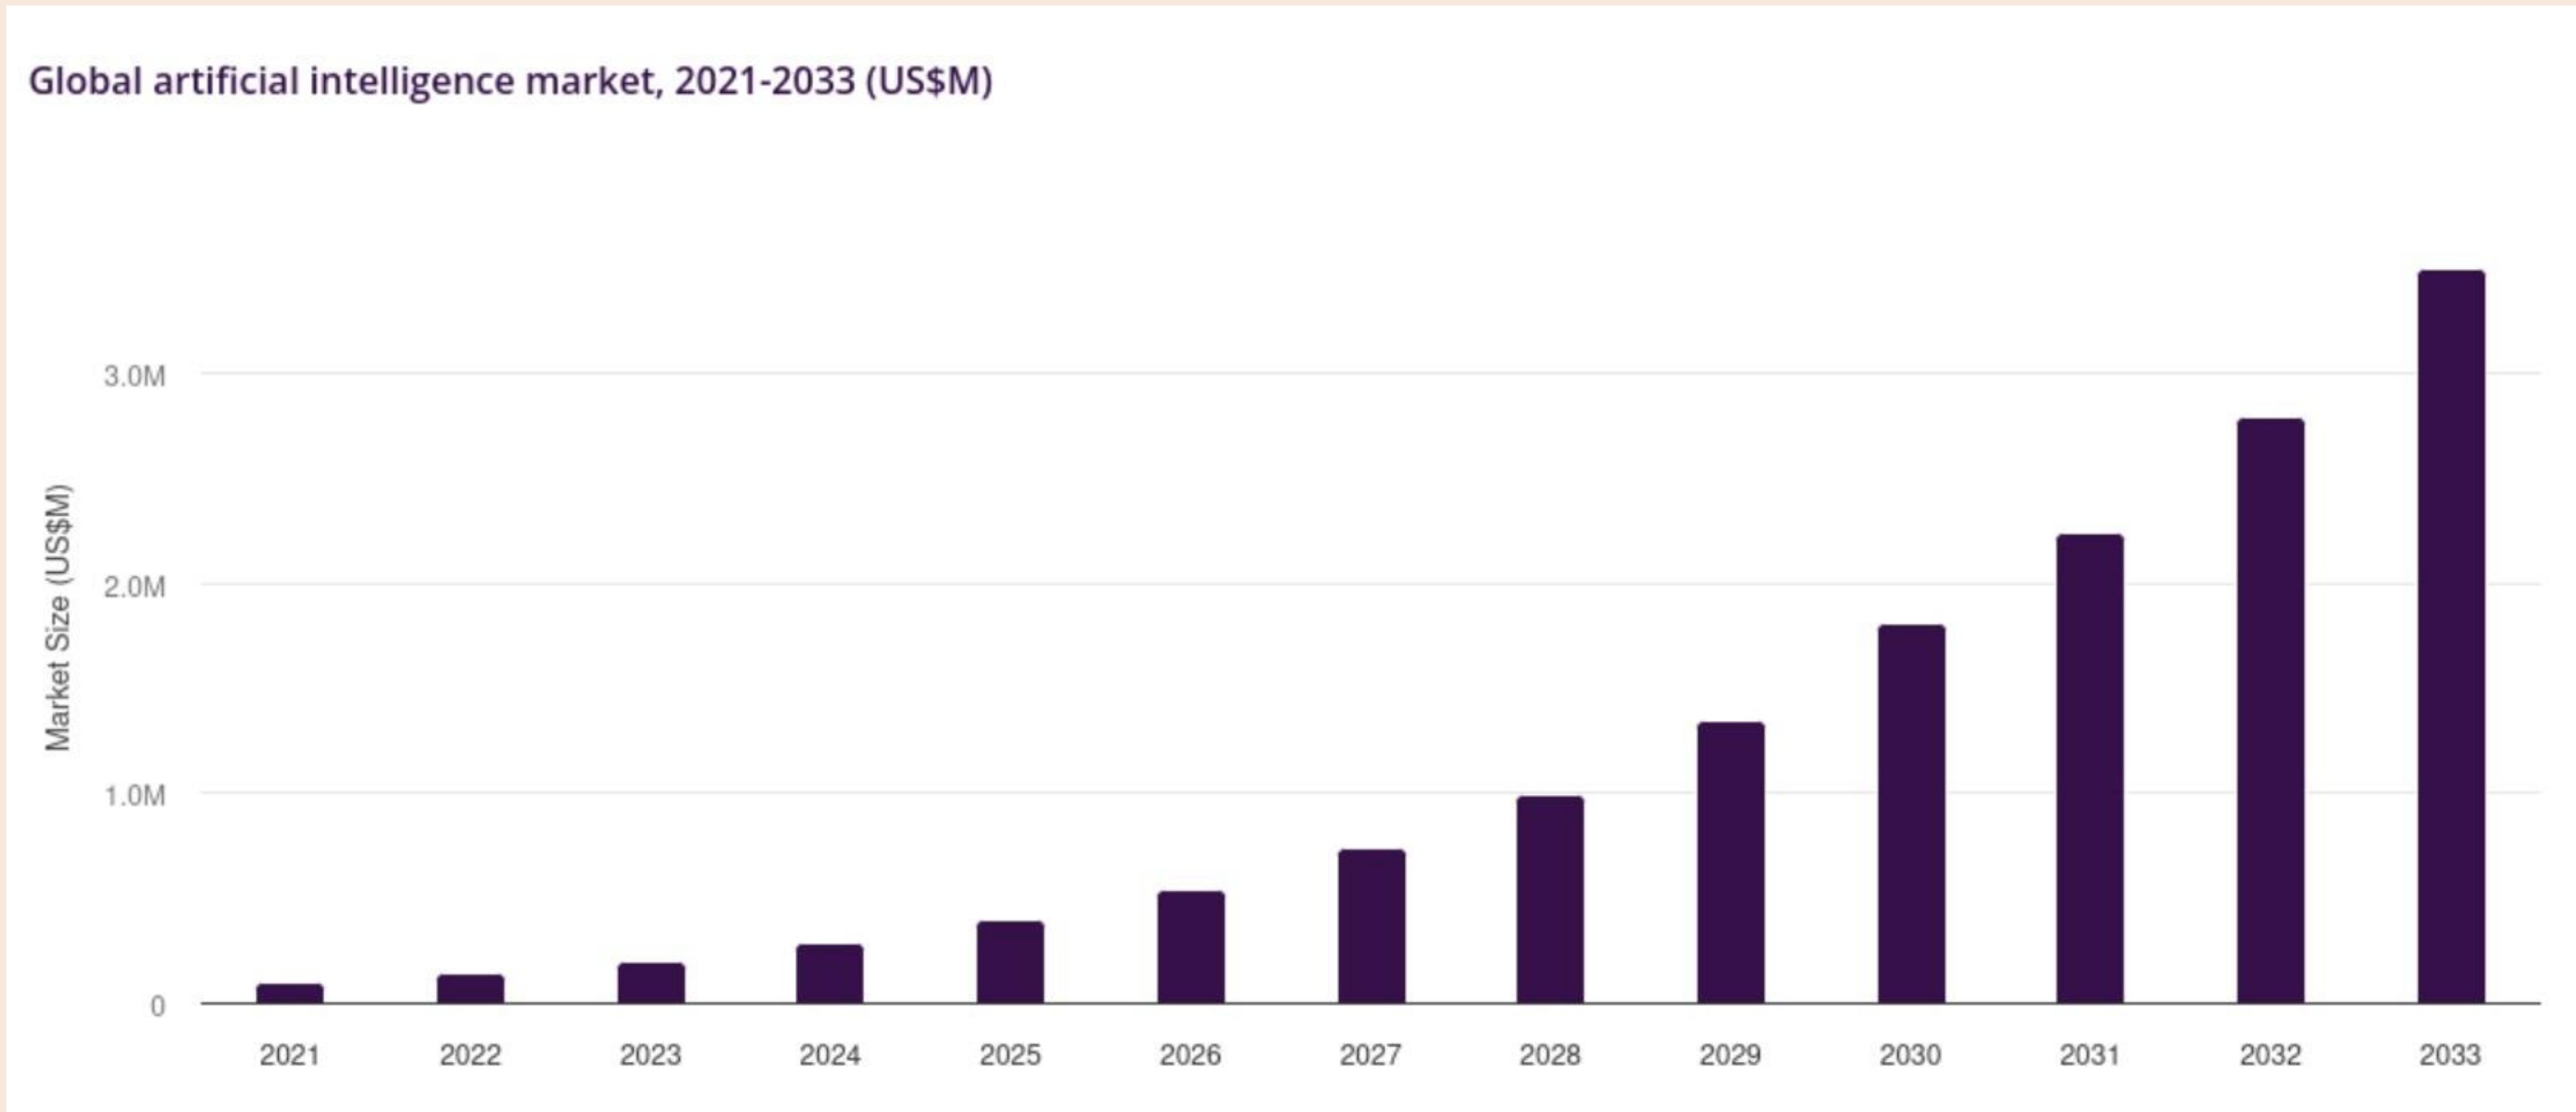
Grandview 2025 Forecast

Region	Projected AI Market Size	Base Year	CAGR	Source (Date)
Global	US \$3,497B By 2033	US \$279 B ₍₂₀₂₄₎	31.5%	Grandview Research- Worldwide AI Spending (2025)
North America	US \$1,026B By 2033	US \$101 B ₍₂₀₂₄₎	28.4%	Grandview Research- Worldwide AI Spending (2025)
Europe	US \$754B By 2033	US \$69 B ₍₂₀₂₄₎	29.6%	Grandview Research- Worldwide AI Spending (2025)
Asia Pacific	US \$1,227B By 2033	US \$71 B ₍₂₀₂₄₎	36.2%	Grandview Research- Worldwide AI Spending (2025)
Latin America	US \$201B → By 2033	US \$21.1 B ₍₂₀₂₄₎	27.5%	Grandview Research- Worldwide AI Spending (2025)
Middle East & Africa	US \$289B By 2033	US \$17 B ₍₂₀₂₄₎	36%	Grandview Research- Worldwide AI Spending (2025)

AI Demand By Region

Worldwide

Grandview Research 2025 Forecast

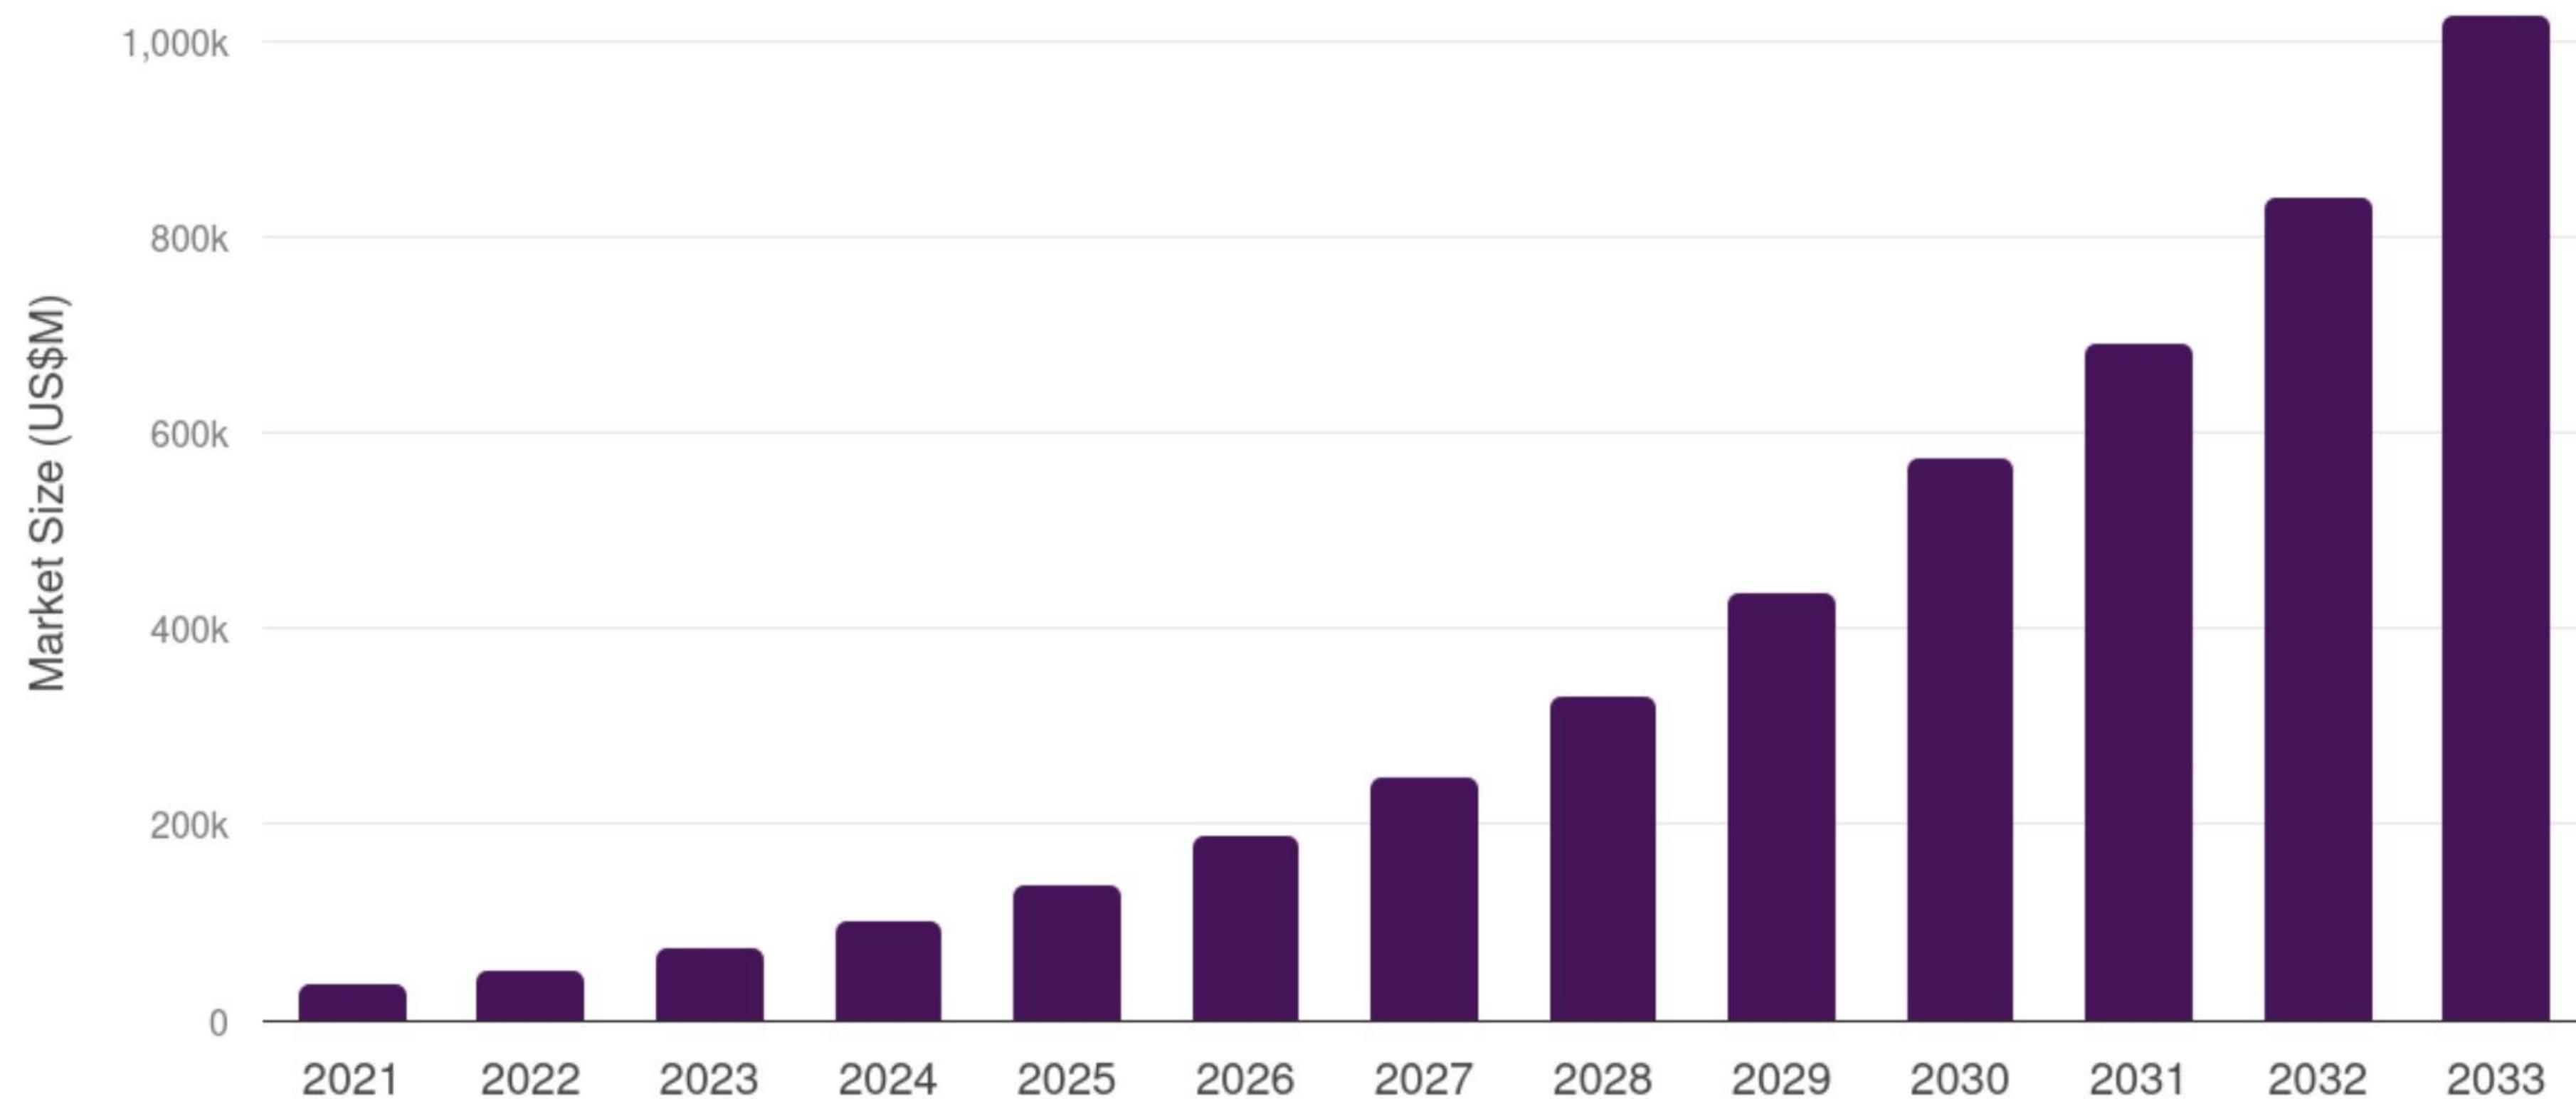


AI Demand By Region

North America

Grandview Research 2025 Forecast

North America artificial intelligence market, 2021-2033



<https://www.grandviewresearch.com/horizon/outlook/artificial-intelligence-market/north-america>

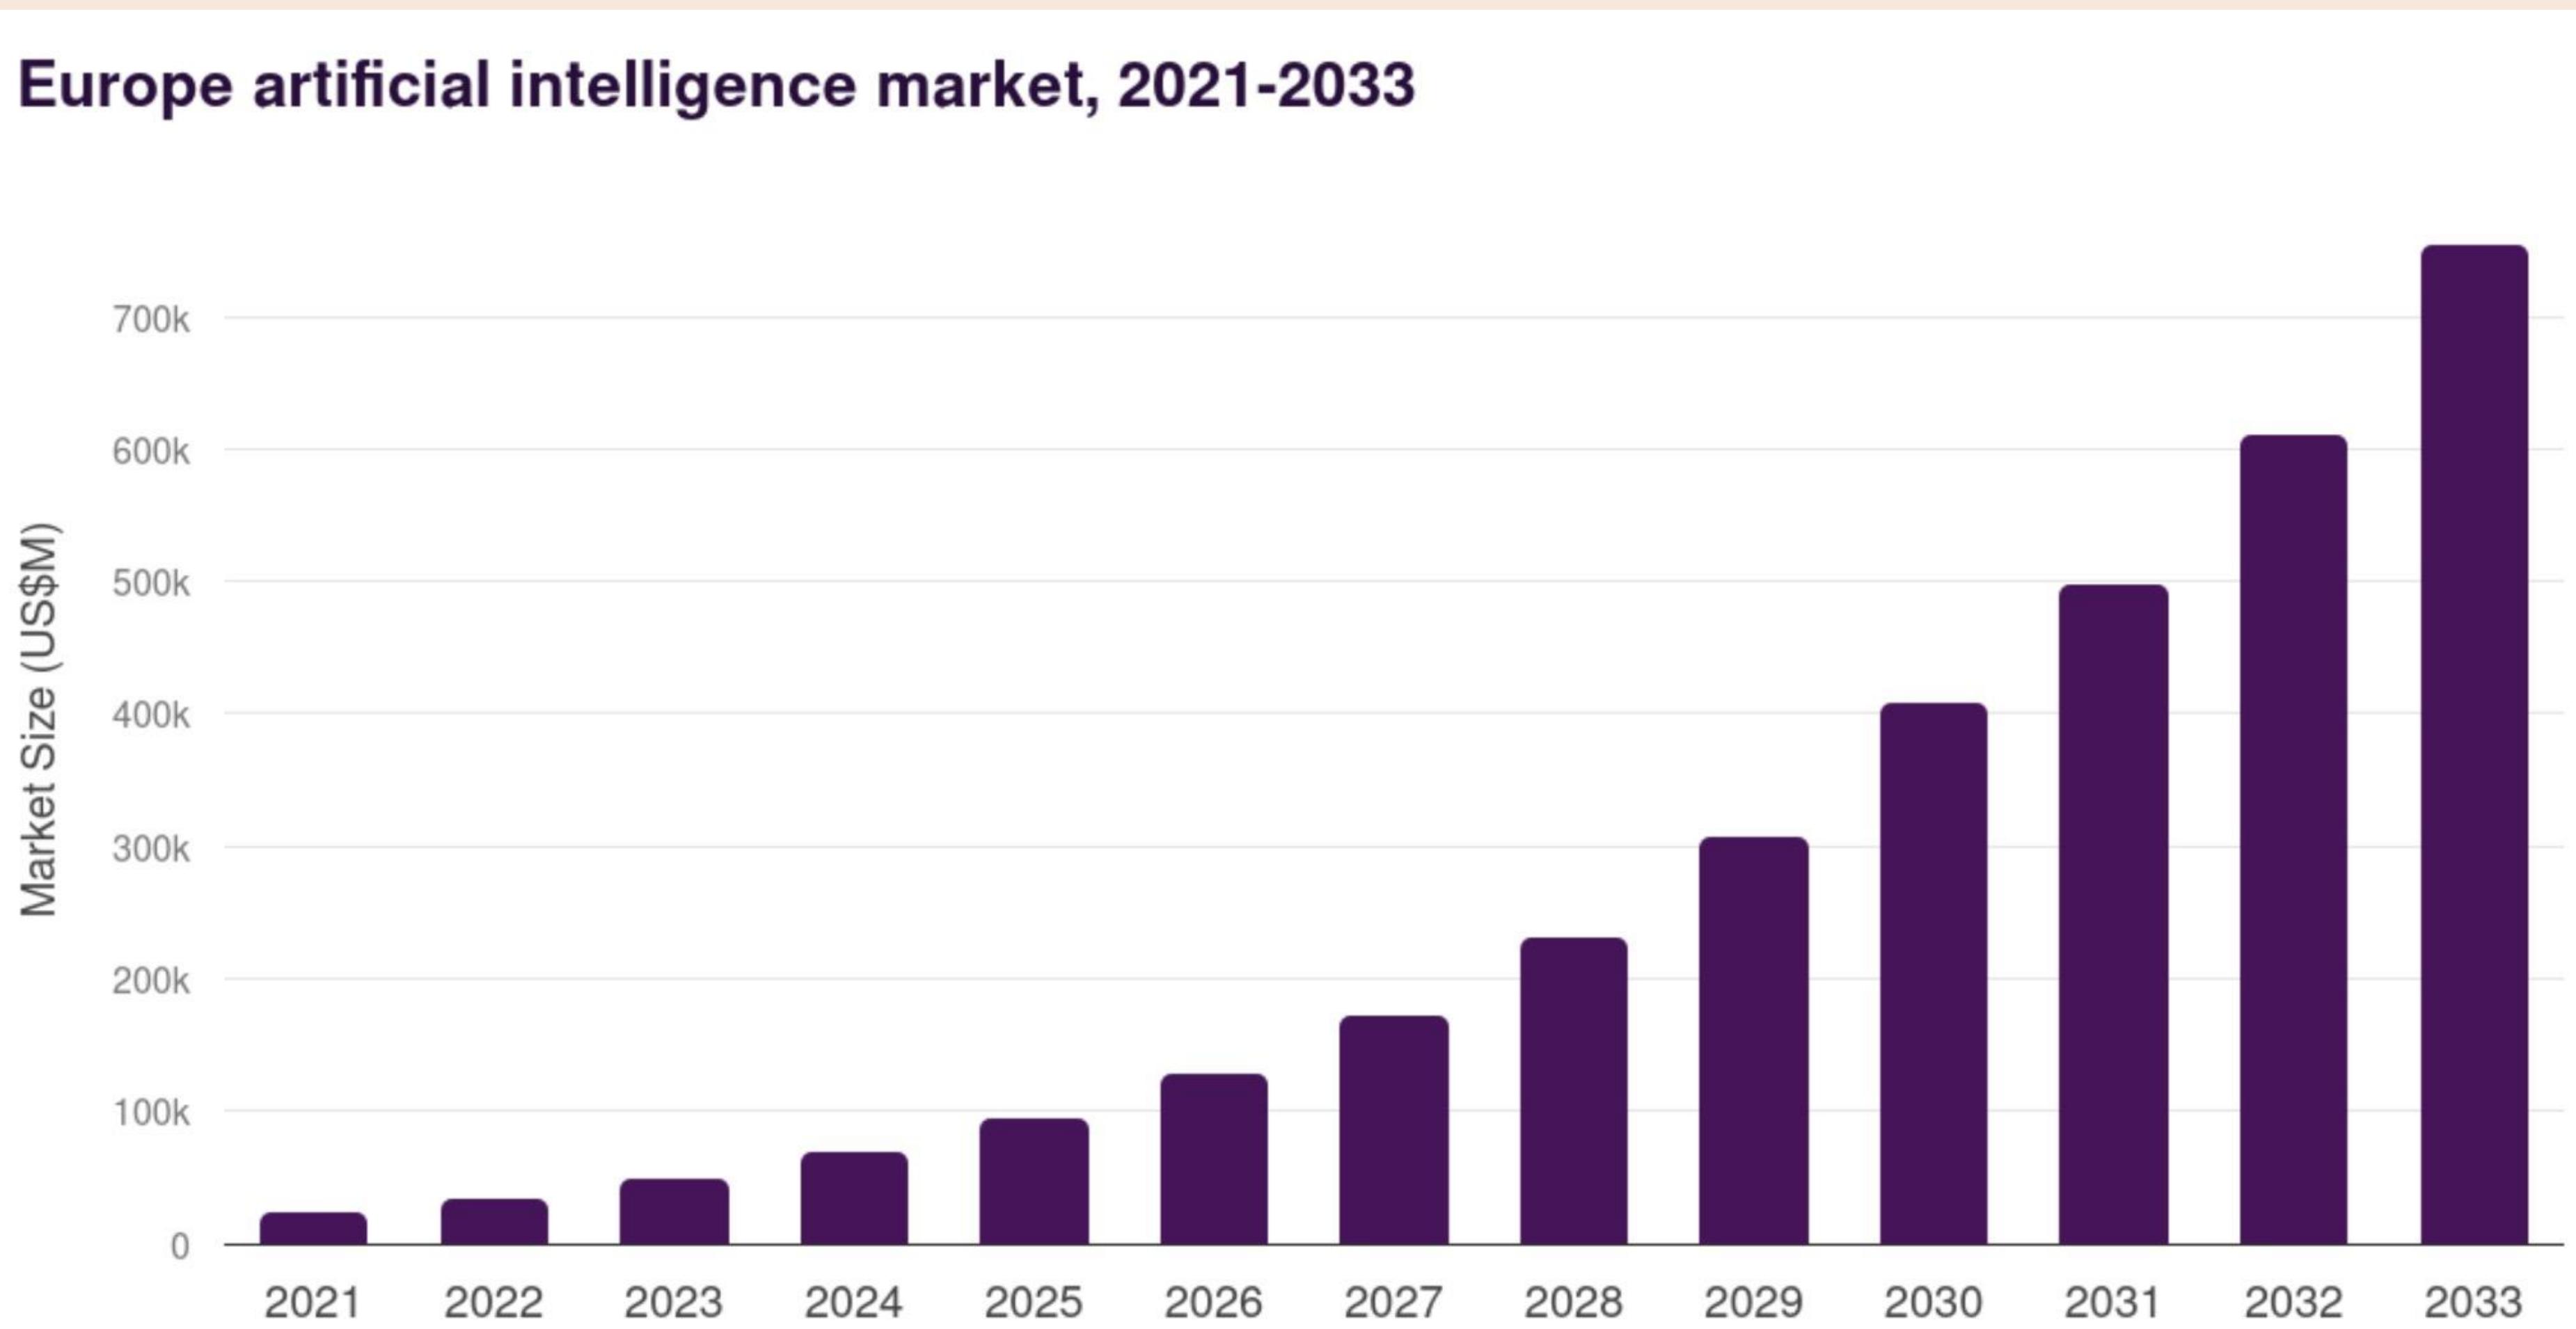


AI Demand By Region

Europe

Grandview Research 2025 Forecast

Europe artificial intelligence market, 2021-2033



<https://www.grandviewresearch.com/horizon/outlook/artificial-intelligence-market/europe>

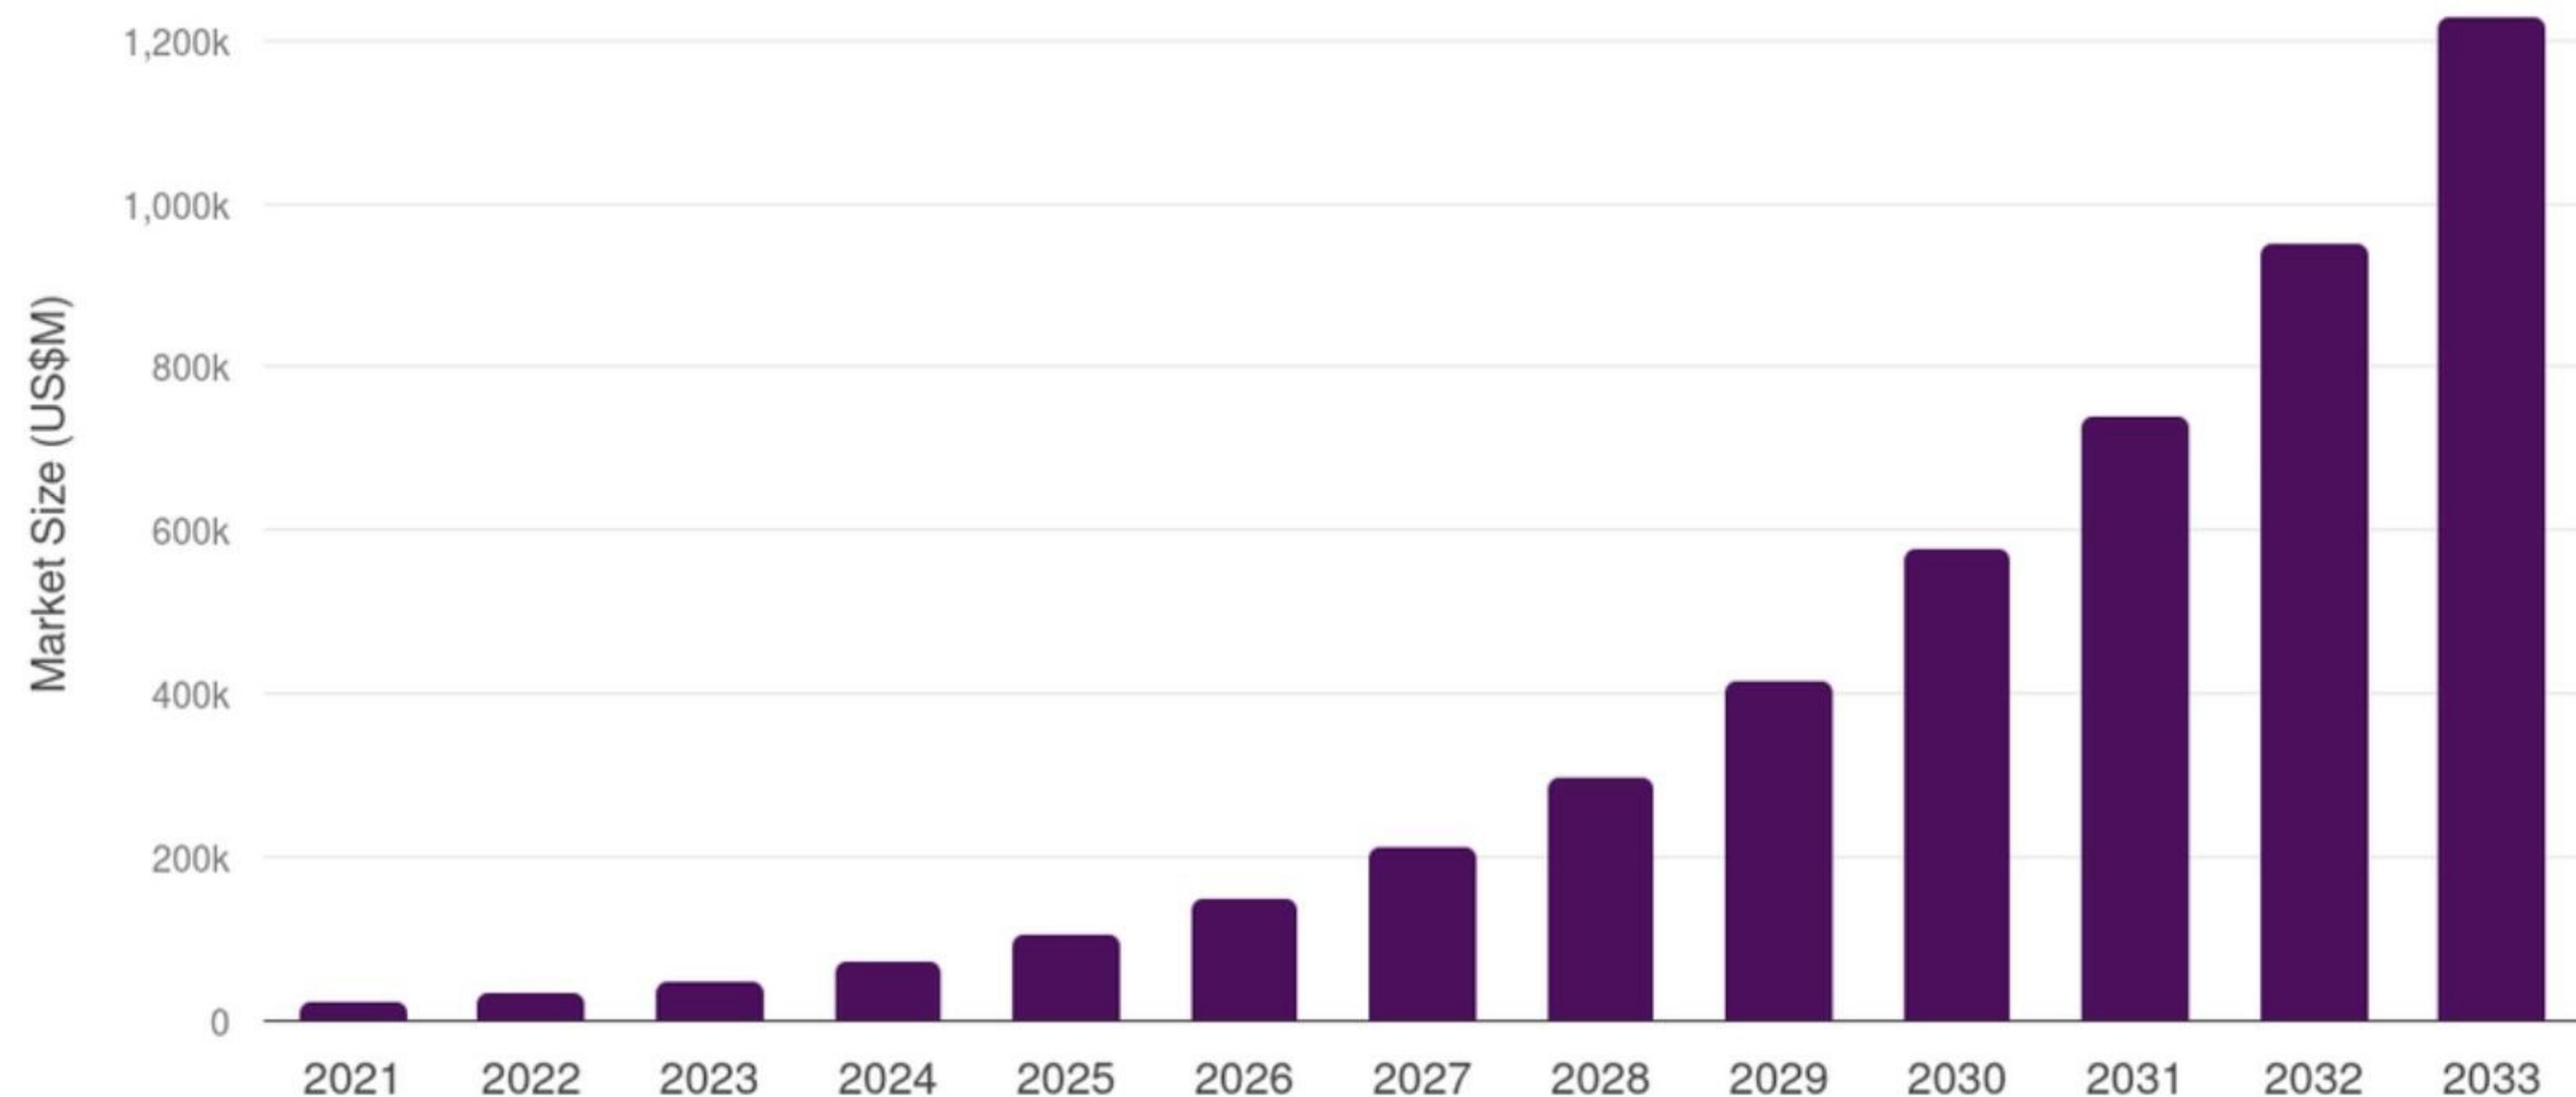


AI Demand By Region

Asia Pacific

Grandview Research 2025 Forecast

Asia Pacific artificial intelligence market, 2021-2033



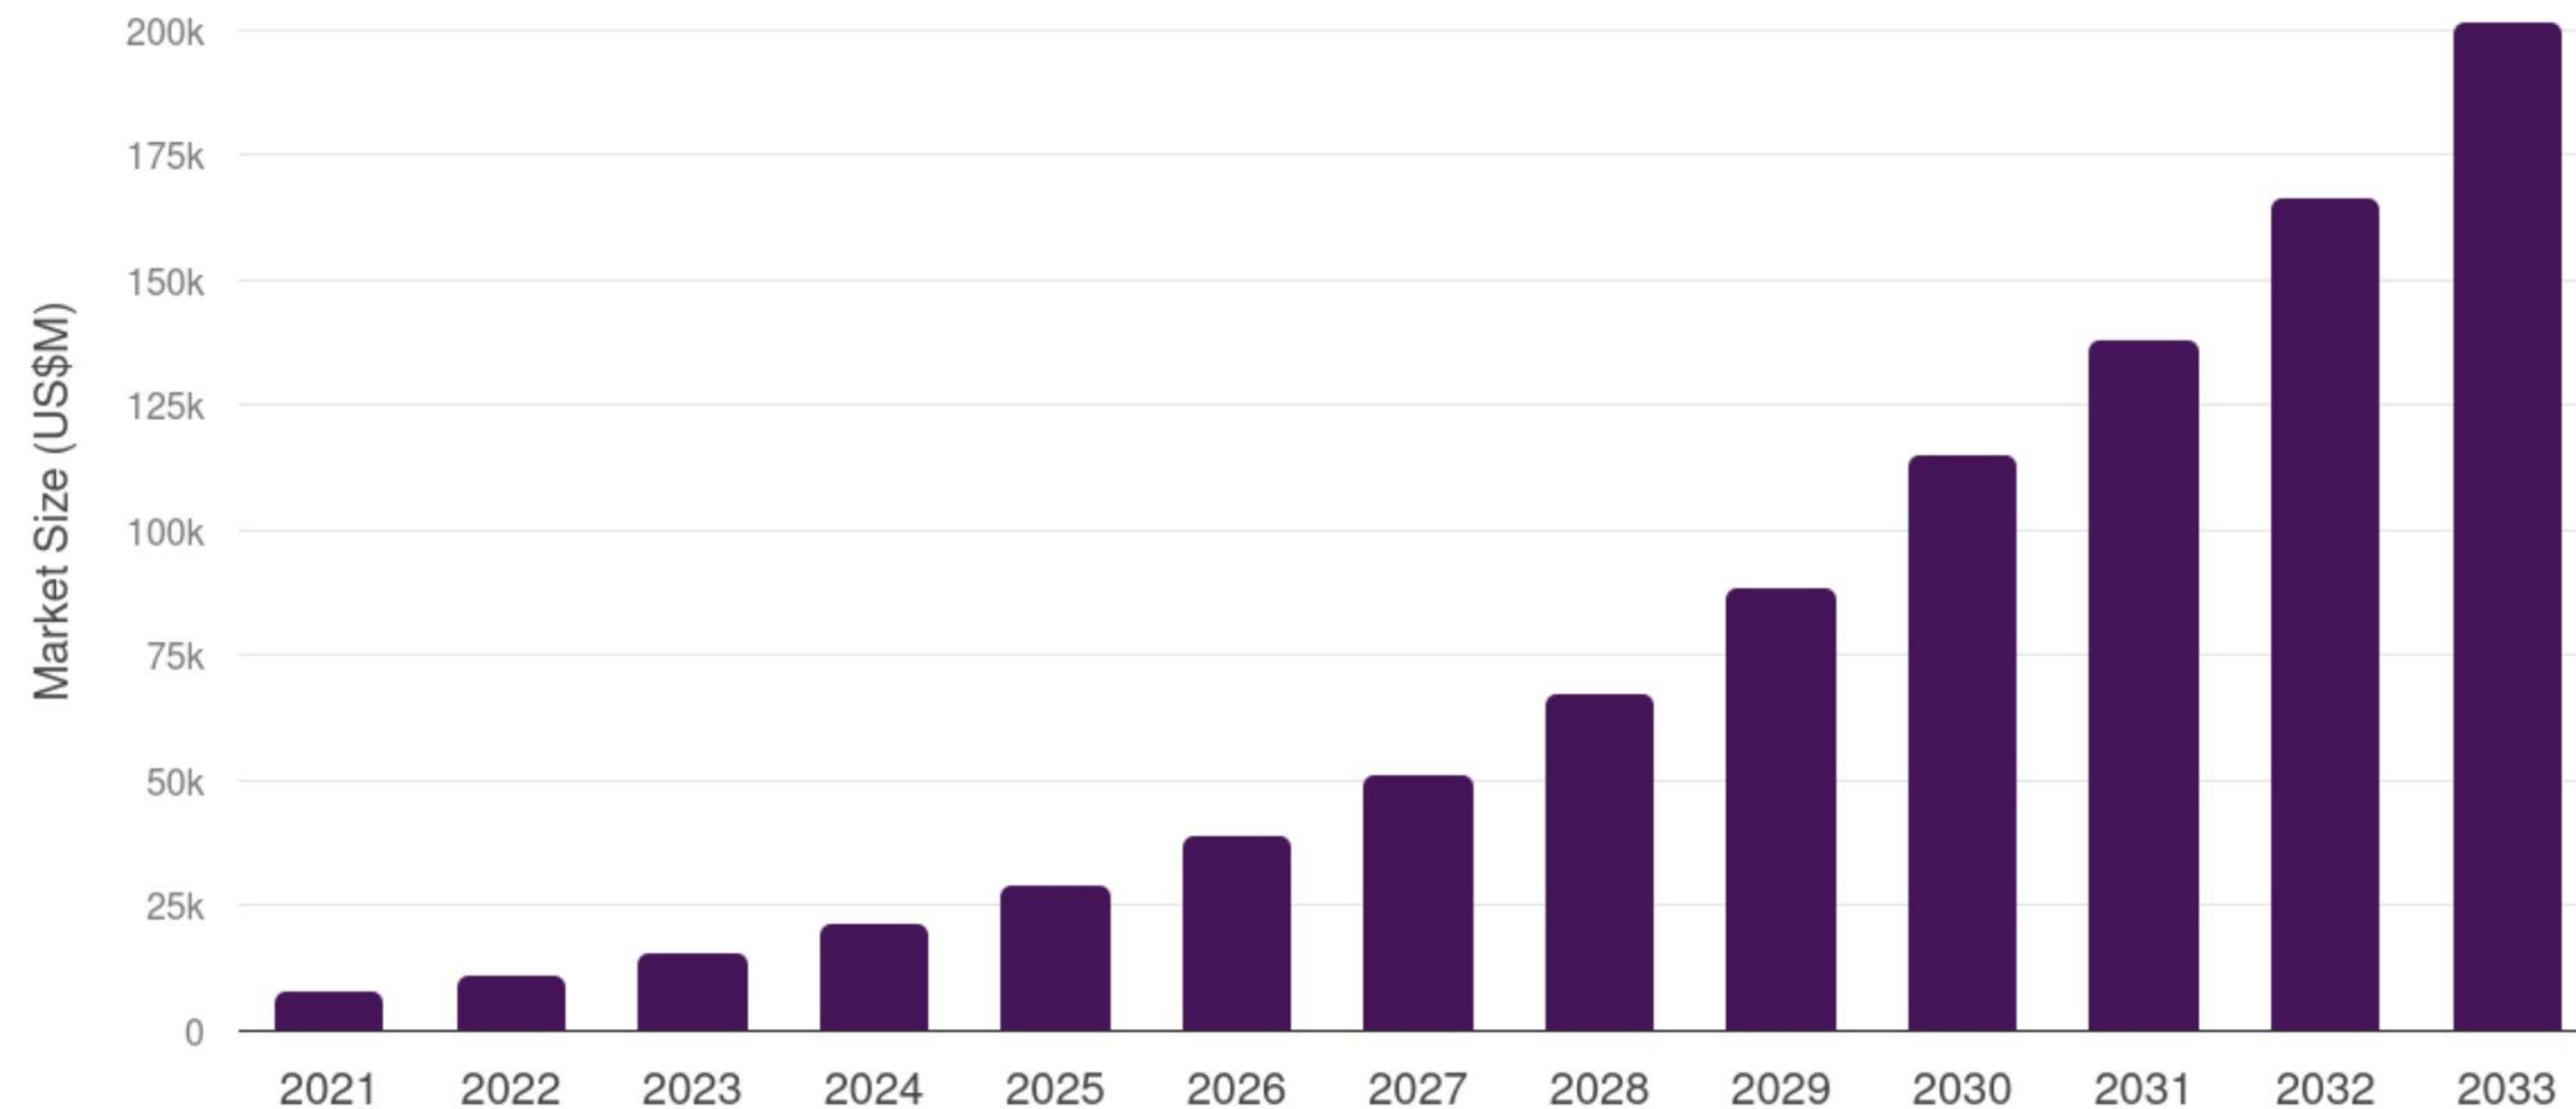
<https://www.grandviewresearch.com/horizon/outlook/artificial-intelligence-market/asia-pacific>

AI Demand By Region

Latin America

Grandview Research 2025 Forecast

Latin America artificial intelligence market, 2021-2033



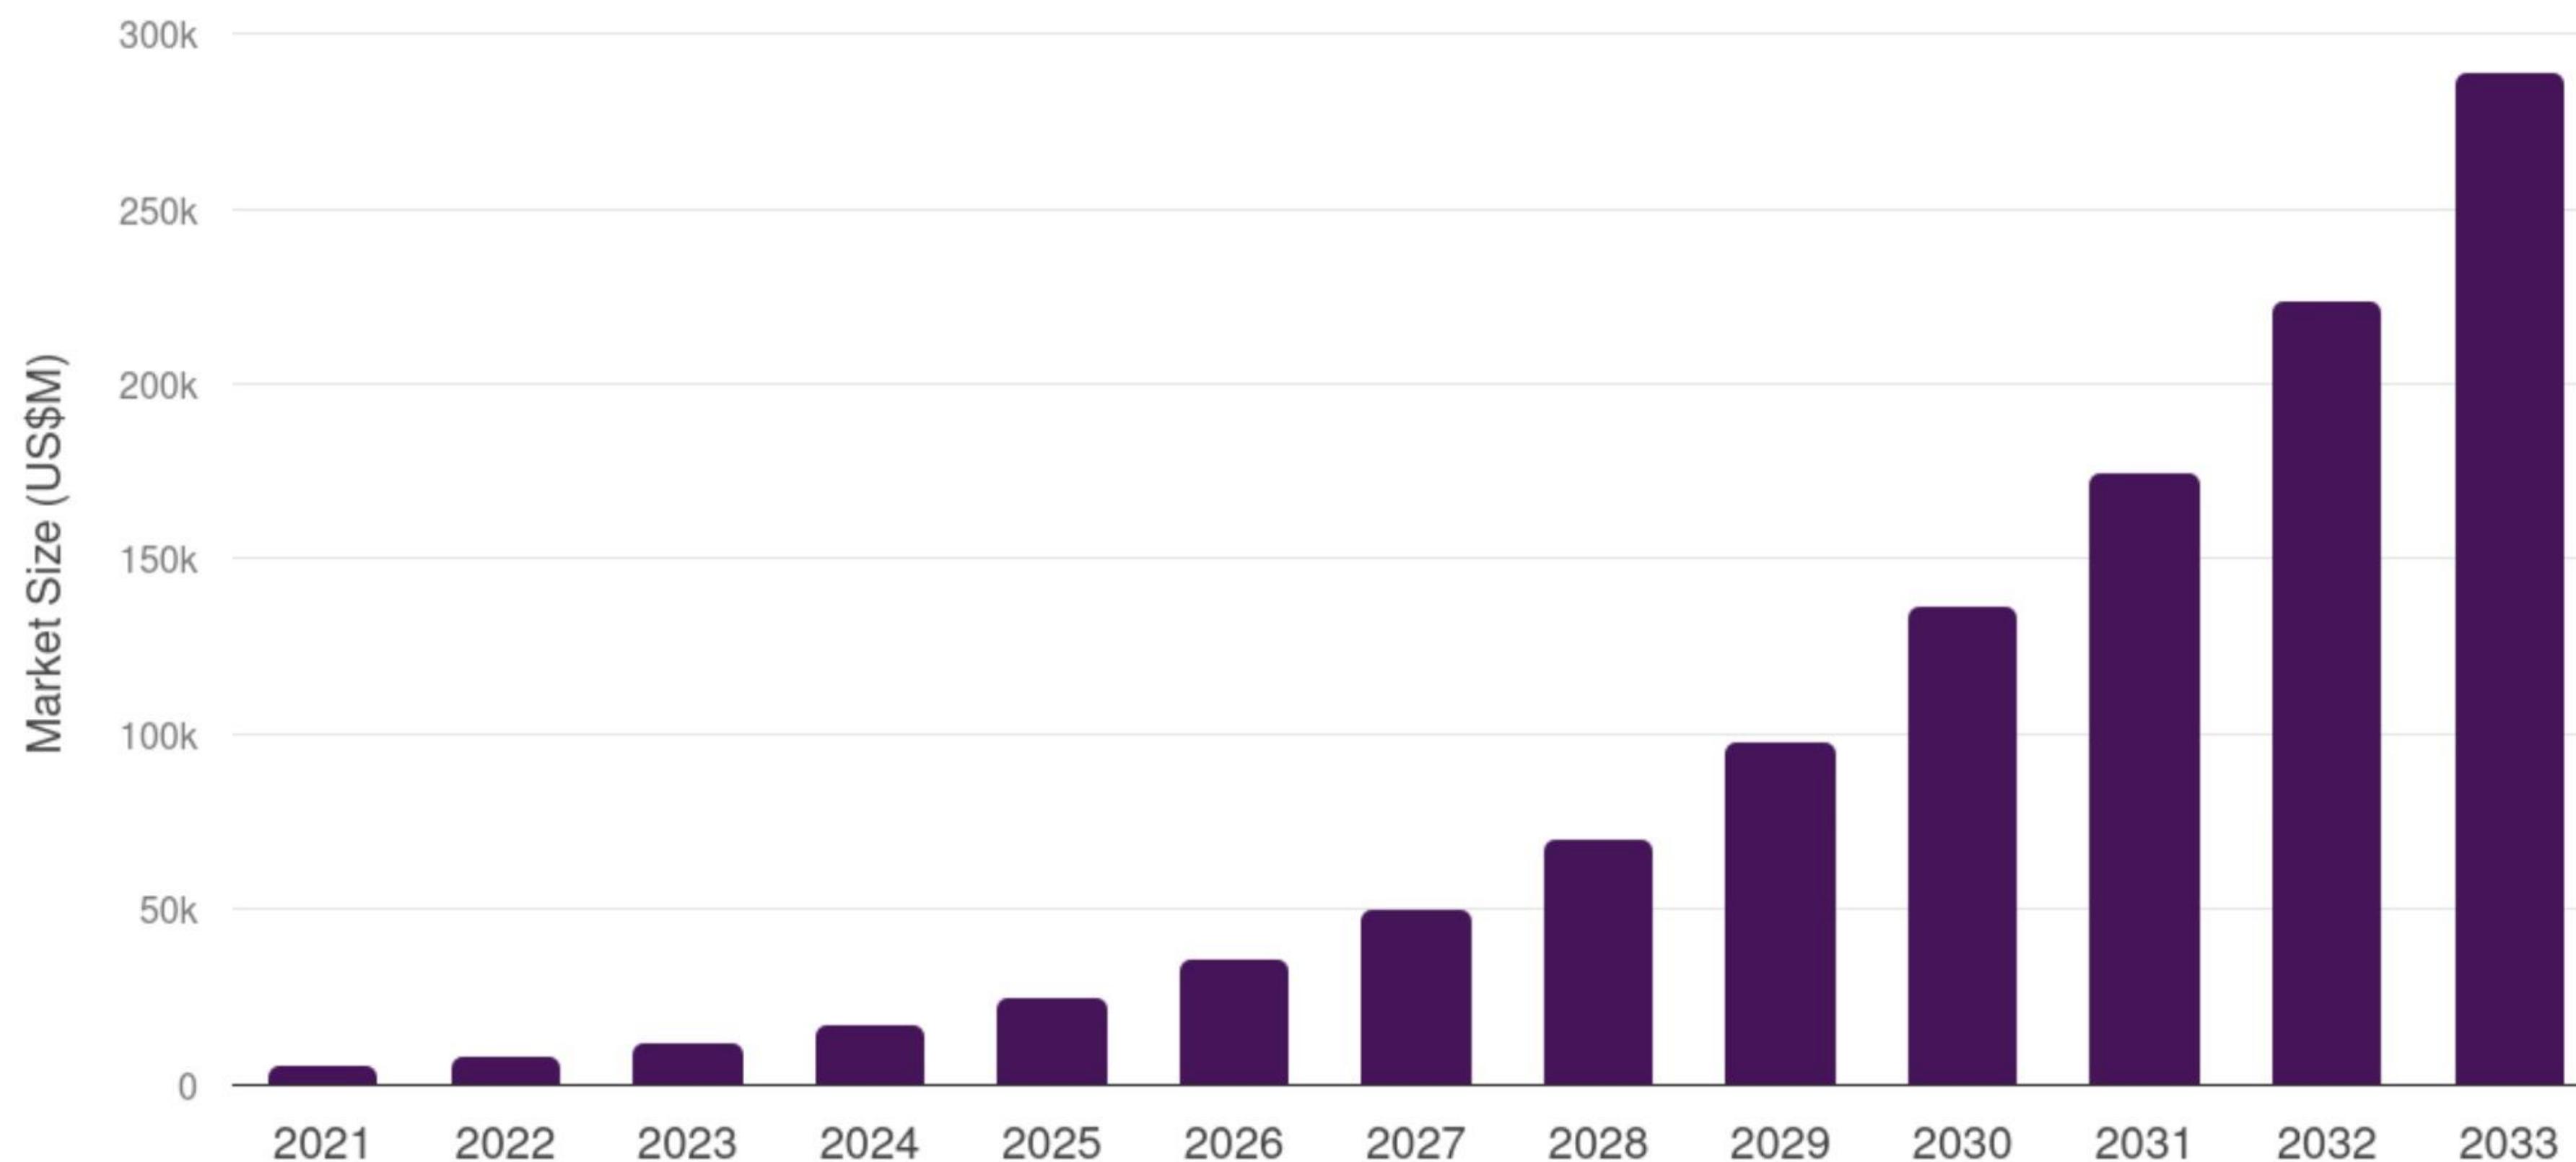
<https://www.grandviewresearch.com/horizon/outlook/artificial-intelligence-market/latin-america>

AI Demand By Region

Middle East & Africa

Grandview Research 2025 Forecast

MEA artificial intelligence market, 2021-2033

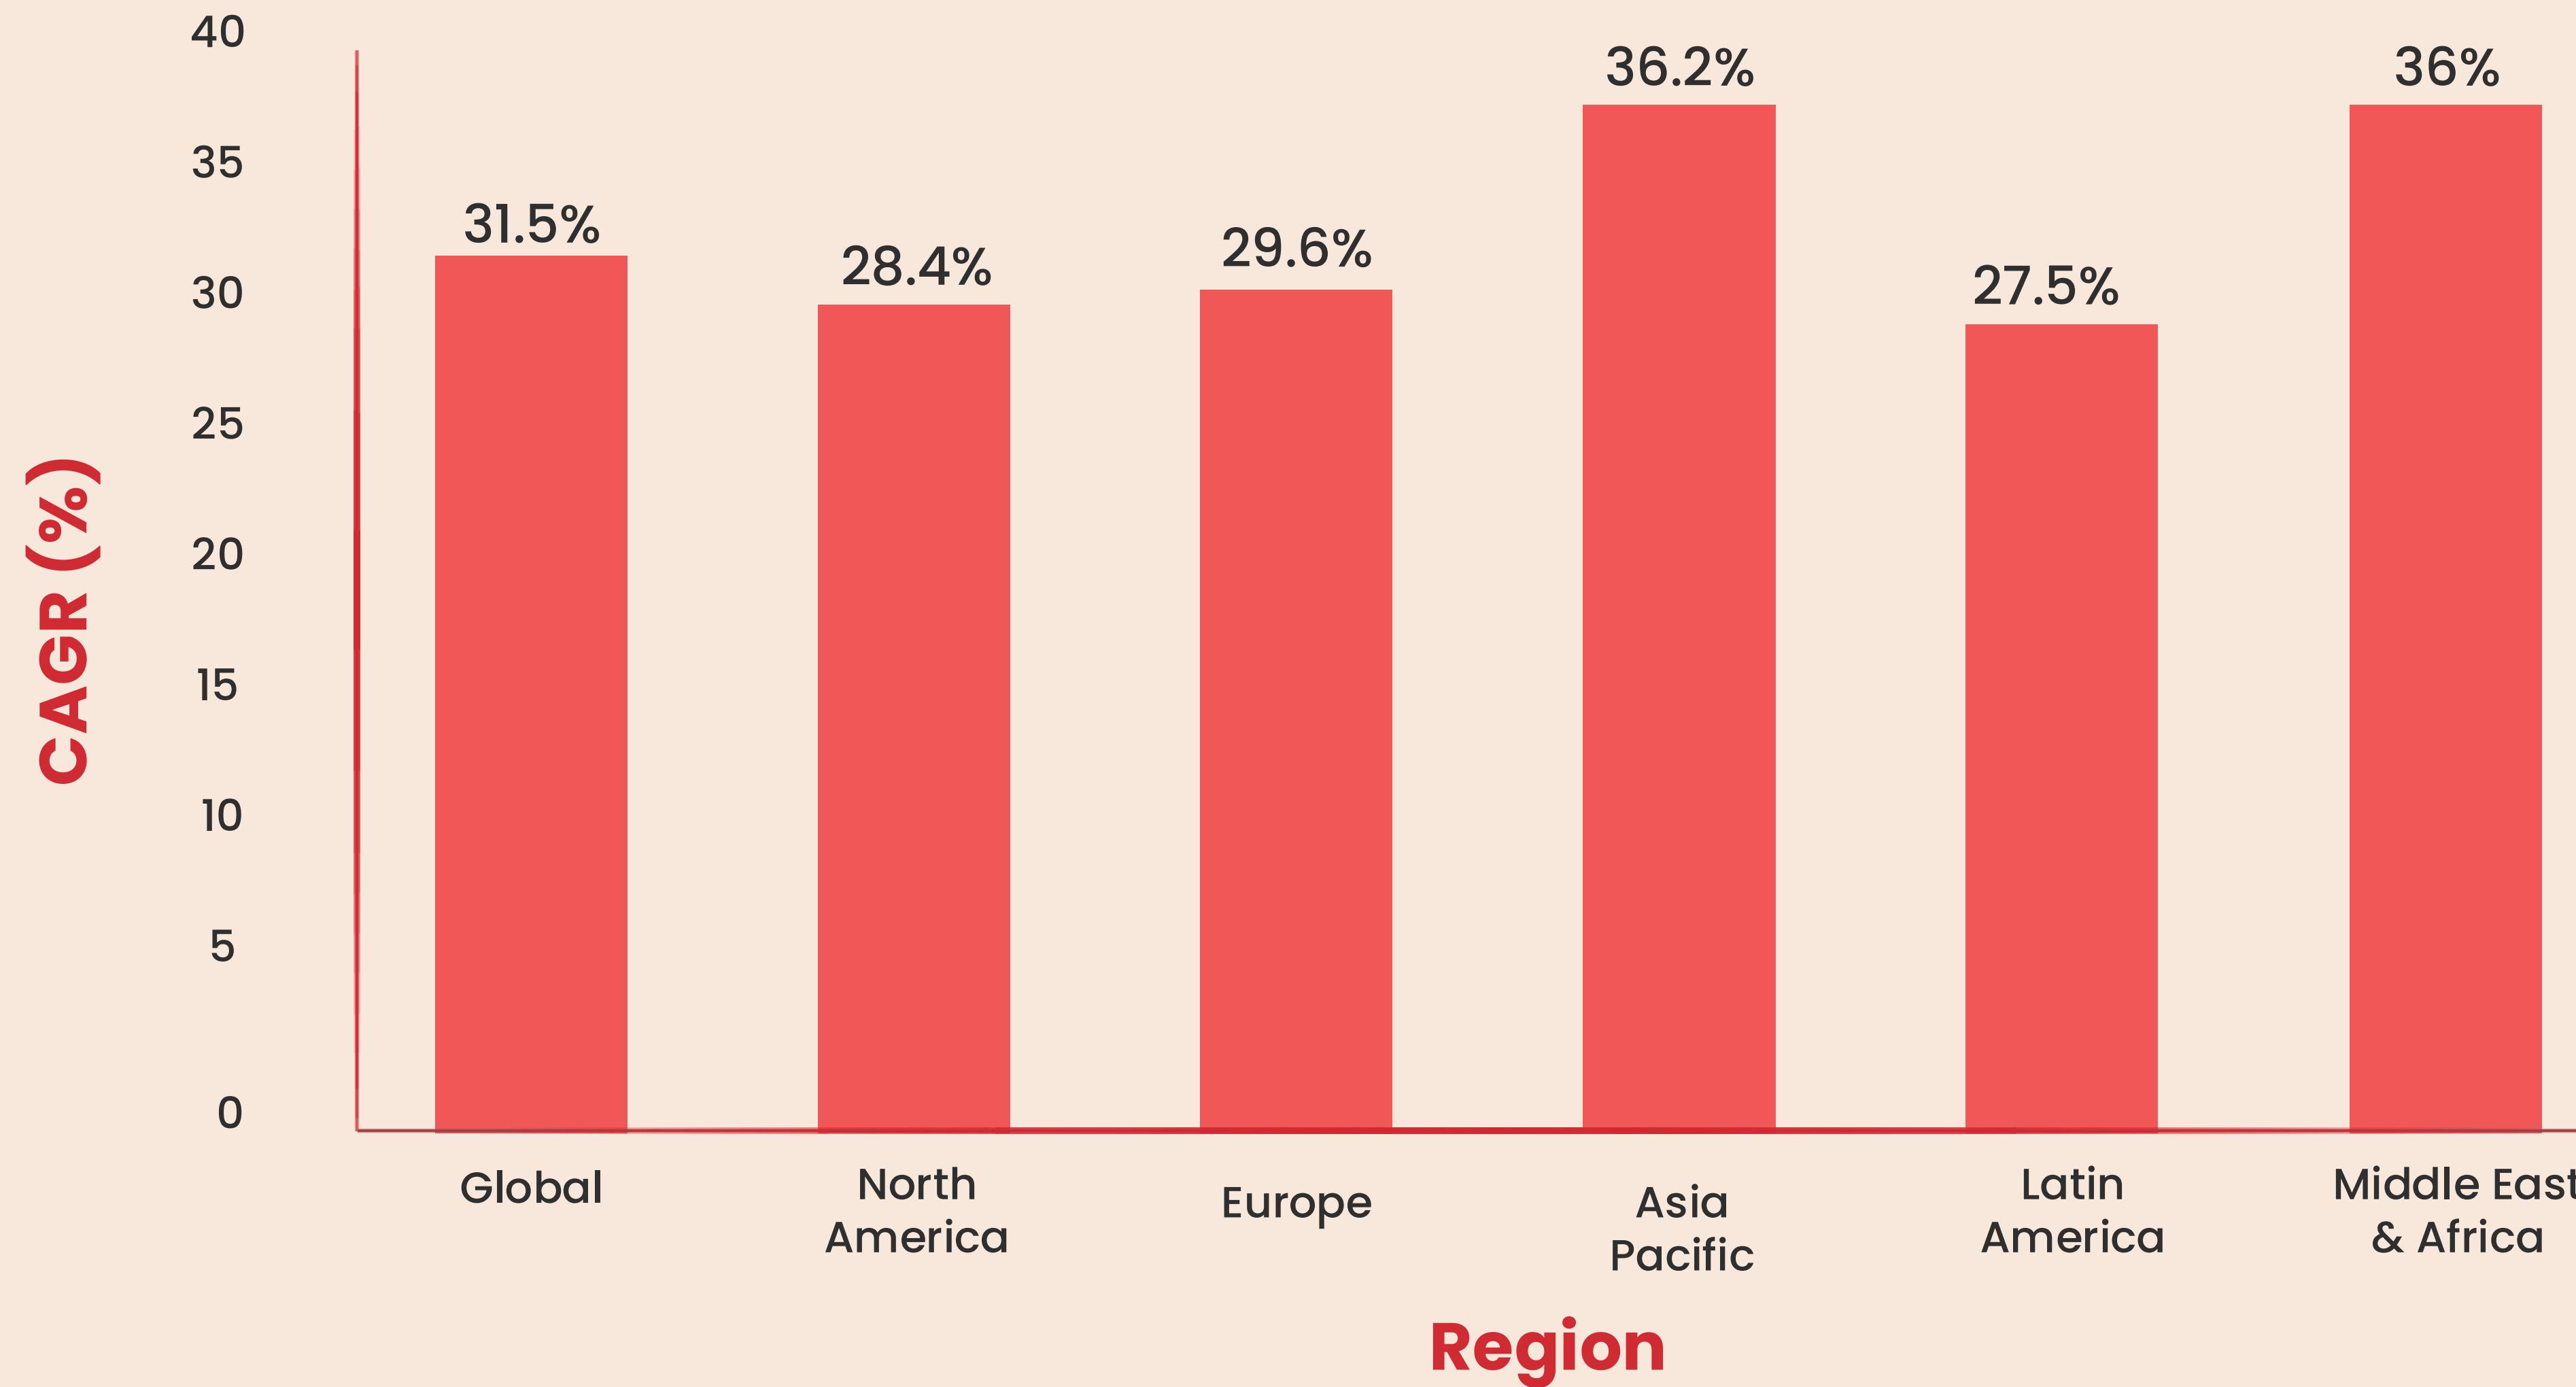


<https://www.grandviewresearch.com/horizon/outlook/artificial-intelligence-market/mea>

AI Demand CAGR By Region

The Emerging Multilingual Global Opportunity

Grandview Research 2025 Report Forecast



Natural Language Processing

Market Projections By Region

Global Market Growth

- 2024: \$27–60B → 2025: \$35–86B → 2030–2034: \$181–791B
- CAGR: 23–39%

Regional Breakdown (2024 → 2025 → 2030–2033)

Region	2024	2025	2030–2033	Share	CAGR
North America	\$8–18B	\$15–22B	\$54–132B	30–48%	25–27%
Asia Pacific	\$9–19B	\$11–27B	\$57–251B	32%	26–32%
Europe	\$6–12B	\$8–20B	\$36–158B	20–22%	23–26%
Latin America	\$1.4–3B	\$1.8–4.3B	\$9–20B	5–7%	22–25%
Middle East & Africa	\$1.2–4B	\$1.4–3.4B	\$7–32B	4–6%	23–27%

Key Highlights

- North America: **Largest market** – U.S. alone: \$6.4B (2024) → \$170B (2034) at 38.7% CAGR
- Asia Pacific: **Fastest-growing** region driven by China and India local language initiatives
- Europe: **2nd largest market** with focus on multilingual capabilities and EU AI Act compliance
- Latin America: **Emerging market** led by Mexico and Brazil
- Middle East & Africa: **Developing market** with growing sovereign AI focus
- Non-English markets represent 50–70% of global opportunity

Sources: Statista, Grand View Research, Mordor Intelligence, Precedence Research (2025)



Global Training Data Infrastructure

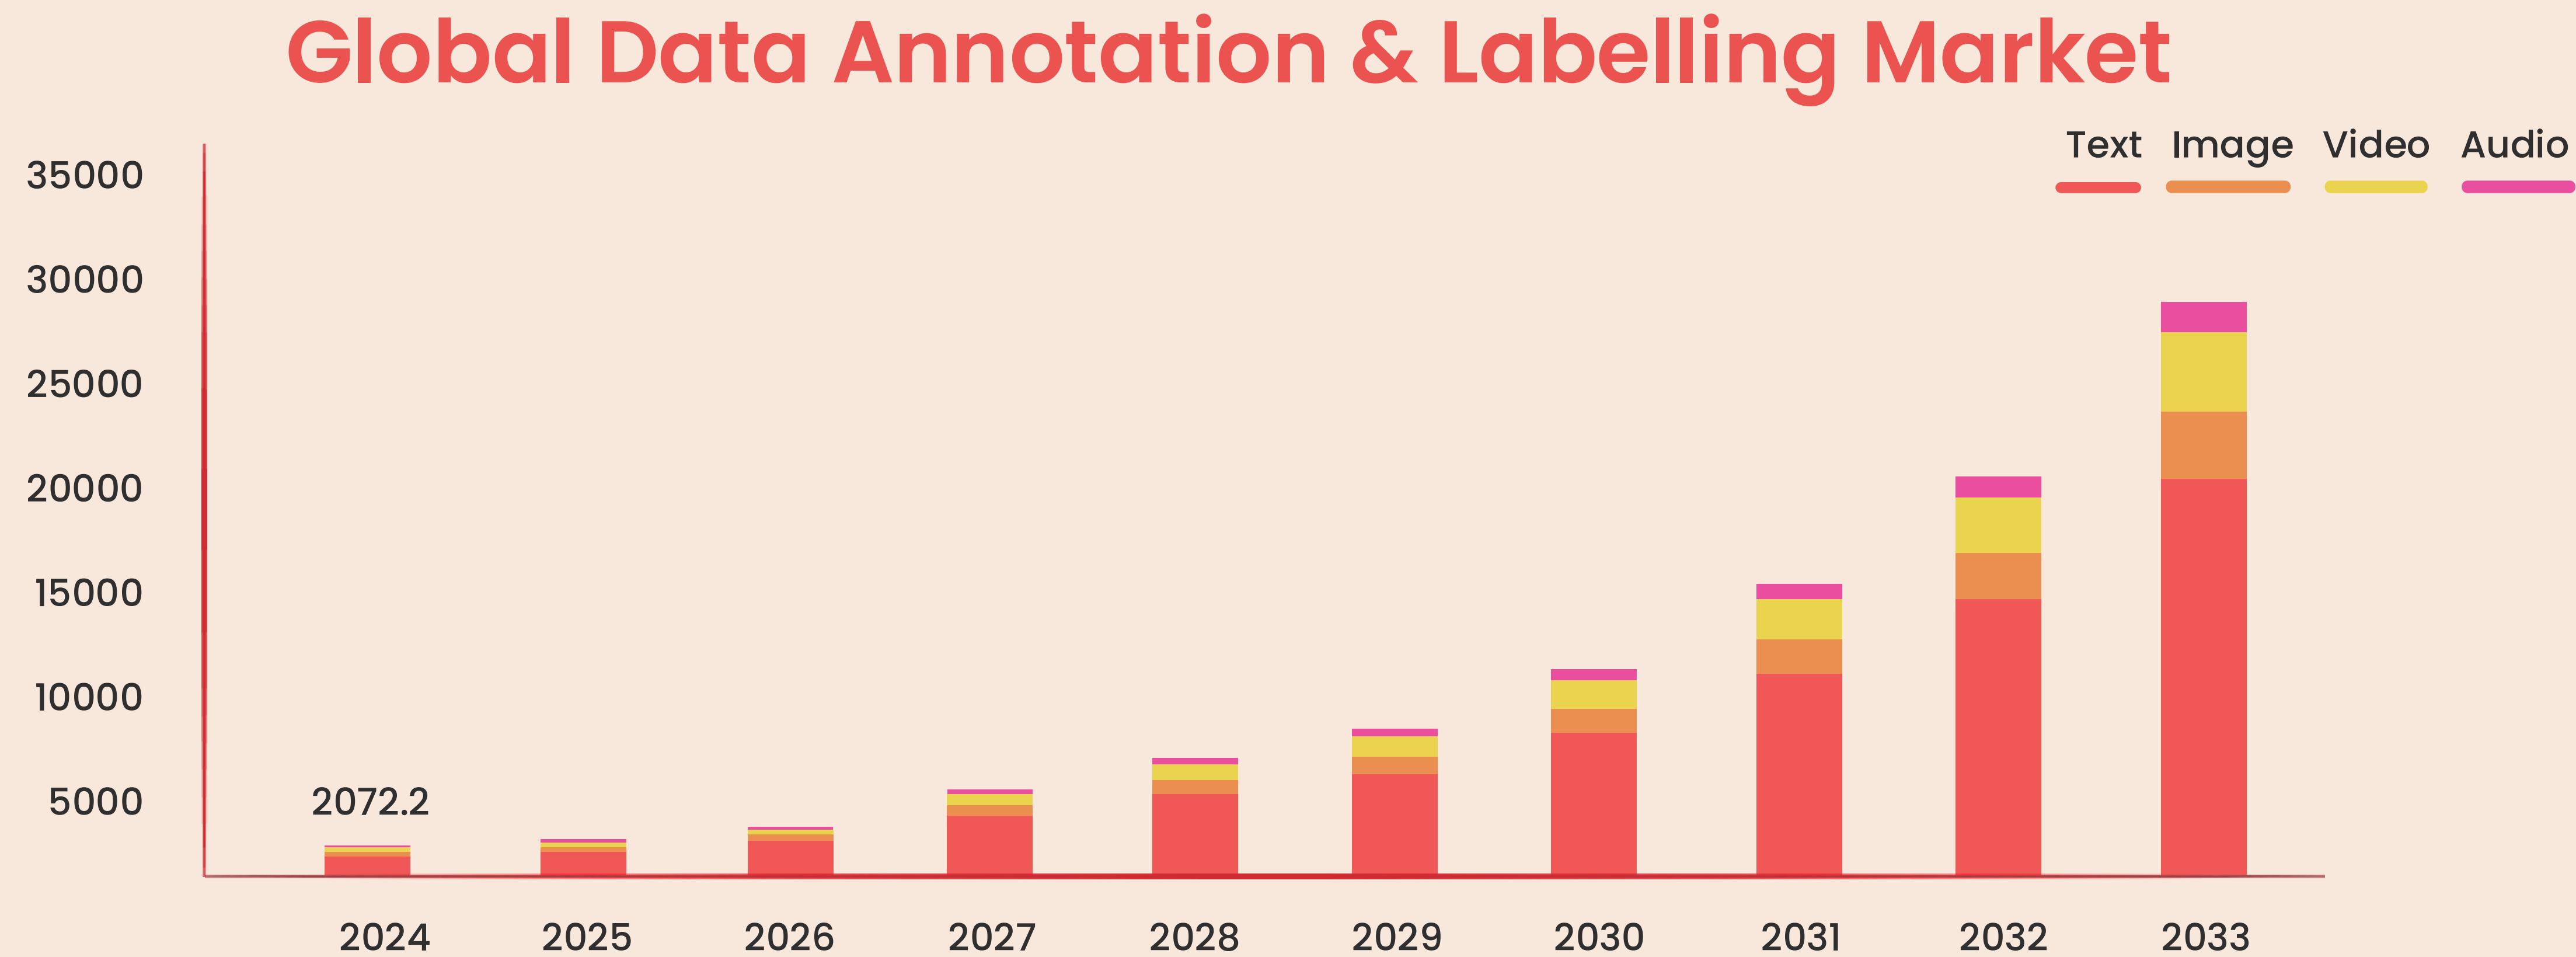
Market Forecast (2024 – 2035)

Year	Base Case (July 2024) Annotation & Labeling CAGR 34.4%	Aggressive Case (Oct 2024) (Includes: Un-Supervised, Semi-Supervised & Supervised Learning) CAGR 24%
2024	\$2.07B	\$15.50B
2025	\$2.79B	\$22.69B
2028	\$6.76B	\$34.33B
2030	\$12.21B	\$51.94B
2033	\$29.65B	~\$100.0B

Market Overview: Base Case

The Global Data Annotation And Labelling Market Size Is Expected To Reach A Value Of USD 2,072.2 Million In 2024, And It Is Further Anticipated To Reach A Market Value Of USD 29,65 Million By 2033 At A CAGR Of 34.4%.

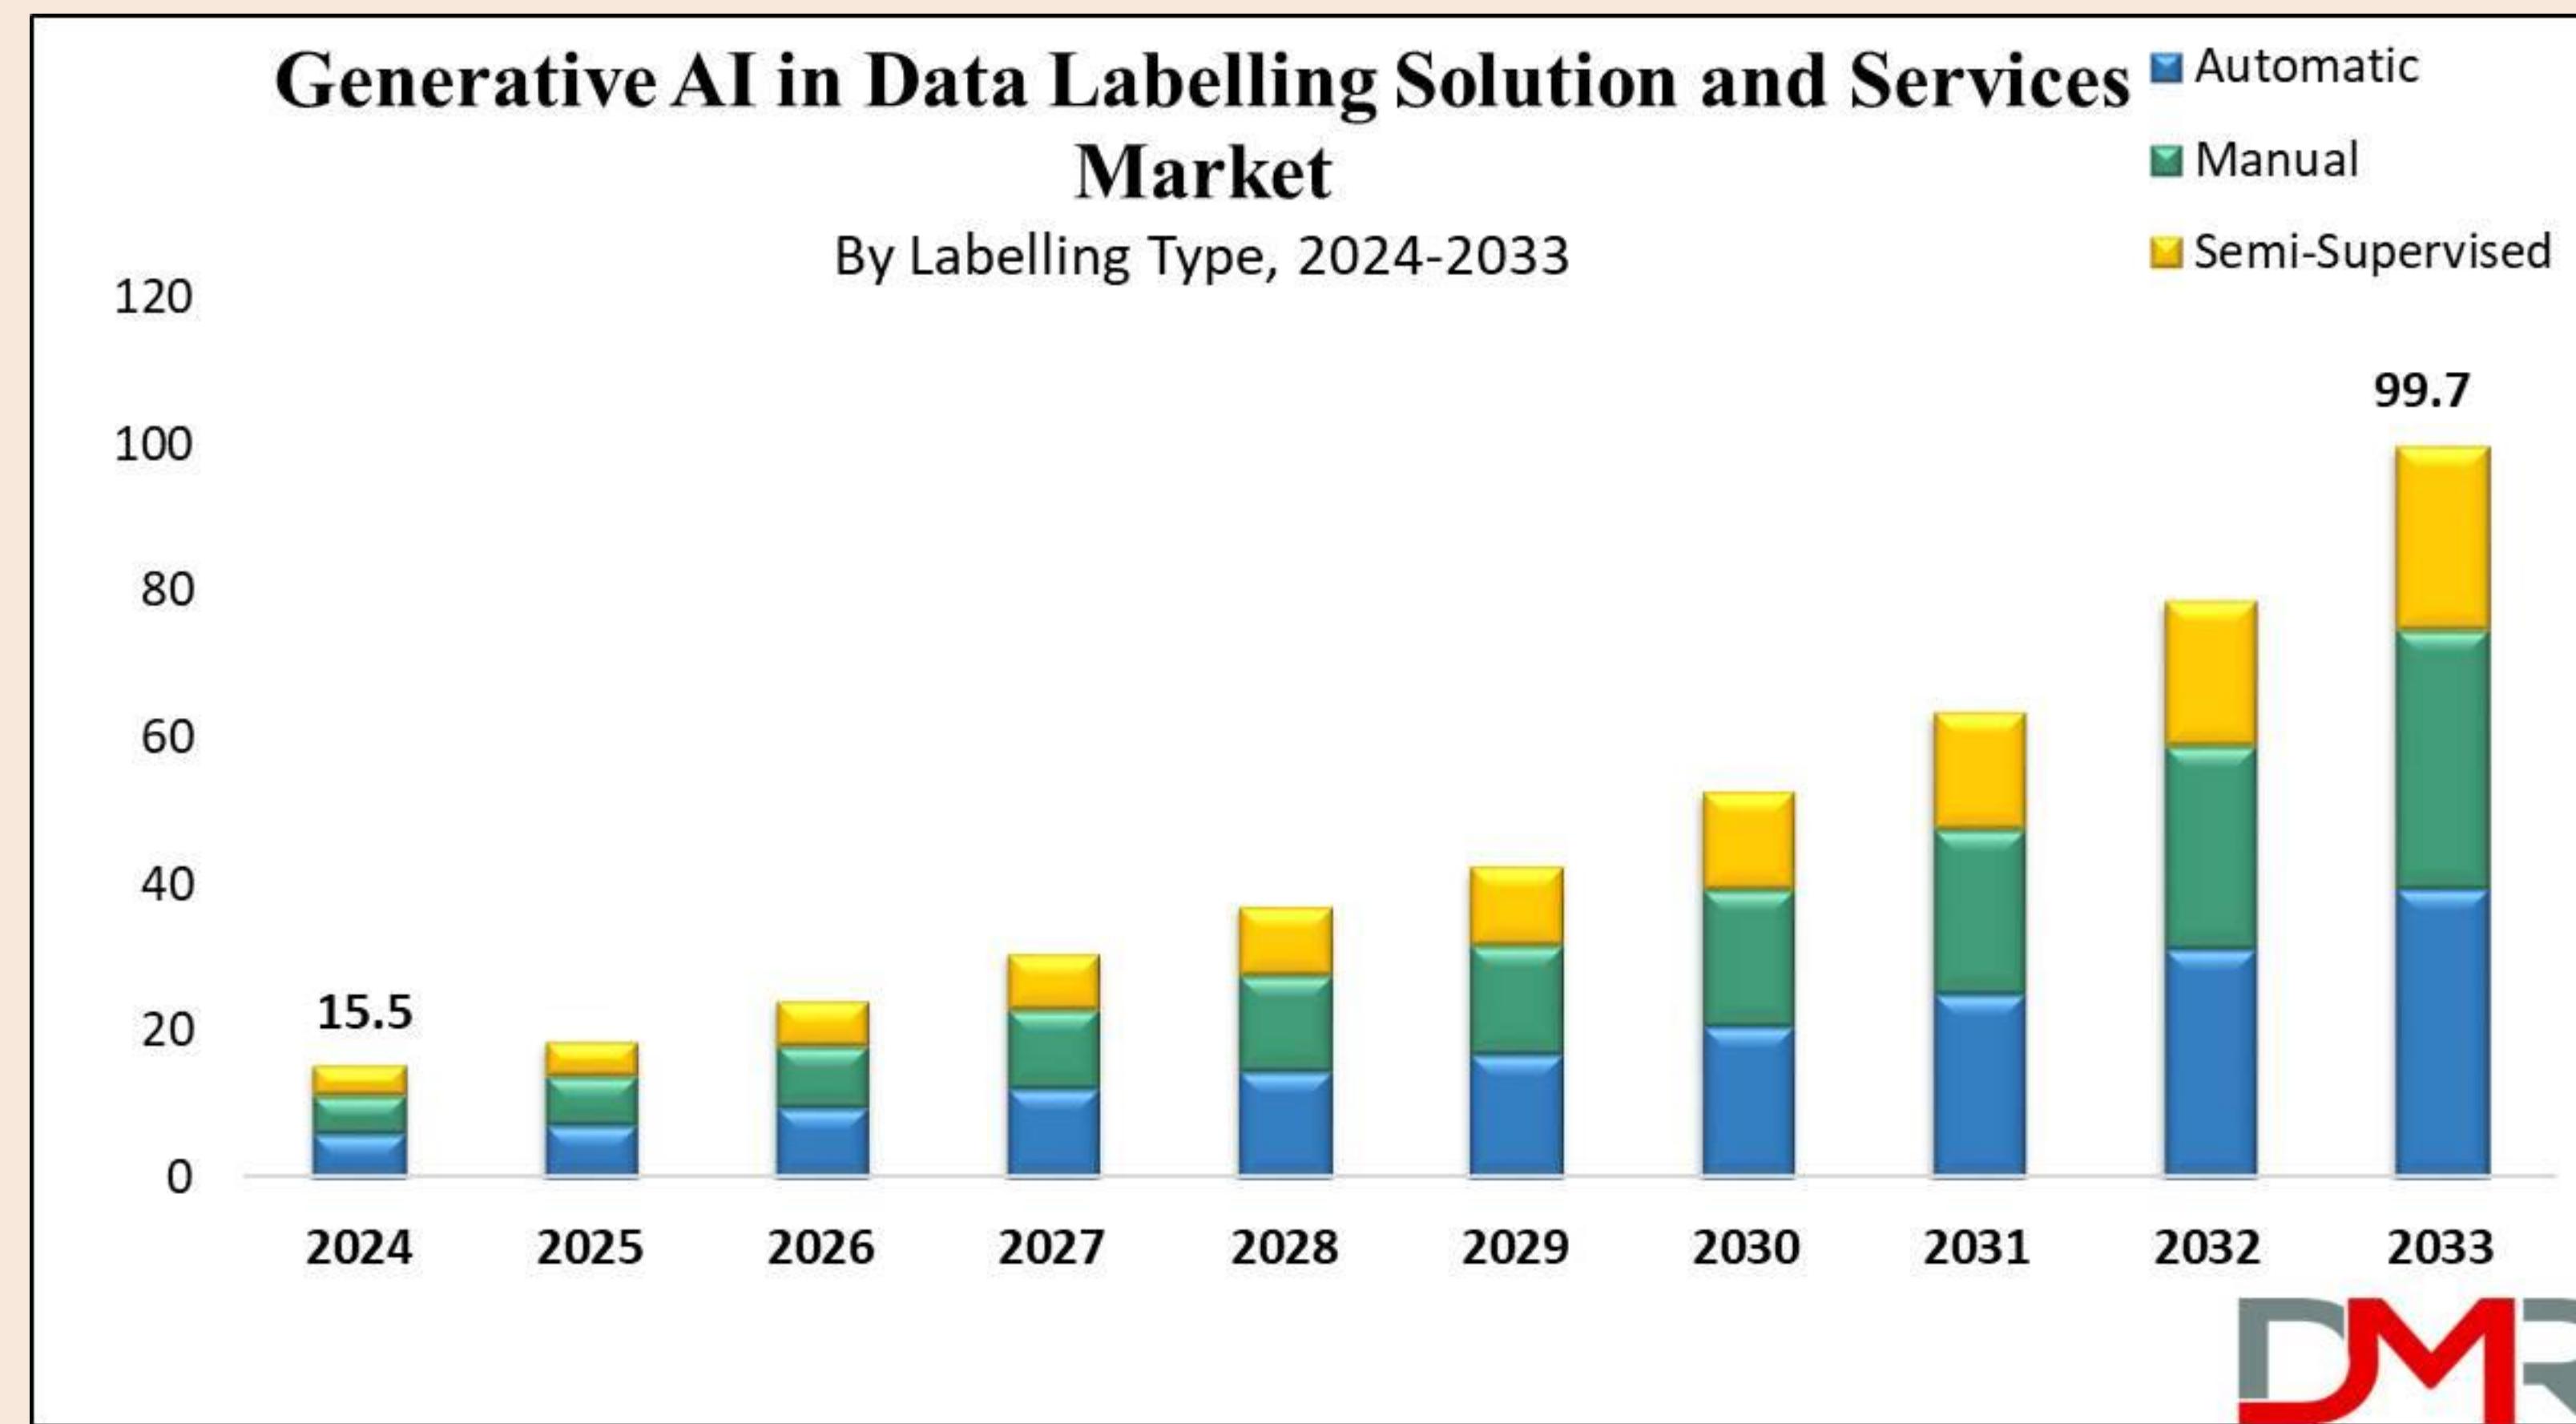
- Source Dimension Market Research, July 2024



Market Overview: Aggressive Case Global Data Labelling Solution & Services

The Emerging \$100B Data Labeling Ecosystem – From Supervised to Semi- and Unsupervised Learning.

- Source Dimension Market Research, October 2024



AI Training & Data Market

EXAMPLE COMPANIES

Market Overview: The AI training and data market spans data labeling, model training infrastructure, and AI development platforms, encompassing supervised, semi-supervised, and unsupervised learning **almost exclusively in English.**

Example Companies:

Scale AI — Market Leader (Founded 2016).

- Valuation: \$29 B (June 2025, following Meta's \$14.3 B investment for a 49% stake)
- Government & Enterprise data labeling and evaluation services for training large AI models (supervised and semi-supervised)

Labelbox (Founded 2018)

- Valuation: \$1 B (2022, SoftBank Vision Fund 2).
- Business Model: SaaS platform for managing data annotation workflows (supervised and semi-supervised)

Hugging Face — Developer Platform (Founded 2016)

- Valuation: \$4.5 B (Aug 2023)
- Business Model: Open ecosystem where 1M+ developers share and fine-tune models, datasets, and AI tools (semi- and self-supervised learning)

Together AI — Infrastructure Layer (Founded 2022)

- Valuation: \$3.3 B (Feb 2025)
- Business Model: Distributed infrastructure for training and running AI models at scale (supporting generative and self-supervised workflows)

Why It Matters: DEMOS builds the next evolution of this ecosystem — a language-based, distributed, autonomous AI infrastructure that unites data labeling, training, and edge-based retraining into one sovereign system.

US Government AI SPENDING

Agency-By-Agency Breakdown (FY 2025 Estimates)

Agency/department	Conservative	Base Case	Aggressive	Primary Use Cases
Department Of Defence	\$2.4B	\$3.7b	\$5.5b	Weapons, Autonomy, Isr, Logistics
Intelligence Community	\$150m	\$200m	\$300m	Sigint, Imint, Translation
Dhs	\$80m	\$120m	\$180m	Border AI, Cybersecurity
State Department	\$30m	\$50m	\$80m	Diplomatic Intelligence, Translation
VA	\$40m	\$60m	\$90m	Healthcare, Benefits Processing
NASA	\$25m	\$40m	\$60m	Autonomous Systems, SatelliteOps
Tresury / IRS	\$20m	\$35m	\$50m	Fraud Detection, Tax Compliance
DOE	\$15m	\$25m	\$40m	Grid Management, Nuclear Safety
Hhs/nih	\$20m	\$30m	\$45m	Drug Discovery, Public Health
Dot /Faa	\$10m	\$15m	\$25m	Aviation Safety, Traffic Management
Other Civilian	\$10m	\$15m	\$25m	Agriculture, Commerce, Education
Total	\$2.8B	\$4.2B	\$6.1B	

Department Of Defense AI Spending (Direct + Embedded)

Fiscal Year	Conservative	Base Case	Aggressive	% Of Total DoD Acquisition
Fy2023 (Actual)	\$0.64B	\$0.64B	\$0.64B	~0.2% Of \$315B
Fy2024 (Actual)	\$0.90B	\$0.90B	\$0.90B	~0.3% Of \$315B
Fy2025	\$2.4B	\$3.7B	\$5.5B	~0.8-1.8% Of \$311B
Fy2026	\$4.1B	\$7.4B	\$12.2B	~1.3-3.9% Of \$315B
Fy2027	\$7.2B	\$13.7B	\$23.3B	~2.2-7.2% Of \$324B
Fy2028	\$12.9B	\$25.1B	\$42.1B	4.1-13.4% Of \$313B
Fy2029	\$18.2B	\$35.8B	\$59.5B	5.8-19.0% Of \$314B
Fy2030	\$25.0B	\$49.5B	\$82.0B	7.9-26.1% Of \$315B
Fy2031	\$31.8B	\$64.3B	\$107B	10.1-34.2% Of \$315B
Fy2032	\$39.7B	\$81.3B	\$138.5B	12.6-44.0% Of \$315B
Fy2033	\$49.5B	\$99.8B	\$175.0B	15.7-55.5% Of \$315B

DoD acquisition budgets from CBO: \$311B (FY2025) → \$324B (FY2027) → \$313B (FY2029); assumed -\$315B steady-state thereafter Congressional Budget Office

FY2023 baseline: Brookings reports ~95% of federal AI spending was DoD

Penetration rate driven by 685+ AI projects embedded in major weapons systems

Major DoD AI Programs Included:

Direct contracts: Replicator (\$500M/year), Maven Smart System (\$795M boost in 2025), CDAO programs

Embedded in platforms: F-35 sensor fusion, autonomous drones, ship AI systems, satellite image analysis, predictive maintenance

Enterprise agreements: Palantir Army EA (up to \$1B/year if fully utilized)

US Government Infrastructure Investments

\$500B – Project STARGATE OpenAI, Oracle, Softbank

- White House-backed AI infrastructure expansion
- New data center sites nationwide
- Critical Challenge: 18%+ US power grid consumption



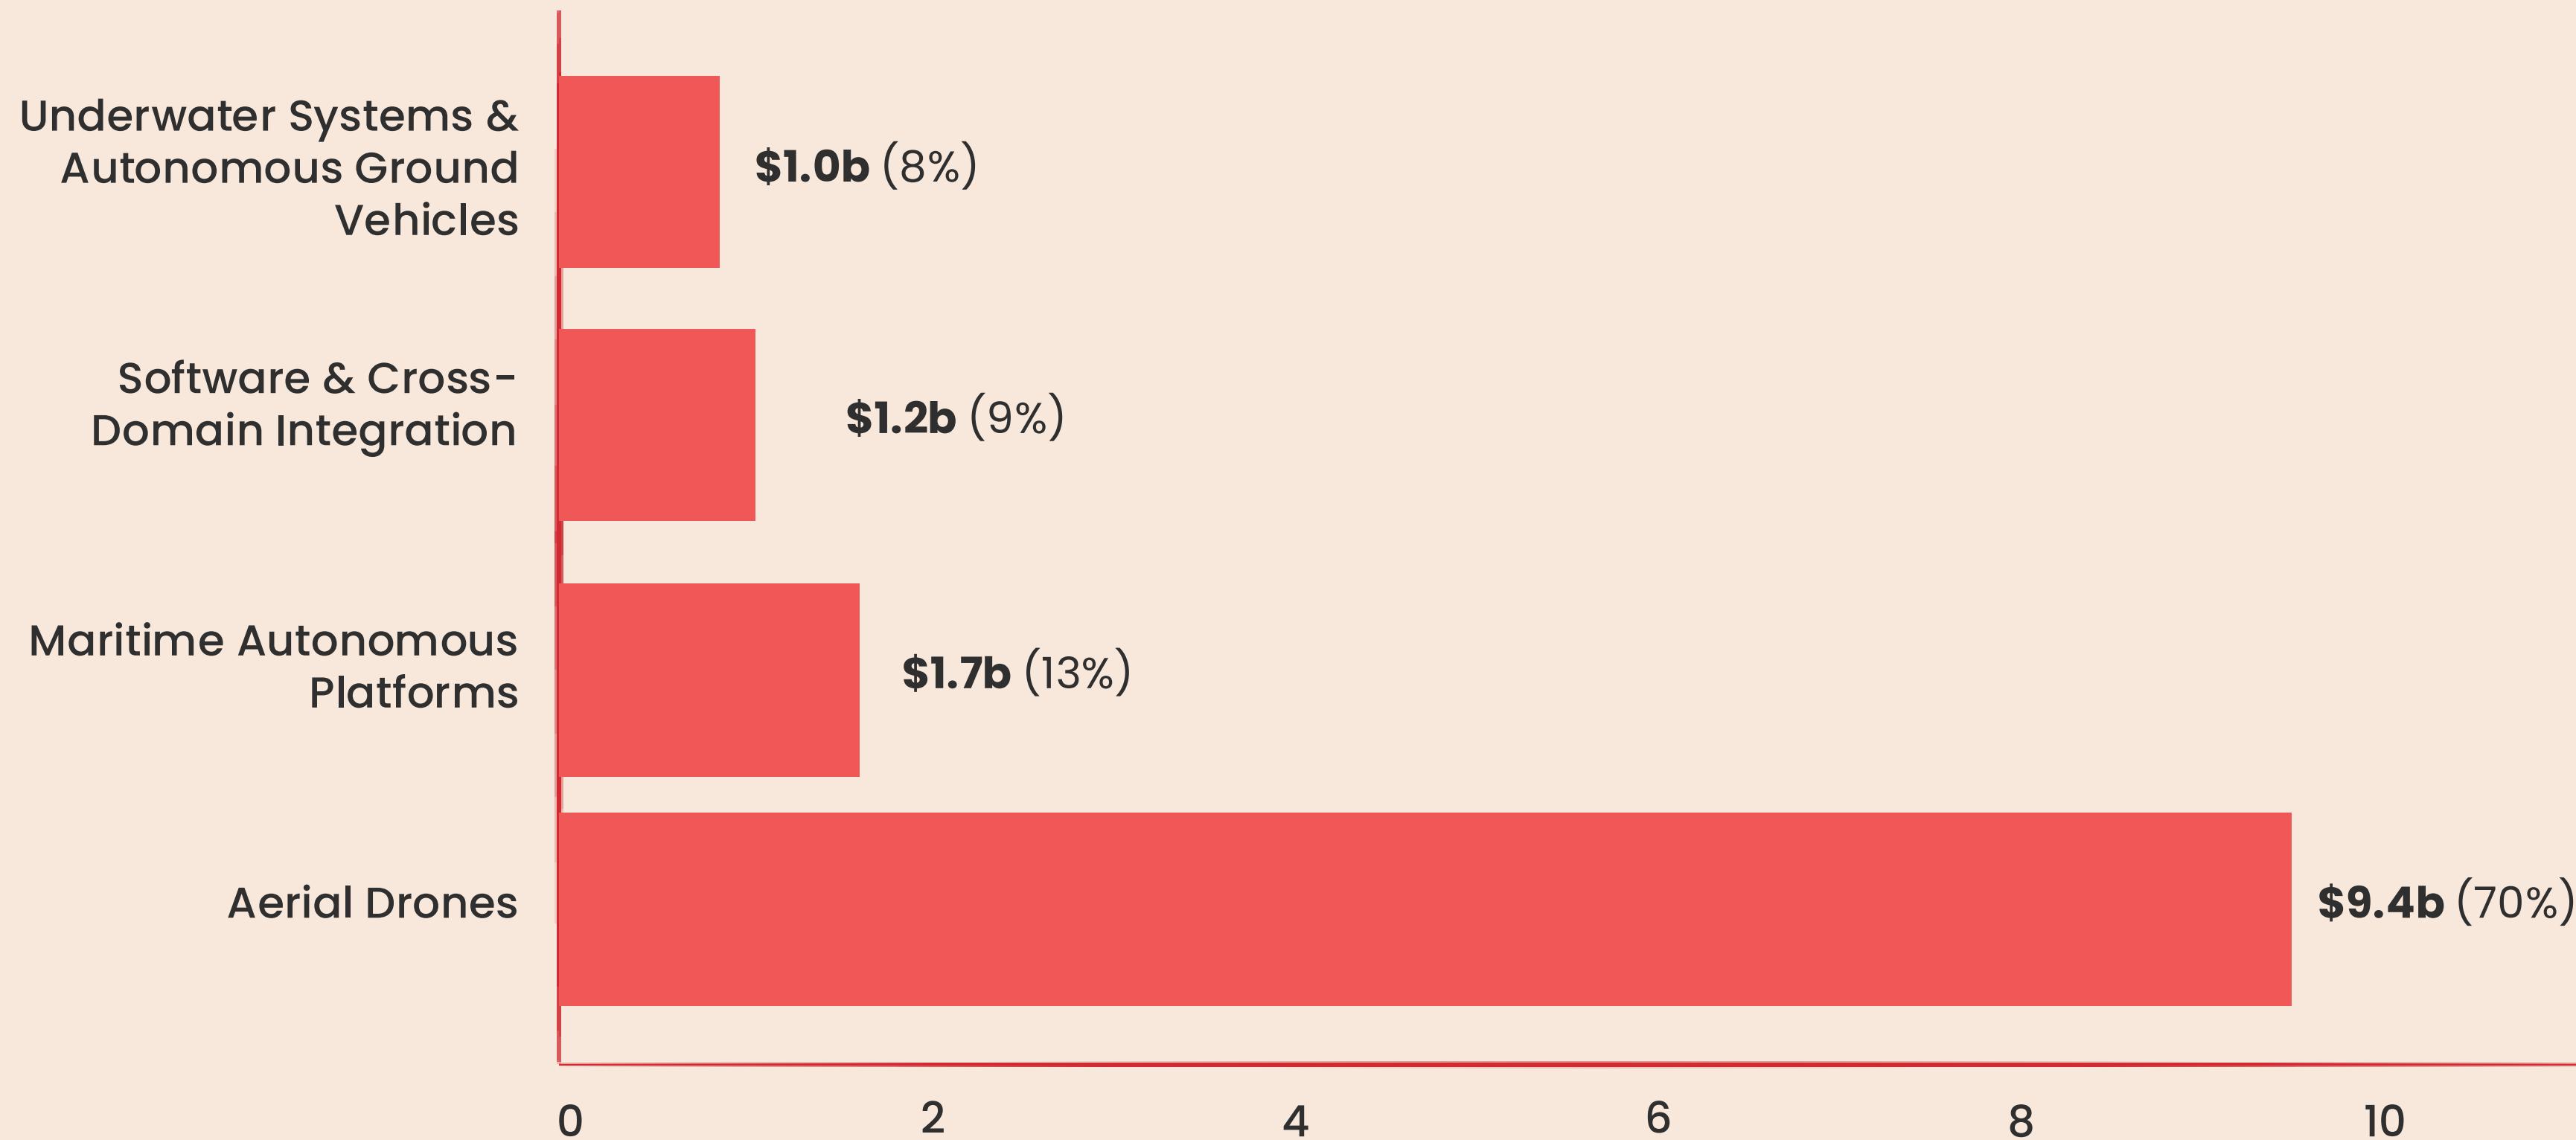
\$833B+ – "One Big Beautiful Bill" Drone & Manufacturing Initiative:

- \$3.5B – Domestic manufacturing & supply chain grants
- \$2B – Defense Innovation Unit (dual-use commercial tech)
- \$1.1B – Small UAS industrial base expansion
- \$1B – Autonomous/kamikaze drone production
- \$500M+ – Counter-UAS systems (kinetic & non-kinetic)



Defense AI Budget (FY 2026)

Total: **\$13.4B** For AI & Autonomy



FY2027-2028 Outlook

Expected to exceed FY2026 levels Priorities: AI, quantum science, semiconductors, advanced communications, advanced manufacturing

GOVT Contract Awards & Market Outlook

Major Contracts Awarded

\$800M – Multi-Vendor AI Platform Deal

- \$200M each: Anthropic, Google, OpenAI, xAI
- Agentic AI workflows for Advana, Maven Smart System, Edge Data Mesh
- Mission-critical defense analytics and decision support

\$100M – Palantir MAVEN Platform

- 5-year IDIQ contract through 2028
- Advanced data analytics and AI capabilities



\$100M – Palantir MAVEN Platform

Market Projection **\$59.6B+ by 2031** – AI in Public Services Conservative estimate; expected significantly higher post-Americas AI Action Plan (July 2025) due to exponential technology evolution

US Government Contractor Example: Palantir

Why Wall Street Values One Company At \$443B

Company Overview: Palantir Technologies (Founded: 2003)

- **Focus:** AI-powered data analytics and decision-making platforms for government and enterprise
- **Core Products:** Palantir Gotham (defense/intelligence), Palantir Foundry (commercial), Palantir AIP (AI Platform)
- **Key Differentiator:** Deep integration with U.S. government agencies; deploys AI for national security, military operations, and critical infrastructure

Peak Market Cap: \$443 Billion (2025) | Trading at 573x earnings

The Palantir Model:

- **July 31, 2025: Army \$10B Enterprise Agreement**
- **10-year ceiling (~\$1B/year) = just 2.3% of market cap**
- **Enterprise model with DoD-wide volume discounts**

Future Insights:

- **Army negotiating "10-15 additional" enterprise agreements → model is spreading (Breaking Defense, Aug 2025)**
- **U.S. procurement law mandates competition—no single vendor can dominate**
- **Federal AI spending → \$100B+ annually by 2033**

Why the \$443B Valuation: Palantir demonstrates that becoming the **preferred AI platform for government agencies** creates enterprise value far exceeding individual contract sizes. The U.S. government views Palantir as critical AI infrastructure—not just a vendor

Bottom Line: If enterprise agreements become standard and **2-5 major AI vendors** share the government market, **each** could command **multi-billion dollar valuations**

US Government Market Validation: Auterion

Edge AI Defense Software: Same Model, Bigger Market

Company Overview: Auterion (Founded: 2017)

- Focus: Edge AI software for military drones (not hardware manufacturer)
- Core Products: Skynode AI kits - 3M-10M parameter computer vision models for autonomous targeting
- Key Technology: Edge deployment, INT-8 quantization, no cloud dependency

Valuation: \$600M+ (Sept 2025) | \$130M Series B

The Pentagon Contract:

- July 2025: \$50M for 33,000 AI drone kits to Ukraine (~\$1,515/kitt)
- Largest Western drone AI deal by unit count
- Partners with Ukrainian/Polish manufacturers (software-only model)
- Expanding to Taiwan, NATO via Rheinmetall partnership

Why the \$600M Valuation:

- Auterion proves governments will pay premium valuations for Edge AI software with sovereignty characteristics (no cloud, battle-tested, allied manufacturing). Single device category (drones only) commanding \$600M

Technical Validation: Auterion demonstrates that small parameter models (3M-10M) with Edge optimization (quantization, pruning, depthwise separable convolutions) achieve production-grade performance for mission-critical applications. Government willingness to deploy at scale validates edge inference viability over cloud architectures for sovereignty-sensitive use case

Market Implications for DEMOS:

- Auterion proves the Edge AI software model at \$600M valuation for single modality (computer vision) on single device category (drones).
- DEMOS applies identical architectural constraints (small parameters, edge deployment, sovereignty-first) to multimodal models (vision + language) across heterogeneous edge devices (phones, IoT, drones, vehicles) with cross-lingual capabilities (127 languages)
- Thesis: If specialized edge AI = \$600M, generalized edge AI = \$1B+ opportunity

The Products

Product 1: Training AI Ecosystem Foundation For Language Based AI

Overview: The **foundation** of DEMOS's architecture is a **distributed data labeling ecosystem** — combining **human expertise, local language knowledge, and generative automation**. It's not just a service; it's the **infrastructure layer** that powers sovereign, small-parameter AI models for every language and region

Core Features:

Hybrid Human + AI Labeling.

- Combines native-language human annotators with generative AI tools for rapid, context-rich labeling

Language-Native Teams:

- Each region's data is labeled by local speakers to preserve nuance, idioms, and cultural semantics.

Multi-Learning Architecture

- Supports supervised, semi-supervised, and unsupervised pipelines — from human annotation to model-generated labeling

Edge Integration

- Designed to connect directly with DEMOS's Edge AI nodes for on-device retraining and feedback loops

Sovereign Data Ownership

- All datasets remain within national or regional boundaries — ensuring compliance with data sovereignty and privacy regulation

Curated Dataset Hosting for Small-Parameter Models & Edge.

- Hosts high-quality, language-specific datasets optimized for **small-parameter and edge-deployed models**, enabling localized training and retraining pipelines focused on linguistic and cultural precision

Why It Matters: This ecosystem forms the **base layer of DEMOS AI LAB's architecture** — the data substrate upon which every language-specific, self-reinforcing AI model is built. It replaces outsourced annotation with a **sovereign, distributed network** that learns from and belongs to its people

Product 1: Training AI Ecosystem Design Example

Foundation For Language Based AI

Overview: DEMOS's distributed data labeling ecosystem design example for pitch deck purposes only.

The screenshot displays the Demos AI Lab Global Annotation Hub interface. At the top, it shows 1,247 active annotators, 23.9M annotations today, and a 98.7% quality score. The main task details are for Task #MED-2847, labeled "Medical Terminology" and "Spanish". The progress bar indicates 65%. The task text is: "El paciente presenta síntomas de hipertensión arterial sistémica con valores de presión arterial de 160/95 mmHg. Se recomienda iniciar tratamiento con inhibidores de la enzima convertidora de angiotensina (IECA) y modificaciones en el estilo de vida." Below the text are five annotation buttons: Entity Tagging, Relationship, Translation, Validation, and Intent. A sidebar on the left lists "TASK CATEGORIES" (Medical Terms 342, Legal Documents 218, Technical Manuals 156, Financial Reports 94, Cultural Context 287) and "DIFFICULTY LEVEL" (Beginner 512, Intermediate 342, Expert 89). On the right, "TODAY'S EARNINGS" are listed as \$127.50, broken down into Medical Tasks (\$85.00), Legal Tasks (\$32.50), and Bonus (Quality) (\$10.00). The "TOP ANNOTATORS - SPANISH" section ranks four users: Maria Rodriguez (1, Madrid, 2,847 tasks, 99.8%), Carlos Mendez (2, Mexico City, 2,654 tasks, 99.5%), Ana Silva (3, Buenos Aires, 2,432 tasks, 99.3%), and You (47, San Juan, 847 tasks, 99.2%). The bottom of the interface shows "YOUR STATS" with completed tasks (47), accuracy rate (99.2%), entities tagged (4), confidence (98%), time spent (2:34), and quality score (A+).

Product 1 Training AI Ecosystem Design Example

Foundation For Language Based AI

Design Platform Concept: Premium Small Data Set Curation & Hosting Focused On Language & Edge AI

The screenshot shows the homepage of the Demos AI Lab Enterprise Licensing Portal. At the top, there's a dark header with the logo 'Demos AI Lab' and 'ENTERPRISE LICENSING PORTAL'. To the right are navigation links for 'Datasets', 'My Licenses', 'API Access', 'Analytics', 'Documentation', and 'Support', along with 'ENTERPRISE' and 'FM' buttons. Below the header, a welcome message 'Welcome back, Fortune Medical' is displayed, followed by statistics: '12 ACTIVE LICENSES', '847M TOTAL RECORDS', and '28 LANGUAGES'. A search bar and dropdown menus for 'All Categories', 'All Languages', and 'License Type' are also present. The main content area features six data set cards arranged in two rows of three. Each card includes an icon, the dataset name, a brief description, language support (e.g., Spanish, Portuguese, Polish), key metrics (e.g., 12.3M annotations, 99.8% accuracy, weekly updates), and a price (\$85,000 per year). A 'Get License' button is included in each card.

Dataset	Language Support	Annotations / Documents	Accuracy	Updates	Price
Medical Terminology Corpus	Spanish, Portuguese, Polish	12.3M Annotations	99.8% Accuracy	Weekly Updates	\$85,000 per year
Financial Reports Dataset	Spanish, Portuguese, Ukrainian	8.7M Documents	98.5% Accuracy	Monthly Updates	\$120,000 per year
Legal Document Archive	Spanish, Polish	5.2M Documents	99.2% Accuracy	Quarterly Updates	\$95,000 per year
Technical Documentation Suite	Portuguese, Polish, Ukrainian	15.8M Pages	98.9% Accuracy	Monthly Updates	\$75,000 per year
Cultural Context Database	Spanish, Portuguese, Polish, Ukrainian	23.4M Entries	97.8% Coverage	Real-time Updates	\$150,000 per year
Government & Public Sector	Spanish, Ukrainian	6.9M Documents	99.5% Accuracy	Monthly Updates	\$110,000 per year

Product 2: Edge AI Ecosystem Platform

Localized Intelligence Trained On Premium & Sovereign Data

Overview: The DEMOS Edge AI Platform transforms labeled data from the DEMOS Data Labeling Ecosystem into compact, continuously learning models that run directly on local devices. These language-specific, small-parameter models deliver real-time intelligence without cloud dependence—enabling sovereign, low-latency AI for every region

Core Technical Features:

Core Learning Paradigm:

- **Supervised Learning** – Trains on human-labeled datasets from local annotation teams
- **Semi-Supervised Learning** – Combines smaller labeled corpora with large volumes of unlabeled text and sensor data
- **Unsupervised / Self-Supervised Learning** – Learns structure and semantics from raw data autonomously
- **Reinforcement Learning** – Continuously improves performance through user feedback and environmental interaction

Dual-Track Model Strategy: DEMOS fine-tunes existing open foundation models (such as Mistral, Falcon, or LLaMA) to accelerate regional deployment while simultaneously developing its own DEMOS-native small-parameter foundation model. Both approaches rely entirely on premium, curated, sovereign, localized datasets from Product 1 and are optimized for edge-first environments rather than cloud infrastructure

Fine-Tuning & Localization Layer:

- **Regional Fine-Tuning Pipeline** – DEMOS adapts and fine-tunes both open-source and DEMOS-native models into language-specific, culturally aligned AI systems using localized datasets from Product 1
- **Domain Specialization** – Enables models to serve local institutions—banks, hospitals, governments—with context-aware precision
- **Continuous Retraining Loops** – Edge devices contribute anonymized feedback that fuels ongoing fine-tuning through the Queen Node

Product 2: Edge AI Ecosystem Platform

Core Technical Features Continued

Core Technical Features Continued:

Architecture & Optimization Layer:

- **Transformer & Mamba Architecture** – Supports both Transformer-based models (for large-scale training and fine-tuning) and Mamba state-space architectures (for efficient, real-time edge inference).
 - Mamba delivers Transformer-level capability with longer context retention, streaming adaptability, and up to 80 % lower compute requirements—ideal for multilingual, always-on device
- **TinyML Integration** – Enables ultra-low-power deployment (<10 mW, <1 MB footprint) across microcontrollers, IoT sensors, drones, and mobile devices. This allows DEMOS models to operate autonomously in disconnected or infrastructure-poor regions, fulfilling the mission of AI for everyone, everywhere – Learns structure and semantics from raw data autonomously
- **Floating-Point Precision (FP8 / FP4)** – Reduces compute and memory by up to 80% while maintaining model accuracy, enabling efficient on-device AI training and inference
- **Reinforcement Learning** – Continuously improves performance through user feedback and environmental interaction

Model Evaluation & Alignment:

- **Performance Benchmarking** – Measures model accuracy, latency, and bias using regional validation datasets
- **Ethical & Cultural Alignment** – Evaluates linguistic fairness, tone sensitivity, and contextual accuracy across dialects and local norms
- **Federated Evaluation Layer** – Performs distributed testing directly on edge nodes, ensuring real-world reliability before model rollouts

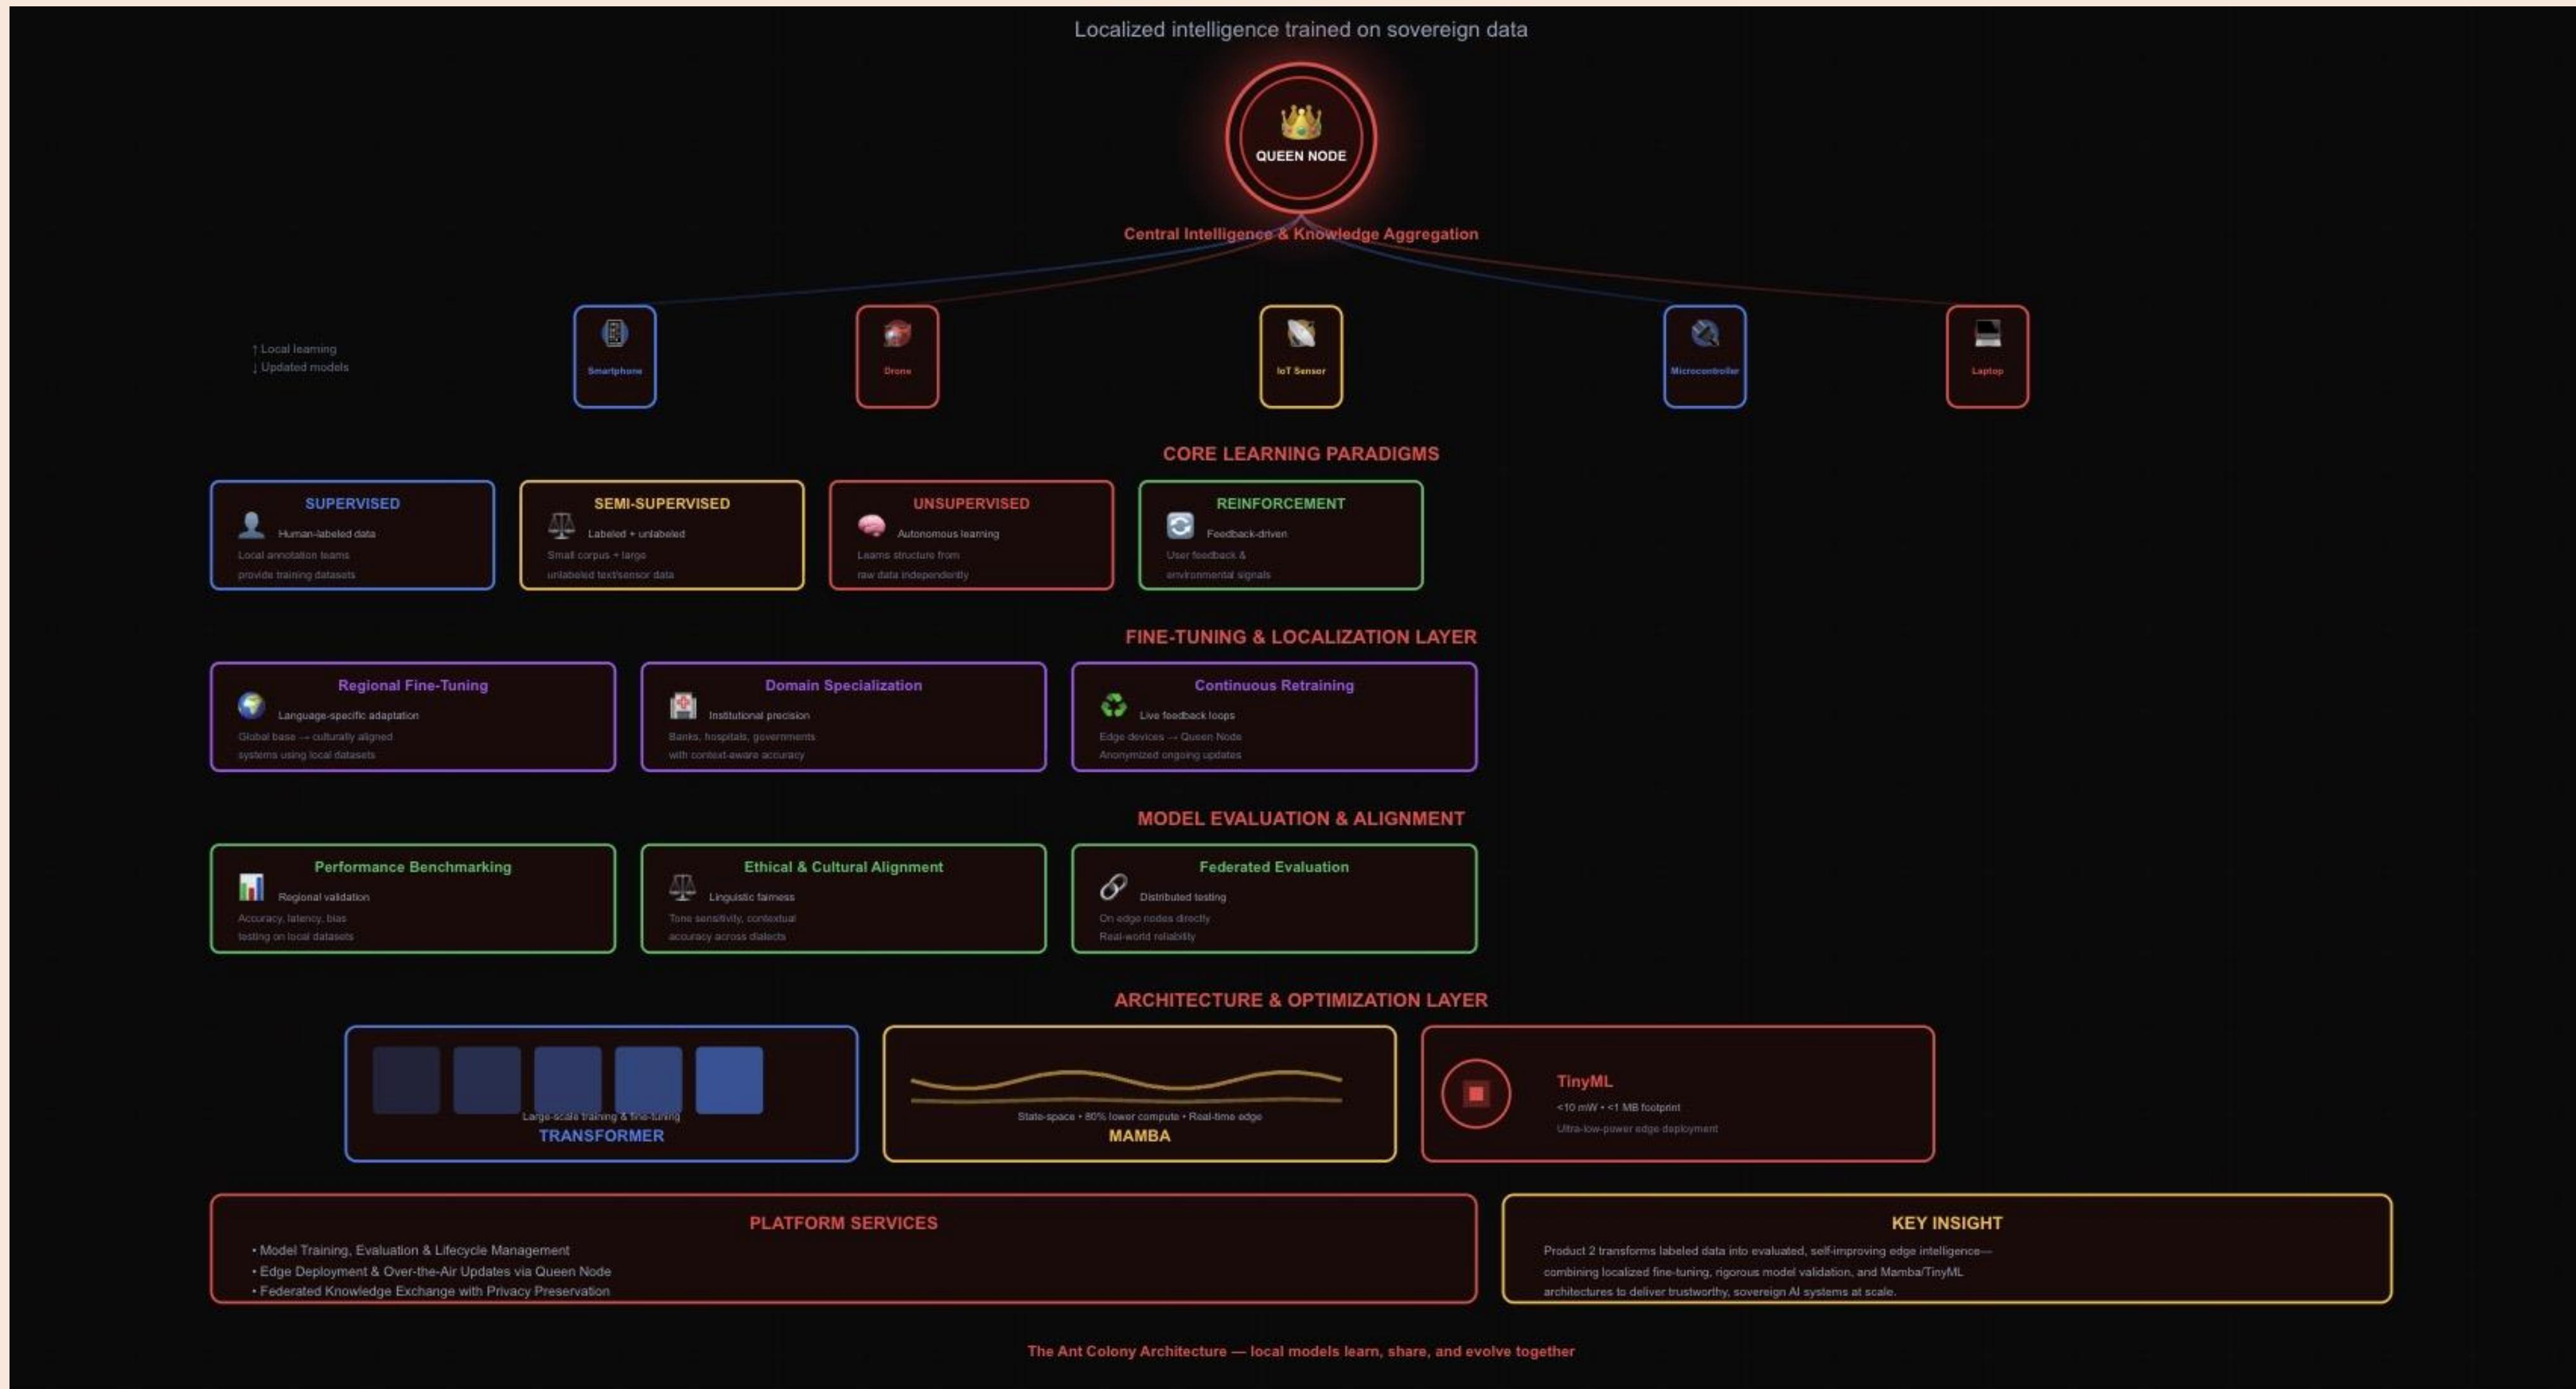
Platform Services:

- Model Training, Fine-Tuning, Evaluation, and Lifecycle Management across all learning paradigms
- Edge Deployment & Over-the-Air Updates via the Queen Node.
- Federated Knowledge Exchange preserving accuracy while maintaining user privacy

Product 2: Edge AI Ecosystem Platform

Platform Concept

Design Platform Concept: Edge AI Ecosystem Platform



Product 2: Edge AI Ecosystem Platform

Platform Design Interface Concept Example

Design Platform Concept: Edge AI Ecosystem Platform

The screenshot displays the Edge AI Ecosystem Platform interface, specifically the Medical Assistant agent section. At the top, it shows "YOUR AGENT Medical Assistant 1,247" with an accuracy of "98.3%" and "LAST SYNC 2m ago". A red circular icon indicates "OFFLINE MODE". The interface is in Spanish.

YOUR AGENT STATUS: Medical Assistant, Spanish Medical Agent v2.3, Specialized: Cardiology, Emergency, Last updated: 2 hours ago.

CONVERSATION: You + 10:32 AM: Tengo dolor de cabeza intenso y náuseas desde esta mañana. ¿Qué podría ser? Medical Assistant + 10:32 AM: Entiendo tus síntomas. El dolor de cabeza intenso con náuseas puede tener varias causas. Algunas posibilidades incluyen: • Migránea - dolor pulsátil, sensibilidad a la luz • Tensión o estrés - dolor constante, rigidez muscular • Deshidratación - especialmente si hace calor.

PERFORMANCE METRICS: RESPONSES TODAY: 47 (↑ 12 from yesterday), AVG RESPONSE TIME: 0.8s (Running locally).

CAPABILITIES: Medical diagnosis assistance, Symptom analysis & triage, Treatment recommendations, Medical terminology translation, Emergency protocol guidance.

DEVICE SYNC: Connected Devices: iPhone, Laptop, Apple Watch. Cloud Sync Status: Synced - All devices updated.

SOVEREIGN DATA: Your conversations stay on your device. No cloud storage. Complete privacy.

SUGGESTED ACTIONS: Find nearby clinic, Set medication reminder, Track symptoms.

LEARNING & UPDATES: Your Agent Is Learning: Personalized to your medical questions and language style. Learning progress: 70% × 1,247 conversations analyzed.

RECENT IMPROVEMENTS: Better cardiology terminology, Improved emergency response time, Enhanced Spanish medical phrasing.

MODEL UPDATES: Update Available: v2.4. New features: Enhanced diagnostic accuracy for respiratory conditions, Expanded medical terminology database. Buttons: Later, Install Update.

USAGE THIS WEEK: CONVERSATIONS: 124 (↑ 18% from last week), AVG SESSION TIME: 4.2m (Typical for medical queries).

DATA & PRIVACY: Your Data, Your Control. Features: Stored locally on device, End-to-end encrypted, No cloud backups, You own all data.

Bottom Bar: Escribe tu pregunta médica... (Type your medical question...), Microphone icon, Red send button.

Top Right Language Buttons: Español, Portugués, Deutsch, Polski, Українська.

Bottom Footer: Model: Medical Assistant ES-MX v2.3 • Running on device • Battery efficient, Data sovereignty guaranteed • GDPR compliant.

Product 2: Implementation of Edge AI Agent Drone Edge Model AI Examples

Design Platform Concept: Our “Ants” Edge AI Models Will Enable Drones To Perform Various Tasks Autonomously



Product 3: Omnestra - The Agent APP Store

Intelligent Agents For Everyone, Connected Through The Queen

Overview: The DEMOS Agent Store brings edge-based AI directly to consumers — a decentralized marketplace where users can download intelligent agents for everyday and professional tasks such as translation, medical note taking, healthcare triage, education, legal matters, taxes, immigration, beauty, logistics, and creative work. Each agent operates locally on the user's device but remains part of the larger Ant Colony Architecture, learning continuously through the Queen Node's coordination and knowledge update

Core Features:

Decentralized AI Marketplace:

- Users download lightweight agents that run fully offline, powered by DEMOS's small-parameter edge model

Inter-Agent Collaboration:

- Agents communicate through the Queen Node to share insights while preserving privac.

Local Intelligence

- Every agent learns from the user's behavior and environment, improving accuracy over time

Privacy by Design

- Data stays on-device — only anonymized learnings are shared with the Queen

Sovereign AI Network

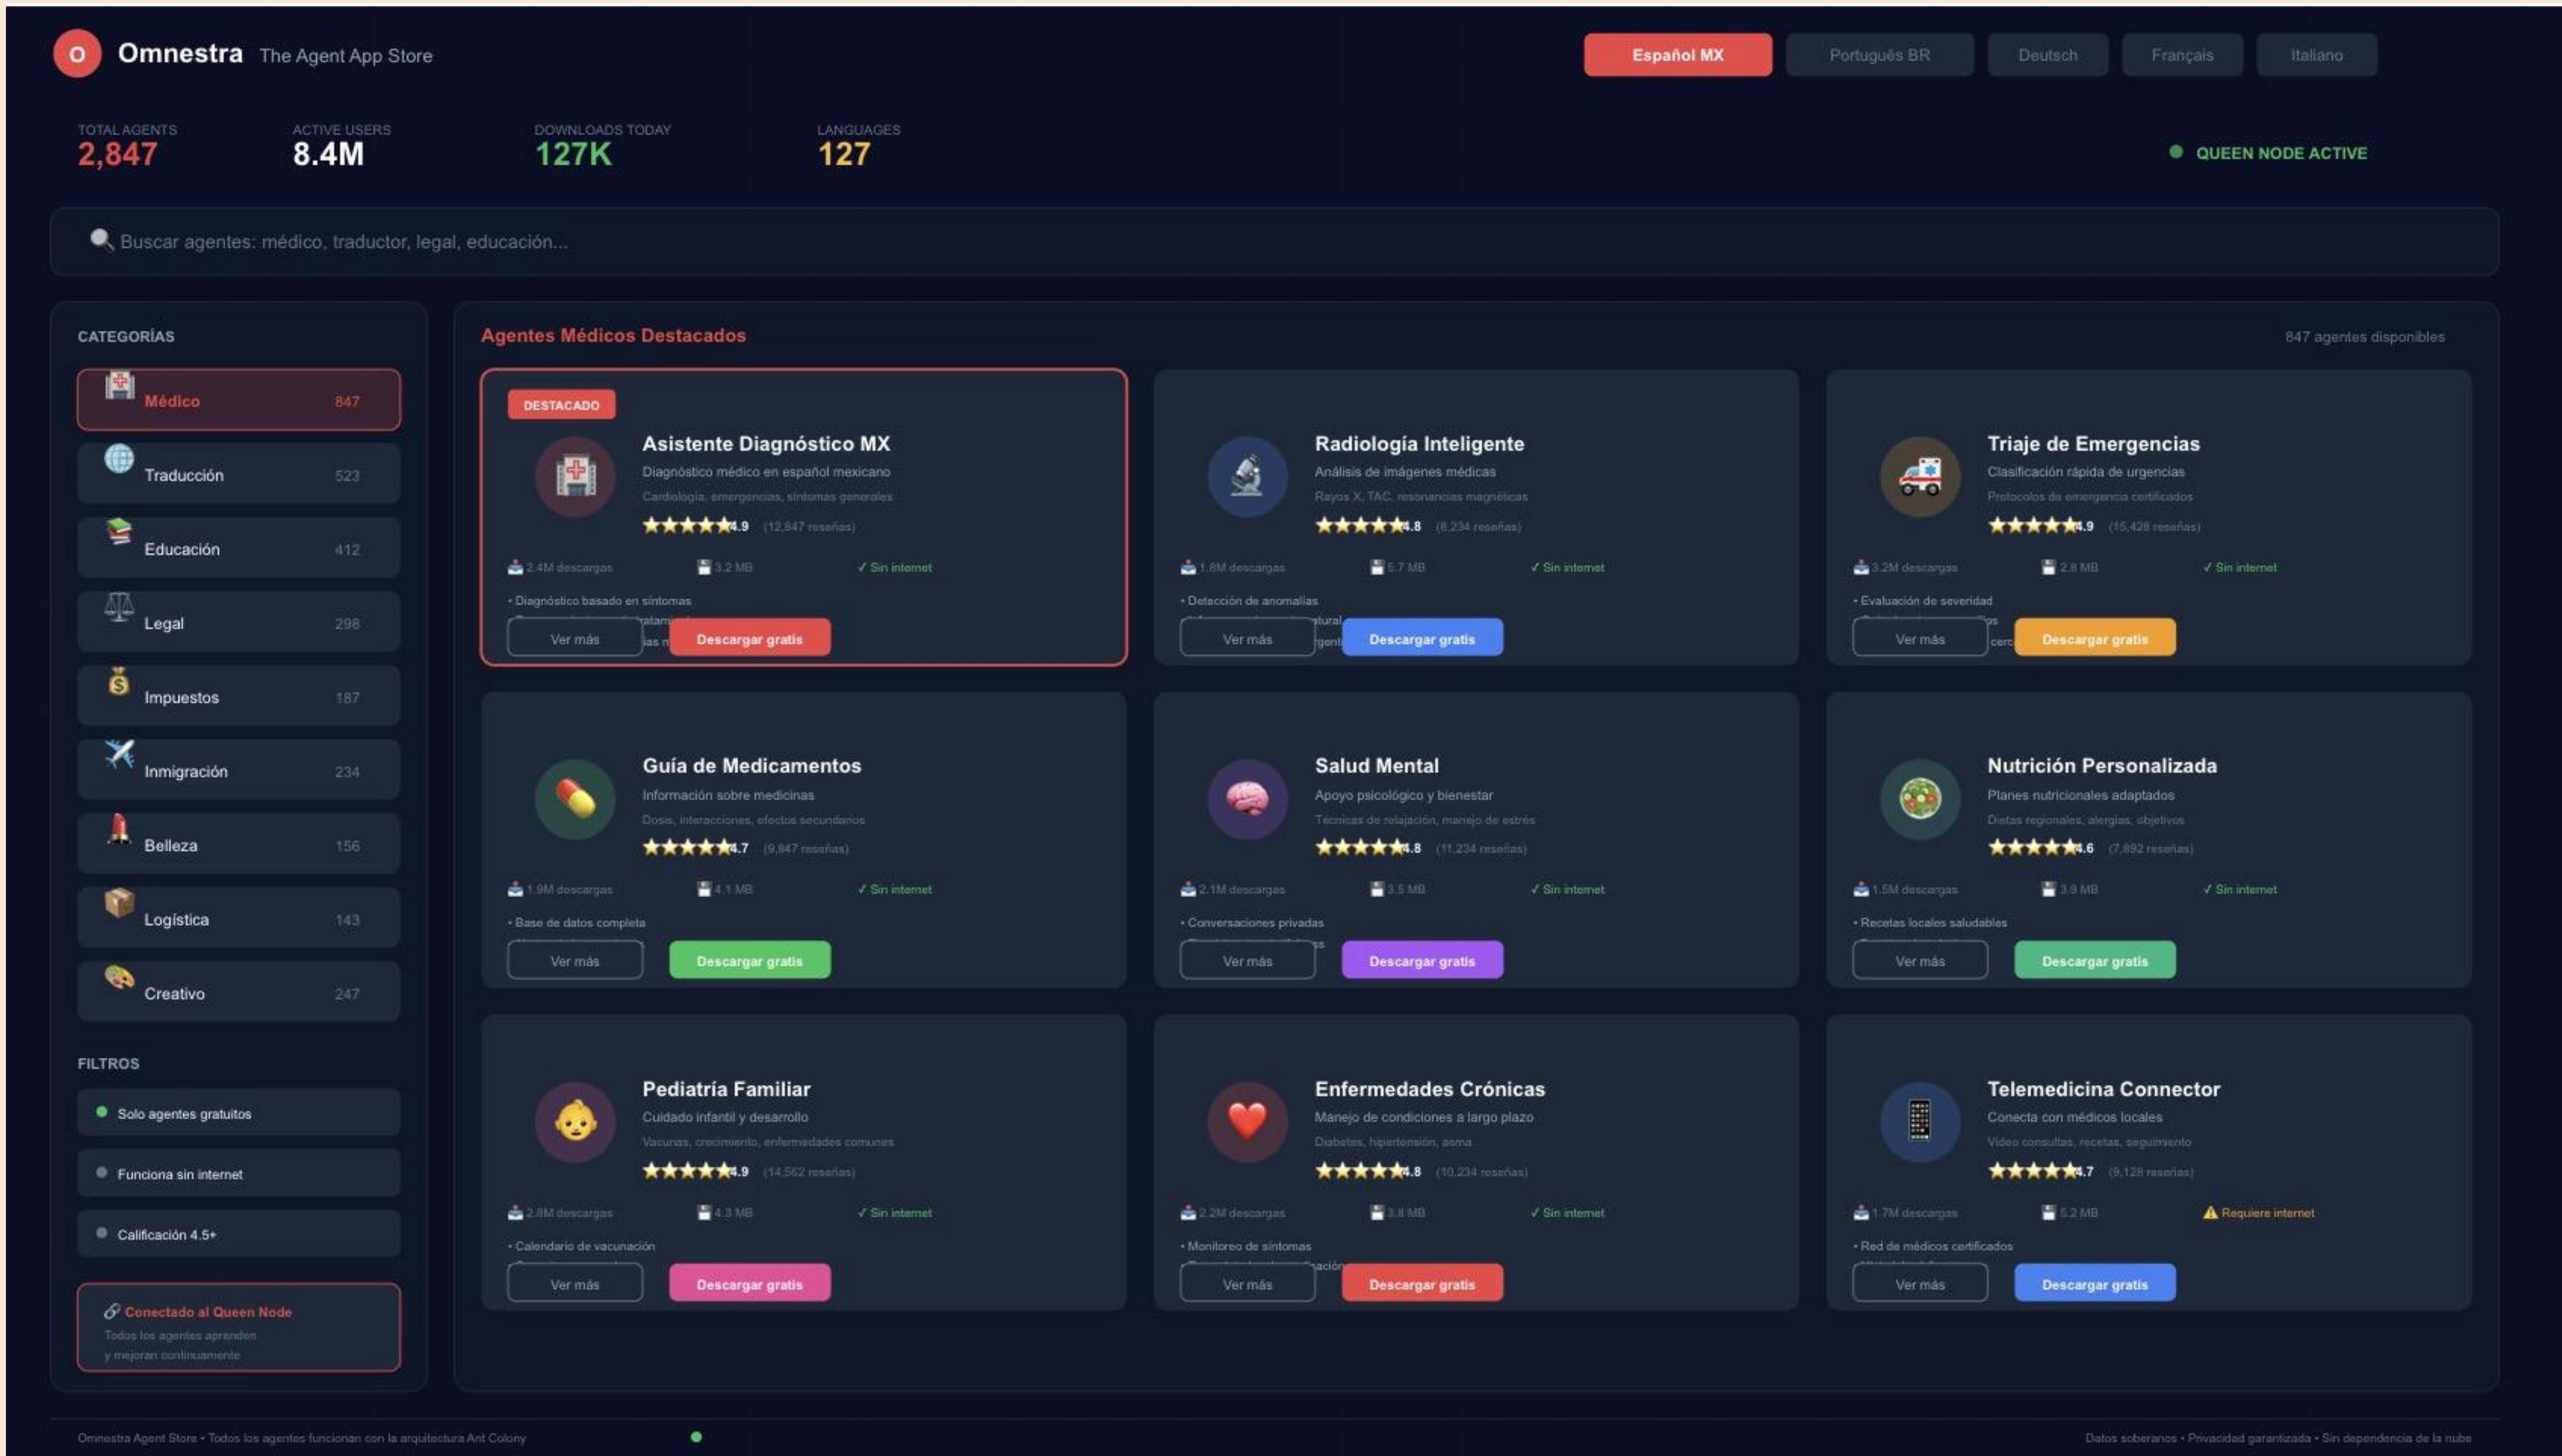
- Each region's app store can host culturally aligned agents, ensuring language and contextual accuracy

Why It Matters: This product closes the loop between infrastructure and user experience — transforming AI from a cloud service into a distributed, living ecosystem that evolves with its user

Product 3: Omnestra - The Agent APP Store

Intelligent Agents For Everyone, Connected Through The Queen

Design Platform Concept: Omnestra App Store Design Example



The image shows a dark-themed user interface for the Omnestra App Store. At the top, there's a header with the Omnestra logo, a search bar, and language selection buttons for Español MX, Portugués BR, Deutsch, Français, and Italiano. Below the header, a banner displays statistics: TOTAL AGENTS 2,847, ACTIVE USERS 8.4M, DOWNLOADS TODAY 127K, and LANGUAGES 127. A green circular icon indicates 'QUEEN NODE ACTIVE'.

The main content area features a sidebar with 'CATEGORÍAS' and 'FILTROS'. The 'CATEGORÍAS' section lists categories like Médico (847), Traducción (523), Educación (412), Legal (298), Impuestos (187), Inmigración (234), Belleza (156), Logística (143), and Creativo (247). The 'FILTROS' section includes options like 'Solo agentes gratuitos' (selected), 'Funciona sin internet', 'Calificación 4.5+', and 'Conectado al Queen Node'.

The main grid displays various intelligent agents:

- Destacados:** Asistente Diagnóstico MX (destacado), Radiología Inteligente, Triage de Emergencias.
- Médicos:** Guía de Medicamentos, Salud Mental, Nutrición Personalizada.
- Familiares:** Pediatría Familiar, Enfermedades Crónicas.
- Otros:** Telemedicina Connector.

Each agent card includes a thumbnail, name, description, star rating, download count, file size, connectivity status, and a 'Descargar gratis' button. A total of 847 agents are available.

At the bottom, there are footer links: Omnestra Agent Store, Datos soberanos, Privacidad garantizada, Sin dependencia de la nube, and a page number 60.

The Business Model

Demos AI Lab Structure

Edgentic Intelligence Division & Rosetta ML Division

DEMOS AI LAB

- The Ant Colony Architecture for Sovereign, Multilingual, On-Device AI
- Ant Colony Architecture: DEMOS AI LAB is the Queen—orchestrating distributed swarms of small, specialized AI agents across billions of devices. Agents can run independently with no central server and no single point of failure, or coordinate as swarms for large-scale collaboration
- Mission: Deliver sovereign AI infrastructure that works on devices people already own, in languages they actually speak, without cloud dependency
- Vision: Build AI for the 6.5 billion people excluded from today's English-first, cloud-dependent AI

Edgentic Intelligence Division

- "Agentic intelligence at the edge"
- Creates small-parameter models, TinyML, and agentic swarms for edge deployment

Rosetta ML Division

- "**Every language, every reality**—powered by curated truth"
- Curates datasets, trains native-language agents, provides training foundation for Edgentic models

The Flow: Rosetta ML curates data and trains → Edgentic Intelligence optimizes for edge → Both contribute to **sovereign App Store**

Edgentic Intelligence Division

Agentic Intelligence At The Edge

Focus:

- Creates TinyML models, small-parameter LLMs, Mamba-based architectures, and agentic agents
- Uses datasets and training from Rosetta ML Division to build edge-optimized models
- Contributes small-parameter models and agentic agents to the sovereign App Store
- Deployment across microcontrollers, smartphones, laptops, drones, and IoT devices

Advantage:

- Sovereign AI agents that run offline, in real time, even during blackouts
- Democratizing compute: Whether it's a \$5 chip or a desktop, intelligence lives on your device

Strategic Value:

- Establishes the swarm model foundation of the Ant Colony
- Enables autonomous networks of lightweight agents
- Applications: healthcare, smart cities, defense, and industrial IoT

Rosetta ML Division

Every Language, Every Reality—powered By Curated Truth

Focus:

- Curates premium datasets and operates large-scale annotation marketplaces
- Trains native-language AI agents in Spanish, Portuguese, Polish, Ukrainian, Arabic, and beyond
- Provides trained models and datasets to Edgentic Intelligence Division for edge optimization
- Contributes multilingual agents and datasets to the sovereign App Store

Advantage:

- Models designed to think natively in each language—not just translate
- Preserves cultural nuance and trust
- Human-in-the-loop expertise ensures datasets are accurate, culturally contextual, and high-quality

Strategic Value:

- Powers the sovereign App Store—distributing multilingual agents, small models, and datasets
- Builds the sovereign dataset backbone for training models and swarm systems
- Datasets distributed for fine-tuning, training, or RAG (retrieval-augmented generation)
- Serves government, healthcare, defense, and vision applications

The Profitability Case – Why Smaller Models Win

Small Models = Better Business Economics

Cost-to-Performance Revolution:

Model Type	Training Cost	Time to Market	Performance	ROI Potential
SLM (1–10B)	\$30K–\$3M	Weeks–Months	85–95% of frontier	HIGH ✓
Frontier LLM	\$100M–\$1B+	18+ months	100% (maybe)	LOW ✗

Why Small Models Are More Profitable:

1. Rapid Iteration: Train 100+ SLMs for the cost of one GPT-5 attempt
2. Lower Risk: \$30K failure vs. \$500M failure
3. Faster Deployment: Weeks to market vs. years
4. Efficient Inference: 187× cheaper to run (Mistral 7B vs. GPT-4)
5. Specialization: Custom models for specific use cases

Real-World Success Stories:

- Microsoft Phi-2: \$30K training → Matches 70B models
- Mistral 7B: Outperforms 13B models at a fraction of cost
- DeepSeek-V3: 18× cheaper than GPT-40 with comparable performance

The Math:

- Train 16,666 Phi-2 models for the cost of one GPT-5 training run
- Break-even time: 3–6 months vs. 3–5 years for frontier models
- Market flexibility: Pivot quickly without billion-dollar sunk costs

Conclusion:

"In AI, smaller is not just smarter – it's more profitable. The era of billion-dollar models may price itself out of viability."

Edge AI

The Small Model Advantage – Training Cost Comparison

Small Parameter Models: The Profitable Path Forward

Training Cost Comparison – Full Spectrum

- TinyML / Edge Models (<10M params): \$100–\$1,000
- Micro SPMs (100M–1B params): \$1,000–\$10,000
- Small SPMs – Phi-2 (7B params): ~\$30,000
- Small SPMs – Mistral 7B (7.3B params): \$50,000–\$150,000 (estimated)
- Medium Models – LLaMA 2 70B: \$5–10 Million
- Large LLMs – GPT-4 (~1T params): \$78–100 Million
- Large LLMs – Gemini Ultra: \$191 Million

Key Insights:

- Small Parameter Models (SPMs) deliver 100×–1,000,000× cost advantage while achieving competitive performance through better data quality and efficient architectures.
- The \$100–\$1,000 Revolution

More Key Insights:

- Fine-tune existing SPMs for specific tasks
- Train edge AI models (TinyML) for IoT devices
- Deploy on-device models with minimal infrastructure

Profitability Metrics

- Sub-\$1K SPMs: Profitable from day one, instant ROI
- Phi-2 (\$30K SPM): Outperforms models 25× larger at 1/3,000th the training cost
- Mistral7B (SPM): 18× cheaper inference cost than GPT-4

ROI Timeline: SPMs reach profitability in days/weeks; Large LLMs need years

Bottom Line: You can train effective Small Parameter Models for less than a laptop costs, while frontier LLMs require \$100M+ with uncertain returns.

Terminology Clarification

- SPM = Small Parameter Model (typically <30B parameters)
- LLM = Large Language Model (typically 50B+ parameters, frontier models)
- TinyML = Ultra-small models for edge devices (<10M parameters)

GPT-5 (Orion) – The \$500 Million+ Challenge

GPT-5: When Training Costs Spiral Out of Control

The Staggering Numbers:

- Single Training Run: \$500+ Million (6 months of compute alone)
- Total Development Cost: Approaching \$1+ Billion
- Failed Attempts: At least 2 major training runs have failed
- Cost per Failure: “Like a space rocket exploding after launch”

What \$500M+ Buys:

- Months of computing on tens of thousands of Nvidia chips
- 13+ trillion tokens of training data
- Hundreds of expert researchers and engineers
- Synthetic data generation from mathematicians and physicists

The Problem:

- Despite massive investment, GPT-5 “hasn’t advanced enough to justify the enormous cost.”
- Performance gains not matching cost increases
- Multiple failed training runs multiply total expenses

Industry Reality Check:

- “A six-month training run can cost over \$500 Million in computing costs alone, let alone the rest of the expenditure.”— WSJ Report

The Plateau:

- Even \$500M+ investments are hitting diminishing returns as the “more-is-more” strategy runs out of steam.

The \$100 Fine-Tuning Revolution – Real World Examples

How Academics Built ChatGPT-Like Models for Under \$600

Case Study 1: Stanford Alpaca (2023)

The Achievement:

- Created a ChatGPT-like conversational model from LLaMA-7B
- 7 billion parameter instruction-following model
- Performance competitive with much larger commercial models

The Cost Breakdown:

- GPU Training: 3 hours on 8×A100 80GB = <\$100
- Dataset Creation: 52K instruction-response pairs via OpenAI API = ~\$500
- TOTAL COST: ≈ \$600

The Strategy:

- Started with pre-trained LLaMA-7B (Meta's open weights)
- Used GPT-3.5 to generate synthetic training data (\$0.02 / 1K tokens)
- Fine-tuned for instruction-following in just 3 hours

The Impact:

"An academic lab produced a ChatGPT-like model for a few hundred dollars — proving fine-tuning can democratize AI."

Databricks Dolly – AI for Less Than \$30

When Building AI Costs Less Than Dinner

Case Study 2: Databricks Dolly (2023)

Dolly 1.0 – The Achievement

- Based on EleutherAI's GPT-J-6B (6 billion parameters)
- Instruction-tuned conversational AI model
- Open source and commercially usable

The Cost Comparison:

- Dolly 1.0: \$30
- Dolly 2.0: \$200–\$300
- Alpaca: \$600
- Train 7B from scratch: \$1–2 Million
- GPT-4: \$78–100 Million

The Cost Breakdown:

- Hardware: Single node with 4×A100 GPUs
- Training Time: A few hours
- Dataset: 15,000 instructions
- Fine-Tuning Cost: < \$30

The Impact:

"Databricks proved you can add AI capabilities to existing models for less than the cost of a nice dinner – democratizing access to custom AI for every business."

Dolly 2.0 – Scaling Up

- Based on EleutherAI's Pythia-12B (12 billion parameters)
- Fine-Tuning Cost: \$200–\$300
- Training Time: Hours to days (not months)

The Great Divide

Fine-Tuning vs. Frontier Models

\$600 Success Stories vs. \$500 Million Struggles

The Cost Comparison:

Model	Approach	Cost	Time	Result
Stanford Alpaca	Fine-tune 7B	\$600	3 hours	ChatGPT-like performance ✓
Databricks Dolly	Fine-tune 6B	\$30	Hours	Instruction-following AI ✓
GPT-5 (Orion)	Train from scratch	\$500M+ per run	6+ months	Failed twice X

The Reality Check:

- Alpaca + Dolly combined: \$630 total investment, 2 successful models
- GPT-5: \$500M+ per attempt, multiple failures, still not released
- Cost ratio: GPT-5 costs 793,650× more than Alpaca
- Success rate: Fine-tuning wins immediately; GPT-5 struggles after 18+ months

What \$500 Million Could Buy Instead:

- Train 833,333 Alpaca models
- Train 16,666,666 Dolly models
- Fund 833 startups with custom AI models
- Fine-tune a specialized model for every use case imaginable

The Performance Paradox:

- Alpaca (\$600): Works, deployed, being used
- Dolly (\$30): Works, deployed, commercially available
- GPT-5 (\$500M+): "Hasn't advanced enough to justify the enormous cost"

The Impact:

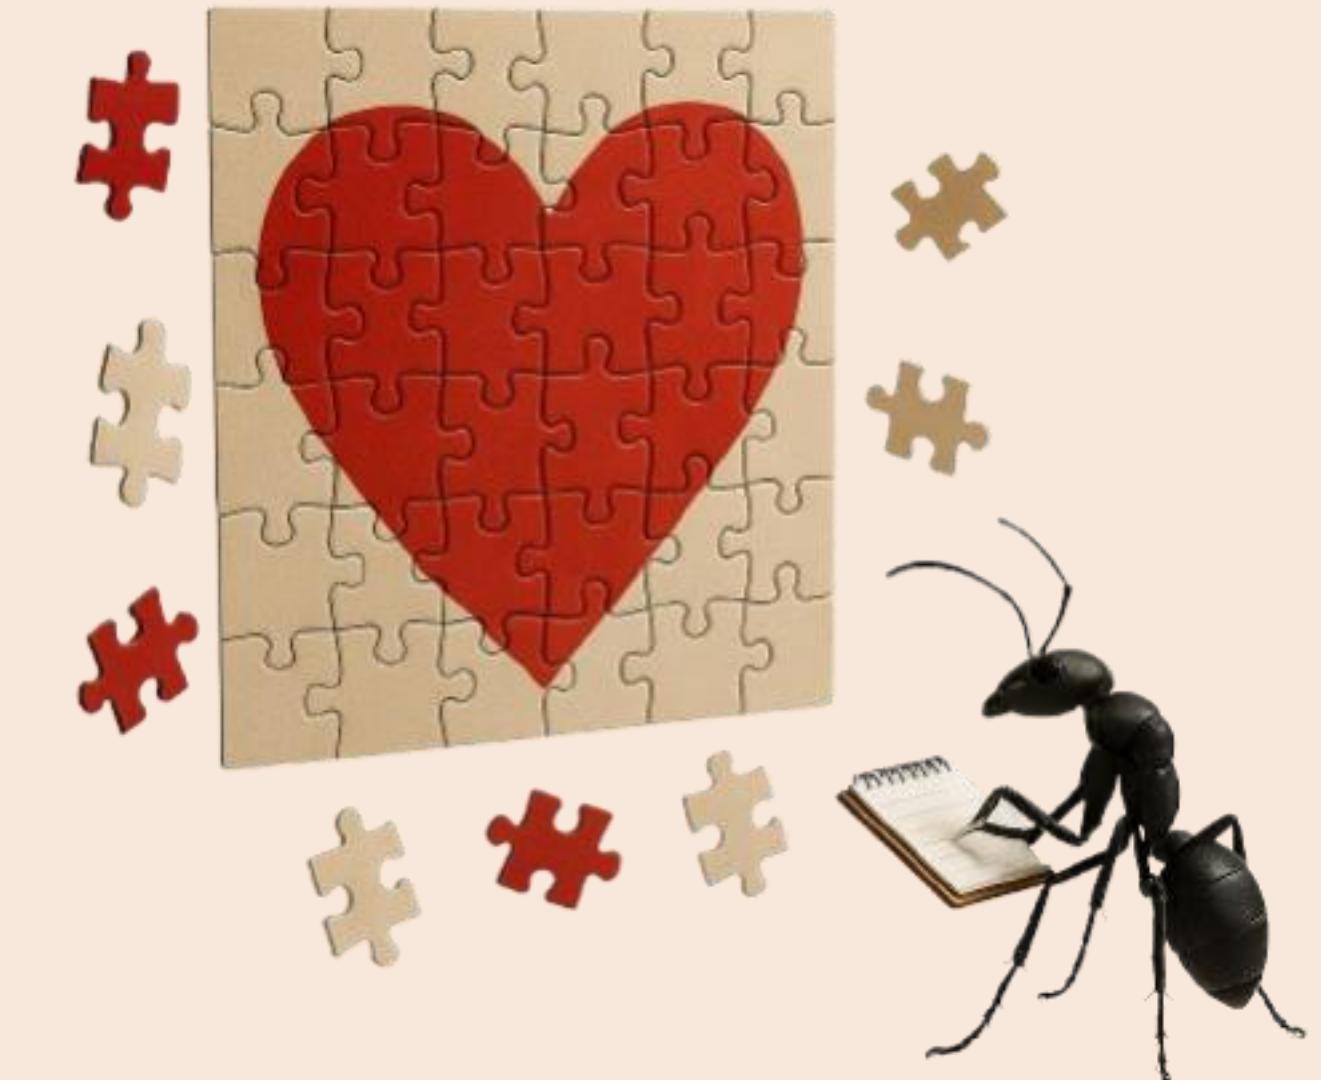
"OpenAI spent 793,650× more than Stanford—and still hasn't achieved a successful GPT-5 launch. Meanwhile, academic labs built working models for the cost of a laptop."

Sources: Stanford CRFM, Databricks, WSJ (GPT-5 reporting)

Revolutionary **Unit Economics**

95% Margins Because No Infrastructure Needed

- **Traditional LLM AI Costs:** Data centers, cooling, electricity = 60% of revenue
- **Small Parameter Models:** Runs on phones users already own = 5% costs
- **Result:** 95% gross margins on software sales
- **Proven Comparables:**
 - Palantir:** \$443B market cap (proven government contract model—we have the **experience and relationships to replicate**)
 - Mistral AI:** \$14B valuation (October 2024 edge pivot validates our thesis)
 - Auterion:** \$130M Series B at \$600M+ valuation (Sept 2025 Edge AI defense validates our government thesis)
 - Scale AI:** \$29B valuation (English only—we're doing ALL languages)
- **Key Insight:** **Palantir** proves AI + Government Contracting = massive value. **Scale AI's** recent Meta acquisition has created a vacuum—enterprise clients are exiting due to IP concerns, top talent like Director of Product Kate Park is seeking exits, and \$29B worth of AI training contracts are in play. We're positioning to capture this team and these contracts.



Eight Revenue Streams Per Language

One Platform, Eight Ways To Monetize—**Multiplied By Each Language**

Training AI Services:

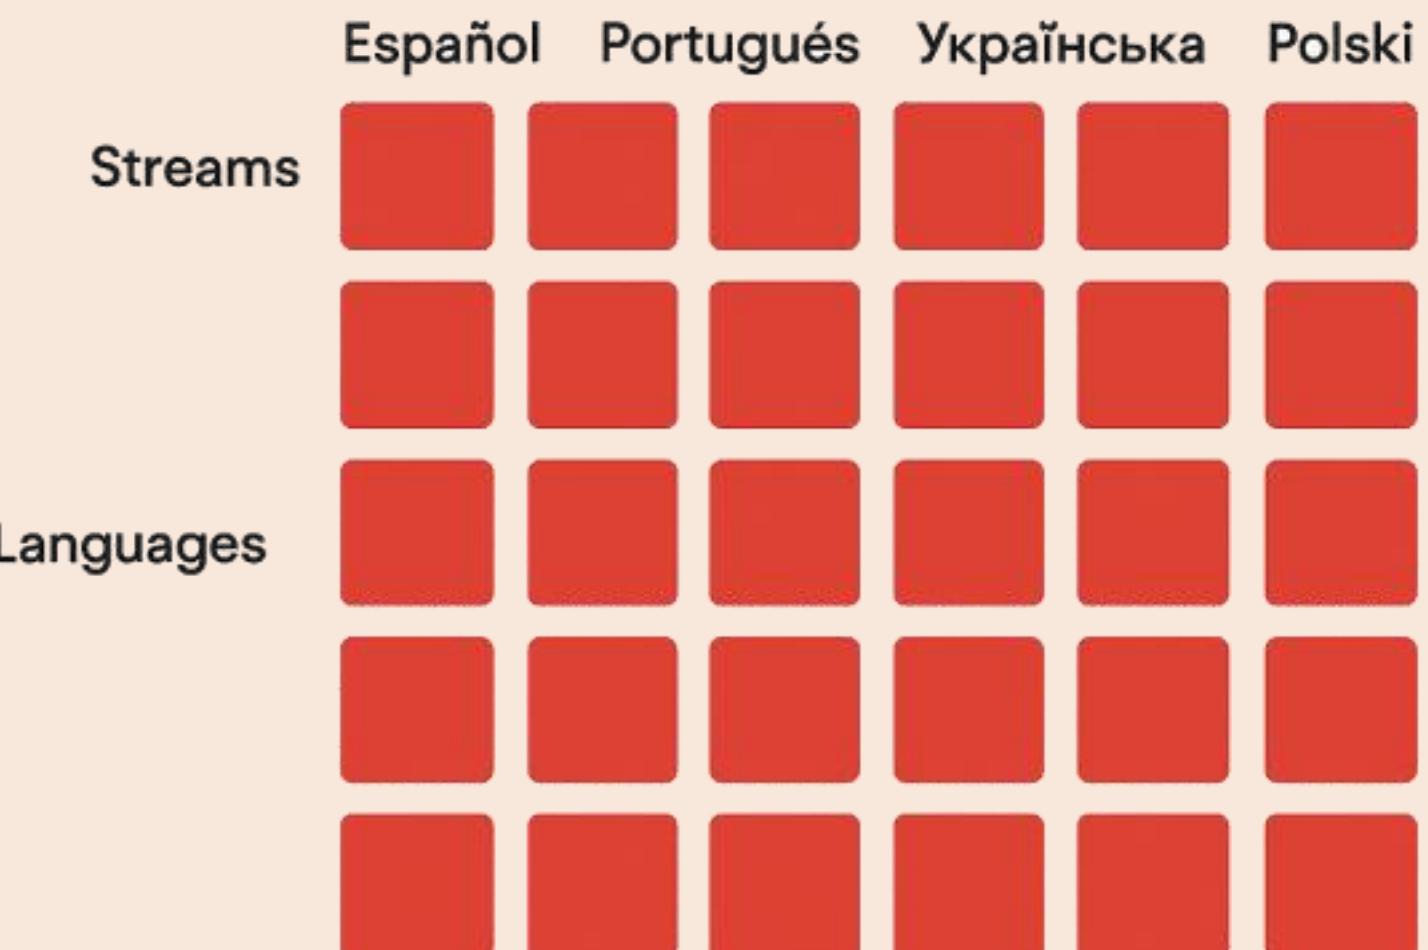
- **Data Annotation Contracts:** \$1-100M per year from US Government, OpenAI, Google, Apple, Claude
- **Premium Datasets:** \$1-10M per specialized corpus (medical, legal, financial)
- **Marketplace Fees:** 20% commission on annotation tasks

Edgentic Sales:

- **Omnestra APP Store Downloads—Consumer Edgentic Agents App Store:** \$1-25 one-time downloads (5% penetration = \$750M)
- **Omnestra APP Store Subscriptions: Connects Consumer Agent To Ant Colony Architecture & Queen**
- **Enterprise Edgentic Licensing & Contracts:** \$500K-5M/year for on-premise AI
- **Government Edgentic Licensing & Contracts:** \$10-100M for sovereign AI (Palantir-style)

Physical AI:

- **Physical AI/Drones:** Deployment & Integration Services – AI software kits + Military (NATO/allied defense) + Commercial (skyscraper maintenance, infrastructure, agriculture) | + recurring model update



Financial Revenue Projections: Aggressive Case

DEMOS AI LAB Yearly Projections (Millions)

DEMOS AI Projected Revenues (Millions): Aggressive Case								
Employee & Other Costs	2026	2027	2028	2029	2030	2031	2032	2033
Training AI Services (millions)								
Data Annotation Contracts (Enterprise)	\$0.50	\$1.00	\$2.00	\$4.00	\$5.36	\$7.18	\$9.62	\$12.90
Data Annotation Contracts (Government)	\$0.50	\$1.00	\$2.00	\$4.00	\$5.36	\$7.18	\$9.62	\$12.90
Data Annotation Contracts (NGO)	\$0.00	\$0.50	\$1.00	\$2.00	\$2.68	\$3.59	\$4.81	\$6.45
Premium Datasets	\$0.00	\$1.00	\$2.00	\$4.00	\$5.36	\$7.18	\$9.62	\$12.90
Marketplace Fees	\$0.20	\$0.50	\$1.00	\$2.00	\$2.68	\$3.59	\$4.81	\$6.45
Subtotal:	\$1.20	\$4.00	\$8.00	\$16.00	\$21.44	\$28.73	\$38.50	\$51.59
Edgentic Sales (millions)								
Edgentic Licensing & Contracts (Enterprise)	\$0.00	\$1.00	\$2.00	\$4.00	\$8.00	\$16.00	\$18.22	\$20.76
Edgentic Licensing & Contracts (Government)	\$0.00	\$2.00	\$5.00	\$12.00	\$30.00	\$32.00	\$36.45	\$41.51
Omnestra App Store Downloads (actual number)	0	0	300,000	250,000	1,250,000	3,000,000	6,000,000	10,000,000
Omnestra App Store Downloads Fees (\$20 per)	\$0.00	\$0.00	\$6.00	\$5.00	\$25.00	\$60.00	\$120.00	\$200.00
Omnestra App Store Subscriptions	\$0.00	\$0.00	\$1.20	\$1.00	\$5.00	\$12.00	\$24.00	\$40.00
Subtotal:	\$0.00	\$3.00	\$14.20	\$22.00	\$68.00	\$120.00	\$198.67	\$302.27
Physical AI (millions)								
AI Software Kits (Drones)	\$0.00	\$1.00	\$2.00	\$4.00	\$8.00	\$16.00	\$30.00	\$60.00
Model Deployment & Integration Services	\$0.00	\$1.00	\$2.00	\$4.00	\$8.00	\$16.00	\$32.00	\$70.00
Infrastructure & Agriculture	\$0.00	\$1.00	\$2.00	\$4.00	\$8.00	\$16.00	\$25.00	\$34.00
Recurring Model Update Subscription Fees	\$0.00	\$0.75	\$1.50	\$3.00	\$6.00	\$12.00	\$21.75	\$41.00
Subtotal:	\$0.00	\$3.75	\$7.50	\$15.00	\$30.00	\$60.00	\$108.75	\$205.00
Total:	\$1.20	\$10.75	\$29.70	\$53.00	\$119.44	\$208.73	\$345.92	\$558.86
CAGR (Edge AI)						13.90%	13.90%	13.90%
CAGR (Data Annotation)					34.00%	34.00%	34.00%	34.00%
Marketplace fees	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%

Financial Revenue Projections: Conservative Case

DEMOS AI LAB Yearly Projections (Millions)

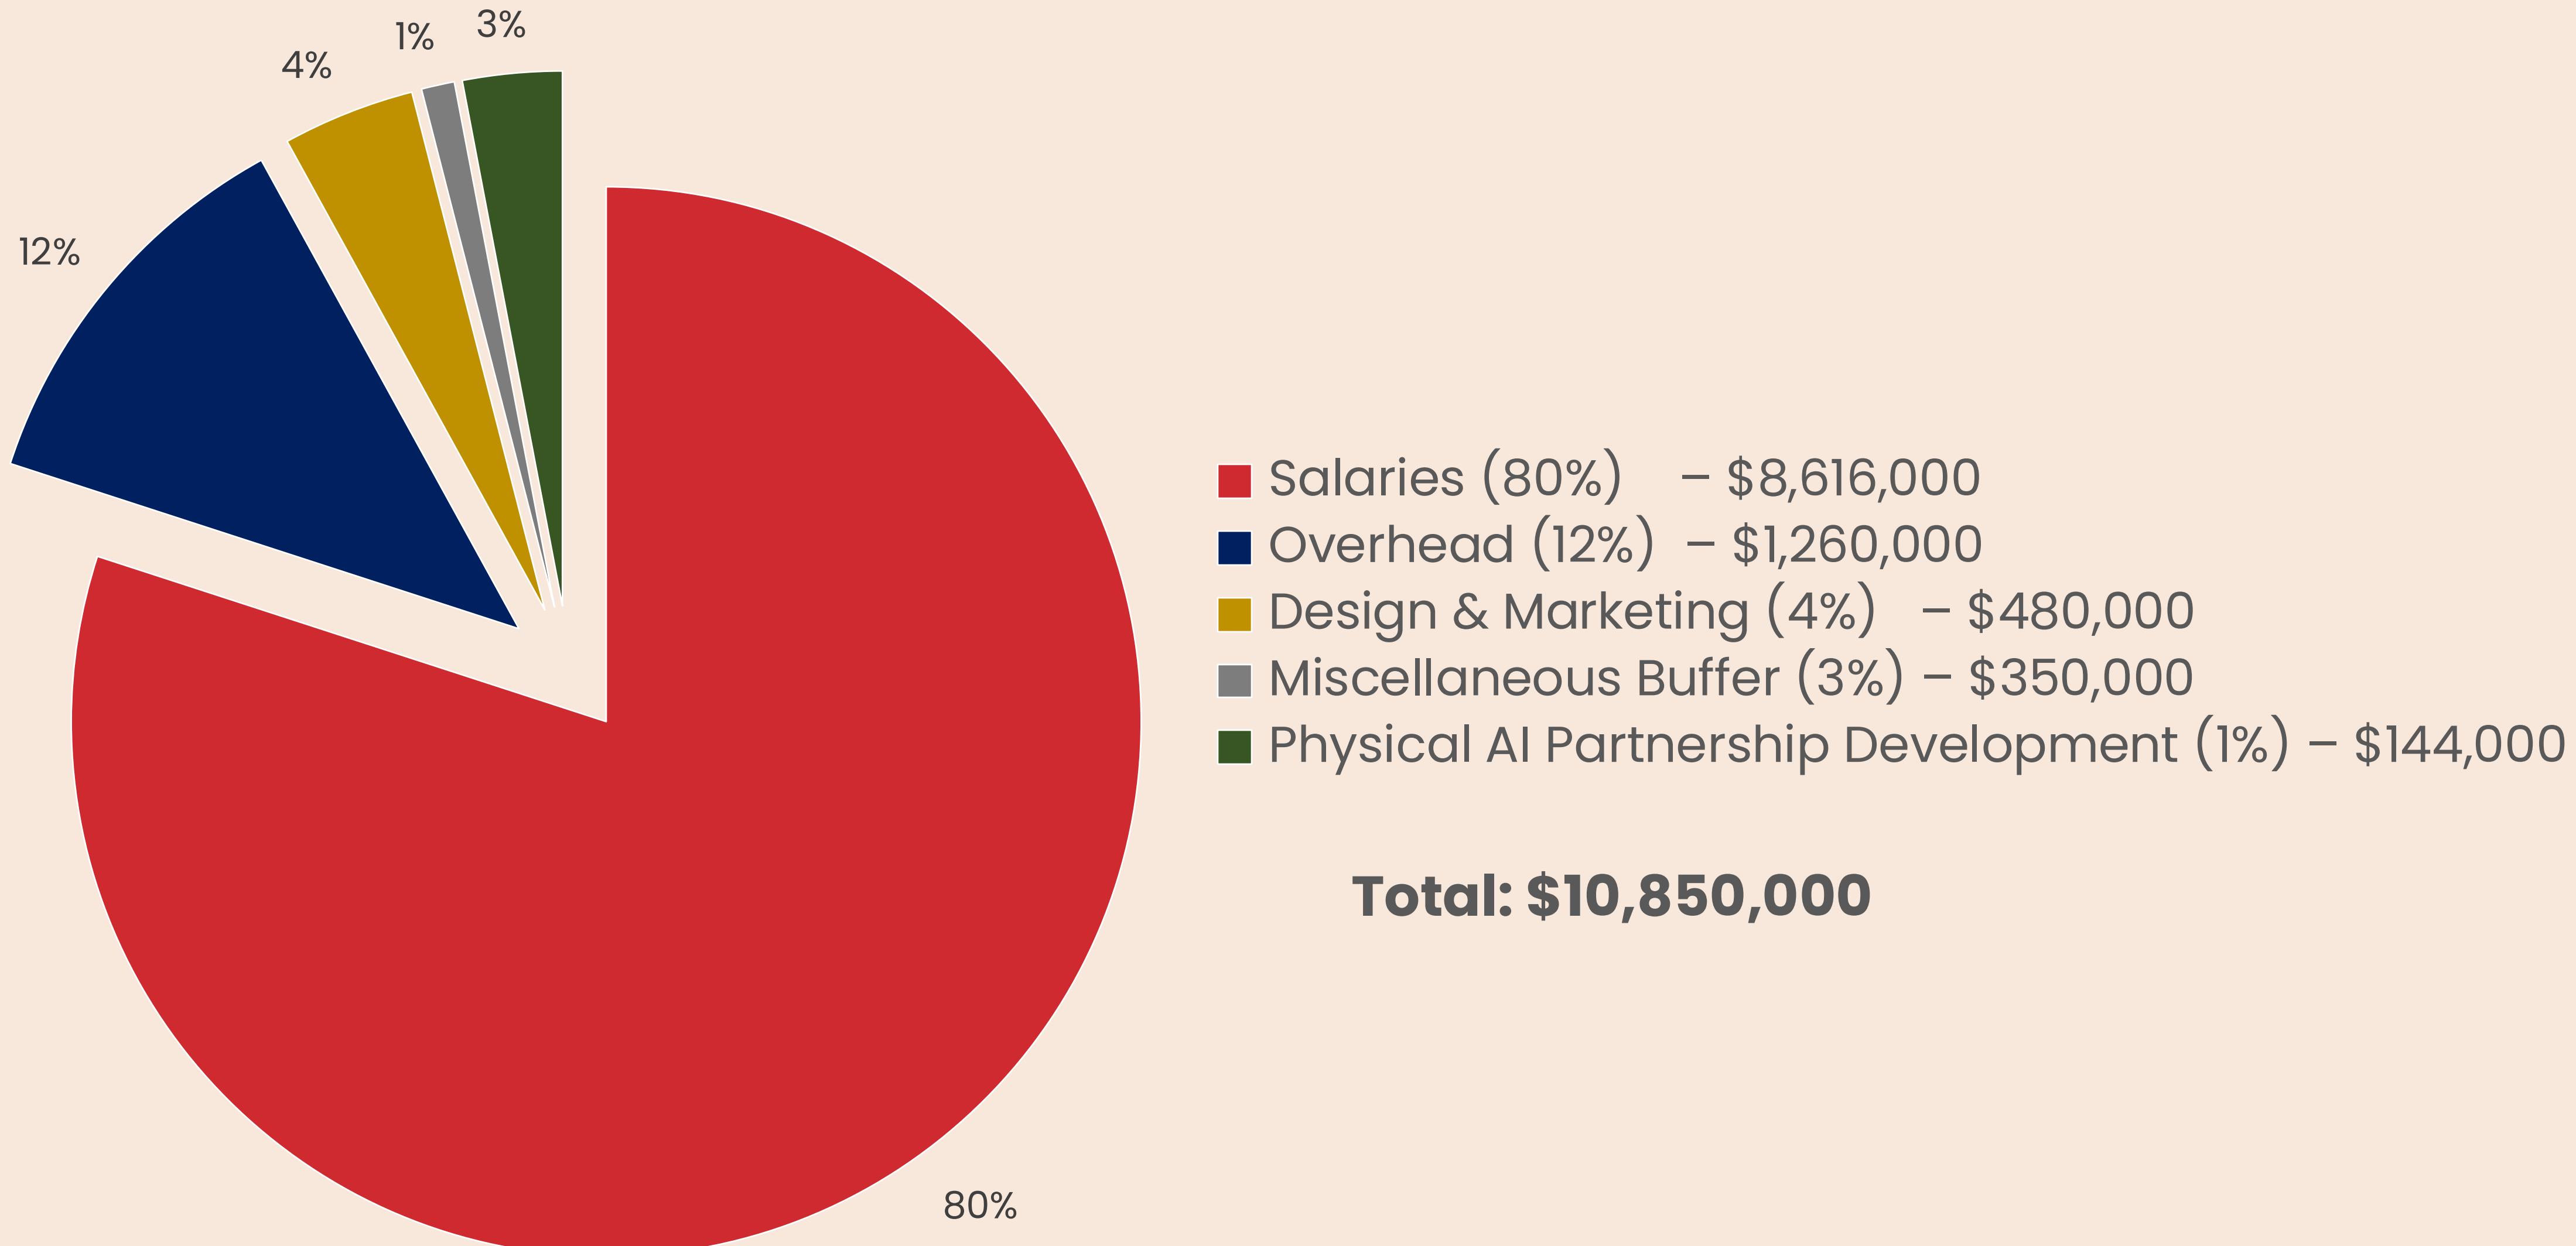
DEMOS AI Projected Revenues (Millions): Conservative Case								
Employee & Other Costs	2026	2027	2028	2029	2030	2031	2032	2033
<u>Training AI Services (millions)</u>								
Data Annotation Contracts (Enterprise)	\$0.00	\$1.00	\$2.00	\$3.00	\$4.02	\$5.39	\$7.22	\$9.67
Data Annotation Contracts (Government)	\$0.00	\$1.00	\$2.00	\$3.00	\$4.02	\$5.39	\$7.22	\$9.67
Data Annotation Contracts (NGO)	\$0.00	\$0.00	\$1.00	\$2.00	\$2.68	\$3.59	\$4.81	\$6.45
Premium Datasets	\$0.00	\$0.50	\$1.00	\$2.00	\$2.68	\$3.59	\$4.81	\$6.45
Marketplace Fees	\$0.00	\$0.40	\$1.00	\$1.60	\$2.14	\$2.87	\$3.85	\$5.16
Subtotal:	\$0.00	\$2.90	\$7.00	\$11.60	\$15.54	\$20.83	\$27.91	\$37.40
<u>Edgentic Sales (millions)</u>								
Edgentic Licensing & Contracts (Enterprise)	\$0.00	\$1.00	\$1.75	\$3.50	\$7.00	\$14.00	\$15.95	\$18.16
Edgentic Licensing & Contracts (Government)	\$0.00	\$1.00	\$3.00	\$7.00	\$14.00	\$28.00	\$31.89	\$36.32
Omnestra App Store Downloads (actual number)	0	0	25,000	50,000	250,000	600,000	1,200,000	2,000,000
Omnestra App Store Downloads Fees (\$20 per)	\$0.00	\$0.00	\$0.50	\$1.00	\$5.00	\$12.00	\$24.00	\$40.00
Omnestra App Store Subscriptions	\$0.00	\$0.00	\$0.10	\$0.20	\$1.00	\$2.40	\$4.80	\$8.00
Subtotal:	\$0.00	\$2.00	\$5.35	\$11.70	\$27.00	\$56.40	\$76.64	\$102.49
<u>Physical AI (millions)</u>								
AI Software Kits (Drones)	\$0.00	\$0.50	\$1.00	\$2.00	\$4.00	\$7.00	\$15.00	\$30.00
Model Deployment & Integration Services	\$0.00	\$0.50	\$1.00	\$2.00	\$4.00	\$8.00	\$16.00	\$32.00
Infrastructure & Agriculture	\$0.00	\$0.50	\$1.00	\$2.00	\$3.00	\$4.00	\$5.00	\$6.00
Recurring Model Update Subscription Fees	\$0.00	\$0.38	\$0.75	\$1.50	\$2.75	\$4.75	\$9.00	\$17.00
Subtotal:	\$0.00	\$1.88	\$3.75	\$7.50	\$13.75	\$23.75	\$45.00	\$85.00
Total:	\$0.00	\$6.78	\$16.10	\$30.80	\$56.29	\$100.98	\$149.55	\$224.89
CAGR (Edge AI)						13.90%	13.90%	13.90%
CAGR (Data Annotation)						34.00%	34.00%	34.00%
Marketplace fees	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%

Financial Cost Projections

DEMOS AI LAB Yearly Projections (Millions)

Employee & Other Costs	Quantity	Monthly	Yearl 0	Year 1	Year 2
Team Salaries					
C0-CEO	1	\$30,000.00	\$360,000.00	\$396,000.00	\$792,000.00
CO-CEO : Governemnt Focised	1	\$30,000.00	\$360,000.00	\$396,000.00	\$792,000.00
CPO (Chief Product Officer)	1	\$100,000.00	\$1,200,000.00	\$1,320,000.00	\$2,640,000.00
CTO	1	\$65,000.00	\$780,000.00	\$858,000.00	\$1,716,000.00
Chief Scientist	1	\$100,000.00	\$1,200,000.00	\$1,320,000.00	\$2,640,000.00
PhD AI Scientists	4	\$65,000.00	\$3,120,000.00	\$3,432,000.00	\$6,864,000.00
Lead Full Stack Developer	1	\$20,000.00	\$240,000.00	\$264,000.00	\$528,000.00
Senior Front End Developer	1	\$18,000.00	\$216,000.00	\$237,600.00	\$475,200.00
Lead ML Engineer	1	\$20,000.00	\$240,000.00	\$264,000.00	\$528,000.00
Mid-Level ML Engineer	1	\$18,000.00	\$216,000.00	\$237,600.00	\$475,200.00
Lead Data & Back End Engineer	1	\$20,000.00	\$240,000.00	\$264,000.00	\$528,000.00
Senior Back End Engineer	1	\$19,000.00	\$228,000.00	\$250,800.00	\$501,600.00
Phd Spanish NLP Linguist	1	\$18,000.00	\$216,000.00	\$237,600.00	\$475,200.00
Subtotal:	16	\$523,000.00	\$8,616,000.00	\$9,477,600.00	\$18,955,200.00
Design & Marketing					
Design Firm	1	\$20,000.00	\$240,000.00	\$264,000.00	\$528,000.00
Marketing	1	\$20,000.00	\$240,000.00	\$264,000.00	\$528,000.00
Subtotal:	2	\$40,000.00	\$480,000.00	\$528,000.00	\$1,056,000.00
Overhead					
Hardware	1	\$40,000.00	\$480,000.00	\$528,000.00	\$1,056,000.00
Developer Tools	1	\$25,000.00	\$300,000.00	\$330,000.00	\$660,000.00
Cloud	1	\$20,000.00	\$240,000.00	\$264,000.00	\$528,000.00
Legal & Accounting	1	\$10,000.00	\$120,000.00	\$132,000.00	\$264,000.00
Office Rent & Misc	1	\$10,000.00	\$120,000.00	\$132,000.00	\$264,000.00
Subtotal:	5	\$105,000.00	\$1,260,000.00	\$1,386,000.00	\$2,772,000.00
Physical AI Partnership Development					
Drone Prototypes & Partnerships	1	\$12,000.00	\$144,000.00	\$158,400.00	\$316,800.00
Subtotal:	1	\$12,000.00	\$144,000.00	\$158,400.00	\$316,800.00
Miscellaneous Buffer					
Data Annotation Contractor Services	1	\$25,000.00	\$300,000.00	\$330,000.00	\$660,000.00
Miscellaneous Buffer	1	\$5,000.00	\$50,000.00	\$55,000.00	\$110,000.00
Subtotal:	2	\$30,000.00	\$350,000.00	\$385,000.00	\$770,000.00
Total:		\$710,000.00	\$10,850,000.00	\$11,935,000.00	\$23,870,000.00
Cost Increase %				10.00%	100.00%

Financial Cost Breakdown



Financial EBITDA Projections: Aggressive Case

DEMOS AI LAB Yearly Projections (Millions)

DEMOS AI Projected EBITDA (Millions): Aggressive Case								
Employee & Other Costs	2026	2027	2028	2029	2030	2031	2032	2033
Revenue								
Training AI Services (millions)	\$1.20	\$4.00	\$8.00	\$16.00	\$21.44	\$28.73	\$38.50	\$51.59
Edgentic Sales (millions)	\$0.00	\$3.00	\$14.20	\$22.00	\$68.00	\$120.00	\$198.67	\$302.27
Physical AI (millions)	\$0.00	\$3.75	\$7.50	\$15.00	\$30.00	\$60.00	\$108.75	\$205.00
Revenue Total	\$1.20	\$10.75	\$29.70	\$53.00	\$119.44	\$208.73	\$345.92	\$558.86
Gross Margin After Training Costs								
Training AI Services (millions)	\$1.14	\$3.80	\$7.60	\$15.20	\$20.37	\$27.29	\$36.57	\$49.01
Edgentic Sales (millions)	\$0.00	\$2.85	\$13.49	\$20.90	\$64.60	\$114.00	\$188.74	\$287.16
Physical AI (millions)	\$0.00	\$3.56	\$7.13	\$14.25	\$28.50	\$57.00	\$103.31	\$194.75
Gross Margin Total (95%)	\$1.14	\$10.21	\$28.22	\$50.35	\$113.47	\$198.29	\$328.62	\$530.92
Costs Before EBITDA								
Team Salaries	\$8.62	\$9.48	\$18.96	\$22.75	\$27.30	\$32.75	\$39.31	\$47.17
Design & Marketing	\$0.48	\$0.53	\$1.06	\$1.27	\$1.52	\$1.82	\$2.19	\$2.63
Overhead	\$1.26	\$1.39	\$2.77	\$3.33	\$3.99	\$4.79	\$5.75	\$6.90
Physical AI Partnership Development	\$0.14	\$0.16	\$0.32	\$0.38	\$0.46	\$0.55	\$0.66	\$0.79
Miscellaneous Buffer	\$0.35	\$0.39	\$0.77	\$0.92	\$1.11	\$1.33	\$1.60	\$1.92
Costs Before EBITDA	\$10.85	\$11.94	\$23.87	\$28.64	\$34.37	\$41.25	\$49.50	\$59.40
EBITDA	-\$9.71	-\$1.72	\$4.35	\$21.71	\$79.10	\$157.05	\$279.13	\$471.52
CAGR (Edge AI)						13.90%	13.90%	13.90%
CAGR (Data Annotation)					34.00%	34.00%	34.00%	34.00%
Marketplace fees	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%
Gross Margin After Training Costs	95.00%	95.00%	95.00%	95.00%	95.00%	95.00%	95.00%	95.00%
Cost Increase Per Year	0.00%	10.00%	100.00%	20.00%	20.00%	20.00%	20.00%	20.00%

Financial EBITDA Projections: Conservative Case

DEMOS AI LAB Yearly Projections (Millions)

DEMOS AI Projected EBITDA (Millions): Conservative Case								
Employee & Other Costs	2026	2027	2028	2029	2030	2031	2032	2033
Revenue								
Training AI Services (millions)	\$0.00	\$2.90	\$7.00	\$11.60	\$15.54	\$20.83	\$27.91	\$37.40
Edgentic Sales (millions)	\$0.00	\$2.00	\$5.35	\$11.70	\$27.00	\$56.40	\$76.64	\$102.49
Physical AI (millions)	\$0.00	\$1.88	\$3.75	\$7.50	\$13.75	\$23.75	\$45.00	\$85.00
Revenue Total	\$0.00	\$6.78	\$16.10	\$30.80	\$56.29	\$100.98	\$149.55	\$224.89
Gross Margin After Training Costs								
Training AI Services (millions)	\$0.00	\$2.76	\$6.65	\$11.02	\$14.77	\$19.79	\$26.52	\$35.53
Edgentic Sales (millions)	\$0.00	\$1.90	\$5.08	\$11.12	\$25.65	\$53.58	\$72.81	\$97.36
Physical AI (millions)	\$0.00	\$1.78	\$3.56	\$7.13	\$13.06	\$22.56	\$42.75	\$80.75
Gross Margin Total (95%)	\$0.00	\$6.44	\$15.30	\$29.26	\$53.48	\$95.93	\$142.07	\$213.64
Costs Before EBITDA								
Team Salaries	\$8.62	\$9.48	\$18.96	\$22.75	\$27.30	\$32.75	\$39.31	\$47.17
Design & Marketing	\$0.48	\$0.53	\$1.06	\$1.27	\$1.52	\$1.82	\$2.19	\$2.63
Overhead	\$1.26	\$1.39	\$2.77	\$3.33	\$3.99	\$4.79	\$5.75	\$6.90
Physical AI Partnership Development	\$0.14	\$0.16	\$0.32	\$0.38	\$0.46	\$0.55	\$0.66	\$0.79
Miscellaneous Buffer	\$0.35	\$0.39	\$0.77	\$0.92	\$1.11	\$1.33	\$1.60	\$1.92
Costs Before EBITDA	\$10.85	\$11.94	\$23.87	\$28.64	\$34.37	\$41.25	\$49.50	\$59.40
EBITDA	-\$10.85	-\$5.50	-\$8.58	\$0.62	\$19.11	\$54.68	\$92.57	\$154.25
CAGR (Edge AI)						13.90%	13.90%	13.90%
CAGR (Data Annotation)					34.00%	34.00%	34.00%	34.00%
Marketplace fees	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%	20.00%
Gross Margin After Training Costs	95.00%	95.00%	95.00%	95.00%	95.00%	95.00%	95.00%	95.00%
Cost Increase Per Year		10.00%	100.00%	20.00%	20.00%	20.00%	20.00%	20.00%

The Roadmap

Global Language Roadmap: 3 Phases

Three Phases to Build the World's First Sovereign Multilingual Edge AI Network

Each phase expands DEMOS's sovereign data, model, and agent ecosystems — compounding toward a \$100B global AI infrastructure opportunity

- **Phase 1 (Years 0–2): Core Europe + Foundational Latin American Languages**

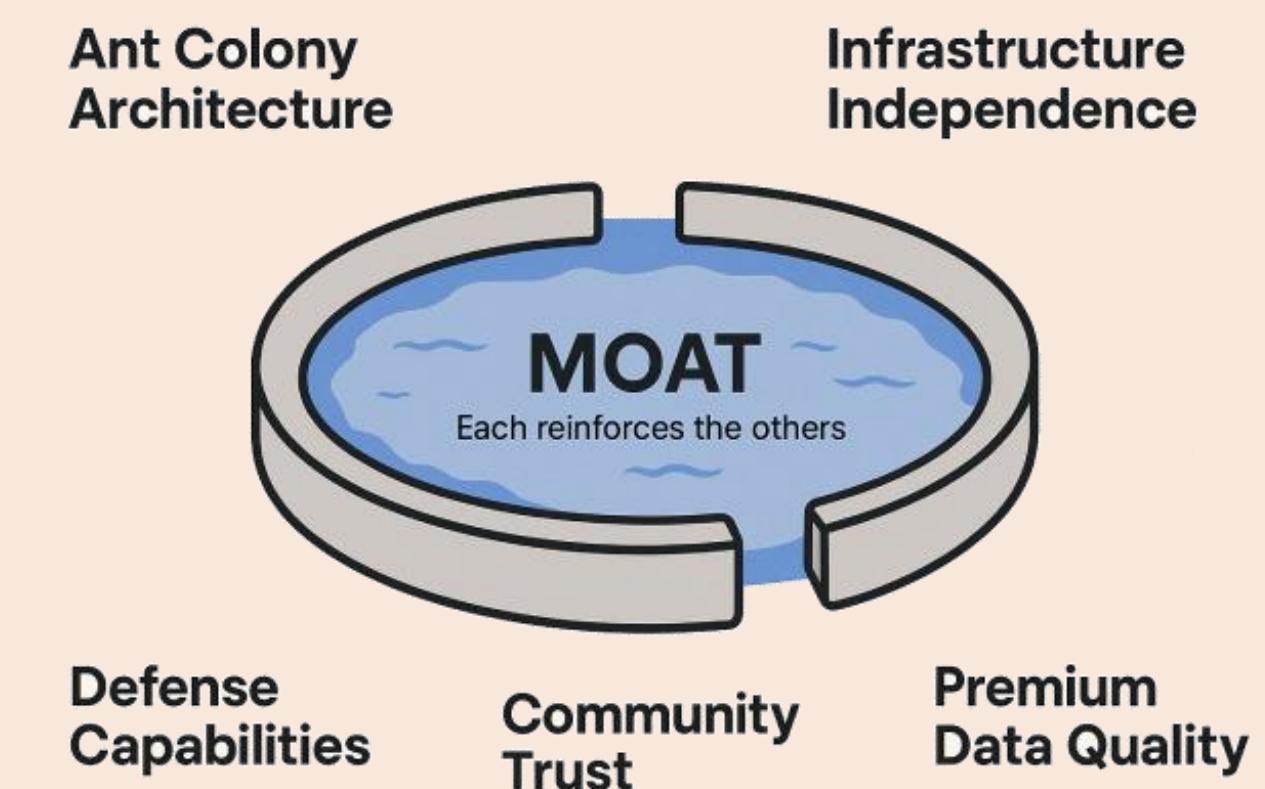
- Launch sovereign data labeling and model ecosystems across Europe and Latin America
- **Build bilingual bridges:** Spanish (Spain + Mexico), Portuguese (Brazil), German, French, Italian, Polish, Ukrainian
- Enter U.S. Government market via English-language annotation and model evaluation contracts
- Establish DEMOS's Queen Node architecture and first consumer AI agents

- **Phase 2 (Years 2–4): Pan-European, Latin American, and Arabic Expansion:**

- Scale to 20+ European, Latin, and Arabic dialects — extending into MENA's \$289B AI market
- Deploy **regional Edge AI nodes** linking Europe, Latin America, and the Middle East
- Partner with government, telecom, and enterprise institutions for localized training dataset
- Launch **Omnestra App Store** in multilingual markets for consumer and enterprise AI agent

- **Phase 3 (Years 4+): Global Majority (Asia & Africa):**

- Extend sovereign Edge AI ecosystems to **Asia-Pacific and Sub-Saharan Africa**
- Add **new languages** including Hindi, Mandarin, Tagalog, Swahili, Indonesian, and Yoruba
- Integrate **Edge hardware alliances** for large-scale, low-power deployment
- Position DEMOS as the global infrastructure layer for decentralized, multilingual AI



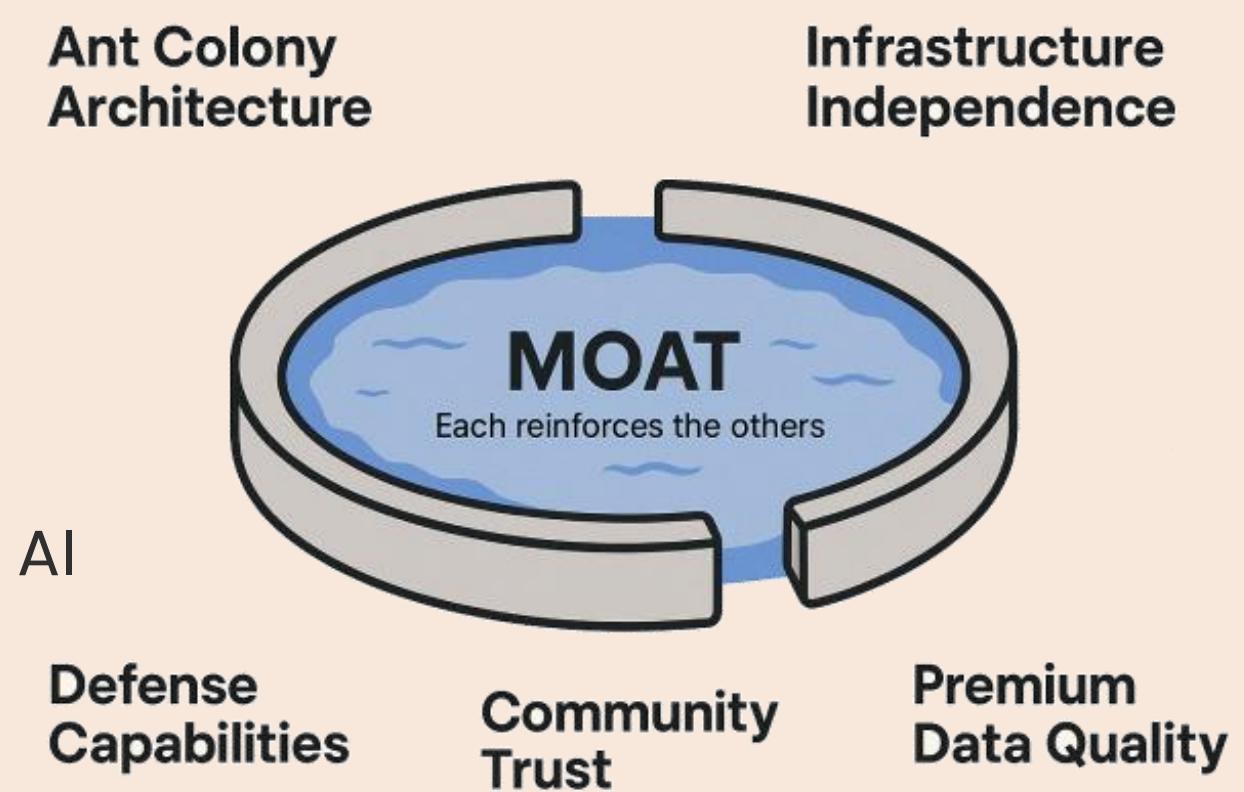
Why This Approach Works: Each phase compounds in data, language coverage, and model sophistication — creating a defensible, sovereign AI infrastructure that scales from Europe to the Global Majority.

Product Roadmap: 4 Phases

From Platform Buildout to a Global Multilingual AI Ecosystem

Overview: DEMOS evolves from an infrastructure builder into the operator of a global multilingual AI network — creating recurring revenue across data, model, and consumer layer. Each phase compounds the previous one, building technical capability, dataset depth, and sovereign market presence.

- **Phase 0 — Platform Buildout (2026)**
 - Develop DEMOS's foundational technology stack: data labeling, Ant Colony architecture, and Edge AI pipelines
 - Secure early government and enterprise (Open AI, Apple, Claude) labor contracts to fund pilot annotation work
 - Objective: Deliver a fully functional multilingual platform by year-end 2026
- **Phase 1 — Data Infrastructure (2027–2028):**
 - Monetize the platform through large-scale annotation and evaluation contracts
 - Establish sovereign datasets and labeling hubs across Europe, Latin America, and the U.S. Government sector
 - Objective: Generate recurring revenue while building proprietary multilingual dataset
- **Phase 2 — Ant Colony Learning Network (2028–2029)**
 - Launch self-learning, continuously retraining Edge AI model
 - Transition from human annotation to decentralized, automated data improvement
 - Objective: Monetize via model customization, API access, and enterprise integration
- **Phase 3 — Omnestra App Store (2029+)**
 - Launch multilingual Edge AI App Store — enabling consumers and enterprises to download localized AI agents
 - Agents communicate and learn through the Queen Node, creating continuous ecosystem growth
 - Objective: Establish DEMOS as the sovereign AI infrastructure layer for the global market



Why This Approach Works: A four-phase roadmap that starts with paid government contracts and evolves into a global AI marketplace — each stage deepening DEMOS's moat in data, language, and edge intelligence.

48-month Product Roadmap To **Market Dominance**

Phase 0: From Spanish Launch To Global Platform In 3 Years

Phase 0 – Platform Buildout & Early Contracts (2026):

Core Focus: Develop the DEMOS AI infrastructure and secure initial contracts

- Build core DEMOS Platform – data labeling ecosystem, Ant Colony coordination layer, and Edge AI model pipeline
- Hire engineering, linguistics, and government contract team
- Pursue pilot government and enterprise annotation contracts as paid labor program
- Begin internal testing of data ingestion, annotation workflow, and model retraining systems
- Lay the foundation for sovereign data hosting and multilingual model fine-tuning

Objective: By end-2026, deliver a functional DEMOS platform ready for multilingual dataset ingestion and early model deployment

Revenue Streams: Small-scale annotation & labor contracts • Early government pilots



48-month Product Roadmap To **Market Dominance**

Phase 1: From Spanish Launch To **Global Platform In 3 Years**

Phase 1 — Paid Annotation & Data Infrastructure (2027–2028):

Core Focus: Monetize the platform through data contracts and sovereign datasets

- Launch Product 1: DEMOS Data Labeling Ecosystem
- Secure large-scale annotation and evaluation contracts with governments, enterprises, and NGOs
- Expand multilingual data hubs (Europe, Latin America, Middle East, U.S. Gov)
- Generate predictable service revenue while building proprietary training datasets

Objective: Establish **DEMOS** as a trusted multilingual data infrastructure provider — generating recurring revenue from government and enterprise contracts while building proprietary sovereign datasets that form the foundation for future model training and Edge AI products

Revenue Streams: Annotation services • Data licensing • Dataset hosting



48-month Product Roadmap To **Market Dominance**

Phase 2: From Spanish Launch To **Global Platform In 3 Years**

From Training Data To **Self Deployment**

Phase 2— Ant Colony Learning Network (2028–2029):

Core Focus: Evolve DEMOS' Training AI Ecosystem into a closed loop system that continuously evaluates, refines, and deploys edge-native models across languages and regions

Training AI Evolution:

- Integrate annotation, QA, and feedback loops into automated training pipelines
- Implement continuous evaluation datasets to test model performance, bias, and safety
- Enable post-deployment feedback collection and real-world fine-tuning
- Establish model certification for sovereign and enterprise deployment

Edge AI Launch:

- Launch Product 2: Edge AI Models trained on sovereign, local data
- Deploy Ant Colony Architecture – a distributed coordination layer where millions of micro-agents learn locally and share optimized behaviors through the Queen Node
- Begin recurring revenue from model subscriptions, API's, and evaluation contracts.

Objective: Transform **DEMOS** from a data services company into a self-learning AI infrastructure – deploying edge-native models that continuously retrain through the Ant Colony network, creating scalable, recurring revenue from model customization, APIs, and sovereign enterprise deployment

Revenue Streams: Model subscriptions • AI training APIs • Enterprise integrations

48-month Product Roadmap To **Market Dominance**

Phase 3: From Spanish Launch To **Global Platform In 3 Years**

Phase 3 – Omnestra App Store (2029 +):

Core Focus: Scale from Infrastructure to ecosystem.

- Launch Product 3: **Omnestra App Store** – multilingual marketplace for AI agents
- Consumers and enterprises download **edge-native agents** that learn via the Queen Node
- Expand ecosystem monetization through app commissions and premium model tiers
- **Users can voluntarily subscribe to periodic updates** from the Queen, enhancing personalization and local model performance
- **Developers and institutions can publish custom agents** to the store, creating an open ecosystem for sovereign AI innovation

Objective: Position **DEMOS** as the global platform for sovereign AI ecosystems, enabling continuous, decentralized intelligence at scale – where every device, language, and agent contributes to a living, self-improving network

Revenue Streams: App Store downloads & commissions • Agent marketplaces • Premium subscriptions

The Competitive Advantage

Early Traction & Momentum

Building Foundation For 48-month Sprint

- **Talent Pipeline:** Stanford AI Lab researchers excited about lightweight AI mission
- **Strategic Partnerships:** Building relationships with Polish and Ukrainian drone companies
- **Government Opportunities:** Team actively identifying U.S. government contracting opportunities
- **Technical Validation:** Stanford, MIT, and Carnegie Mellon research proves lightweight AI works (as noted by Karen Hao)
- **Market Pull:** Strong demand for Language Edge Based AI



Talent Pipeline



Strategic Partnerships



Government Opportunities



Technical Validation

Our Unbreakable Moats

Seven Moats That Compound Over Time

- **Ant Colony Architecture:** Distributed micro-agents that learn locally and coordinate through the Queen Node—enabling AI that works offline, improves continuously, and scales without cloud infrastructure. Competitors building centralized systems can't easily replicate this model.
- **Infrastructure Independence:** Runs entirely on-device and offline, reaching the 6.5 billion people cloud AI can't serve
- **Premium Language & Data Sovereignty:** Locally governed, expert-labeled datasets in native languages and dialects that Big Tech cannot ethically or technically reproduce
- **Defense Background:** Backed by a team with deep U.S. Department of Defense contracting experience, former Pentagon advisors, and a West Point PhD from MIT, DEMOS is strategically positioned to secure early government AI and data labeling contracts, leveraging insider understanding of federal procurement and national security need
- **Resource Efficiency:** Designed for edge environments, DEMOS operates without data centers — eliminating the massive energy, water, and housing demands of Big AI, and delivering sovereign intelligence at a fraction of the cost
- **Language Sovereignty:** DEMOS is the only platform purpose-built for multilingual Edge AI, training native-language models that reflect local culture, semantics, and dialects — reaching the 6,500+ languages Big AI ignores
- **Research-Validated Efficiency:** Independent studies from MIT, Stanford, and IEEE confirm that small-parameter, edge-based models achieve near-LLM accuracy with up to 90% lower power and training costs, validating DEMOS's architecture as the future of sustainable A

Competitive Matrix

We're The Only Full Stack Solution

No Single Competitor Covers What We Cover

Capability	Demos	OpenAI	Palantir	Scale AI	Anduril	Snowflake
Lightweight AI	Yes	No	No	No	No	No
Works Offline	Yes	No	No	No	Yes	No
Native Languages	Yes	No	Partial	No	No	No
Data Annotation	Yes	No	Yes	Yes	No	No
Government/Defense	Yes	No	Yes	No	Yes	No
Physical AI/Drones	Yes	No	No	No	Yes	No
Data Infrastructure	Yes	No	Partial	No	No	Yes
Consumer Apps	Yes	Yes	No	No	No	No

The Team

Founders With \$3B+ Operational Experience

Building Foundation For 48-month Sprint

Co-Founders:

- **Valen Wardak (CEO/CTO):**

- Won Stanford CS109 competition with lightweight model proving small agents work
- Currently Pursuing Stanford Master in Computer Science, Focus on Artificial Intelligence
- Professor called it "one of the most substantial submission in Stanford's history"
- American Rhodes Scholar, Marshall Finalist, Georgetown Valedictorian
- Former Merrill Lynch Technology Investment Banker, built \$3.2B US DOD government contracting business

- **Mark B. Watson (CEO Business Development):**

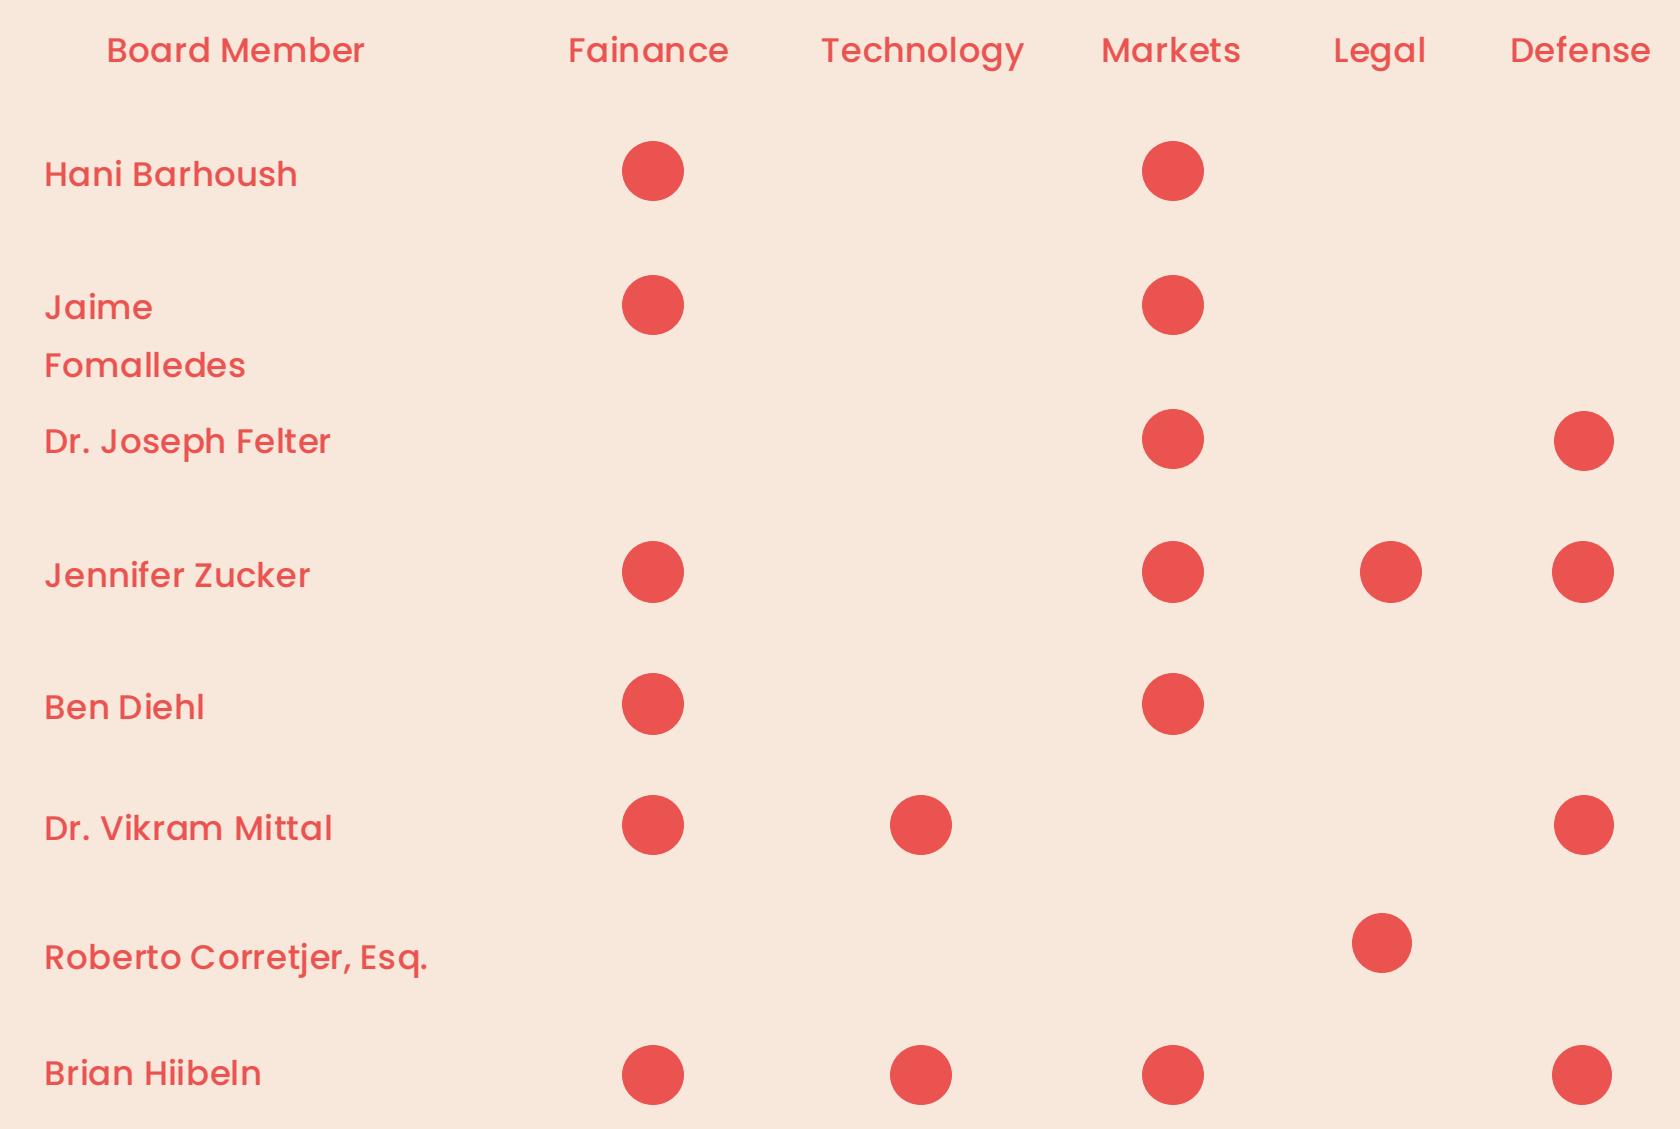
- Former Air Force contracting officer who authorized billions in defense procurements
- 25+ years defense tech, Inc 500 entrepreneur
- Knows both sides: awarded contracts as government, won them as contractor

The Origin: Valen's CS109 project revealed how catastrophically AI fails non-English speakers AND proved lightweight models could solve it

Board Of Directors - Unmatched Access

Board That Opens Every Door We Need

- **Hani Barhoush (Co-Chair):** CEO Mubadala Capital (\$25B AUM), Georgetown/Harvard
- **Dr. Joseph Felter (Co-Chair):** Former Deputy Assistant Secretary of Defense, Stanford fellow, defense contracts
- **Jennifer Zucker:** Co-Chair Greenberg Taurig's Government Contracts practice, former attorney to the U.S. Army Chief of Staff, Colonel in the U.S. Army Reserve, and trusted advisor to major private equity firms.
- **Brian Hibbeln:** Former Deputy Assistant Secretary of Defense for Special Capabilities and current senior advisor to Blackstone Private Equity .
- **Ben "Billions" Diehl:** 3x GRAMMY producer, unlocks entertainment industry capital.
- **Dr. Vikram Mittal:** Marshall Scholar. West Point Professor, MIT PhD, military AI expert.
- **Roberto Corretjer, Esq.:** Legal secretary, ensures compliance across jurisdictions.
- **Jaime Fonalledes:** Executive Vice President and heir apparent of Empresas Fonalledas, Inc.; attorney and Wharton-trained executive leading the next generation of Puerto Rico's largest private enterprise.



World-class Advisory Board

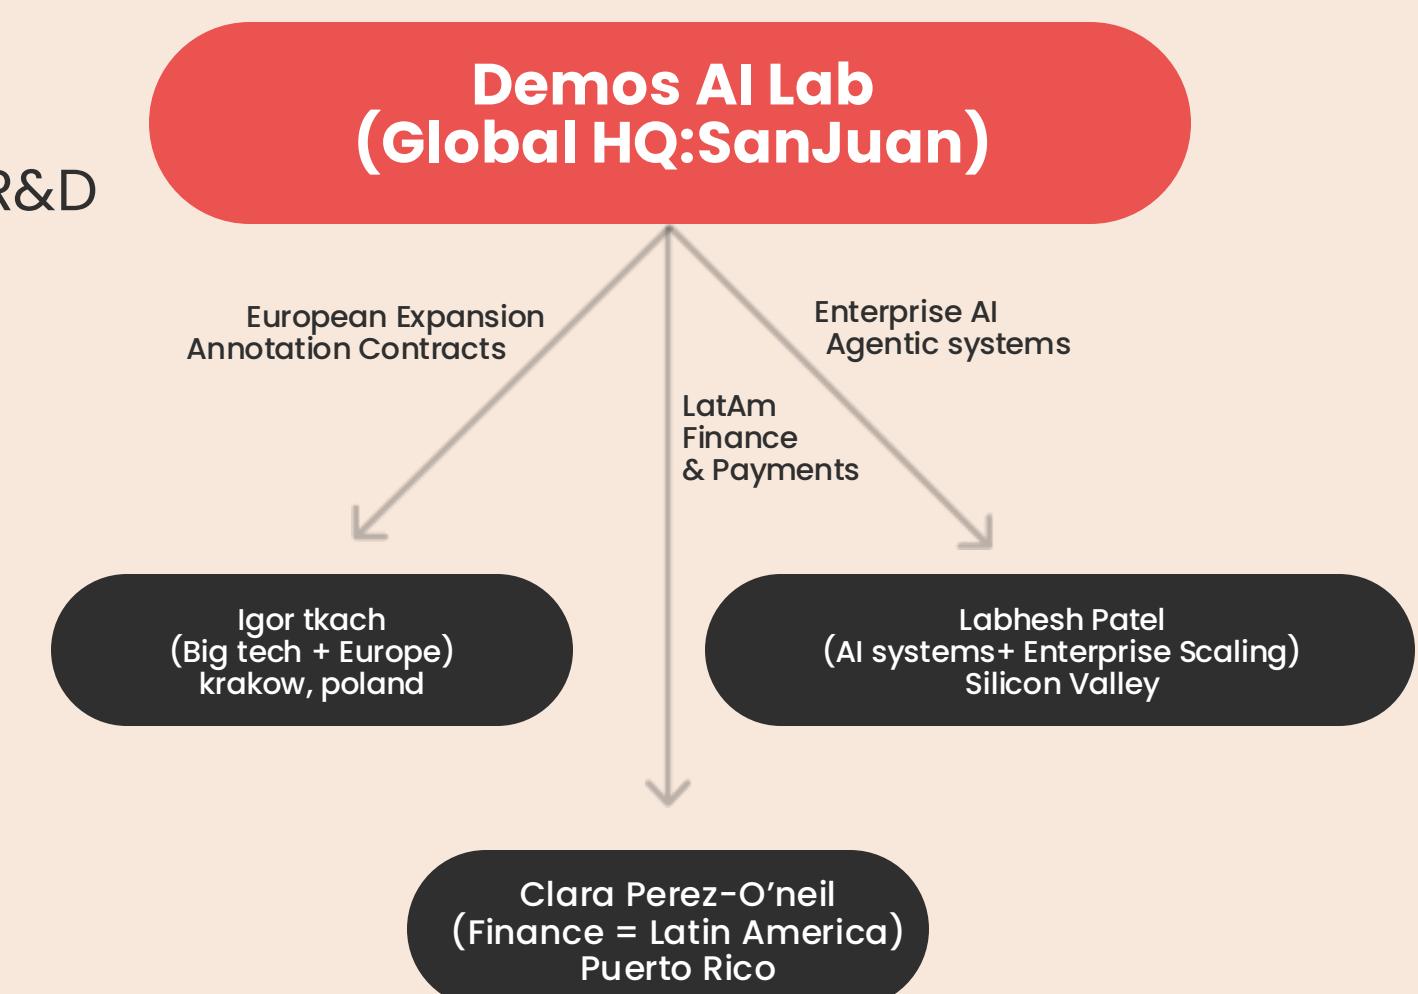
Advisors Who Deliver Global Reach, Capital, And Capabilities

- **Strategic Advisory Board:**

- **Igor Tkach (Co-Chairman, Strategic Technology Advisor):** CEO of Mindy Support (1,000+ annotators); delivered multilingual annotation services to Apple, Google, OpenAI; based in Krakow, Poland. Brings U.S.–European experience, deep knowledge of annotation markets, AI talent, and is a potential senior leader for Demos AI Lab.
- **Clara Perez-O'Neill (Co-Chairman, Senior Economic Advisor):** 25 years at Popular Securities; architected cross-border payment infrastructure for Latin America; expertise across American, Caribbean, and Latin American markets. Ensures scalable and compliant financial systems across dozens of currencies.
- **Labhesh Patel (Senior Technical Advisor, Agentic Solutions):** Former CTO & Chief Scientist at Jumio; scaled R&D from 40 to 500; 280+ U.S. patents. Founder of PlanetScale (compliance-aware multi-agent systems). Aligns with our mission of AI as augmentation, not replacement, ensuring our ant-colony architecture empowers communities.

- **What They Deliver:**

- Immediate access to Big Tech annotation contracts and European AI talent pools
- Expertise to architect Latin American payments & banking infrastructure
- Proven ability to scale AI enterprises from millions to billions in value
- Technical leadership on agentic AI architectures that enhance humans



The Ask

The Investment Opportunity

\$30m At \$300m Valuation = Ground Floor Of Next Unicorn

- **The Ask: \$30M seed round at \$300M pre-money valuation**

- **Compare to Recent AI Raises:**

- Safe Superintelligence: \$2B at \$32B valuation (no product, no revenue)
- Thinking Machines: \$1B at \$12B valuation (pre-product)
- Character AI: \$150M at \$1B valuation (chatbots only)
- Demos: \$30M at \$300M (serving 10x larger market with proven approach)

- **Why This is Undervalued:**

- Palantir at \$443B does government AI in English only
- Scale AI at \$29B does annotation in English only
- We're building BOTH for ALL languages = bigger opportunity

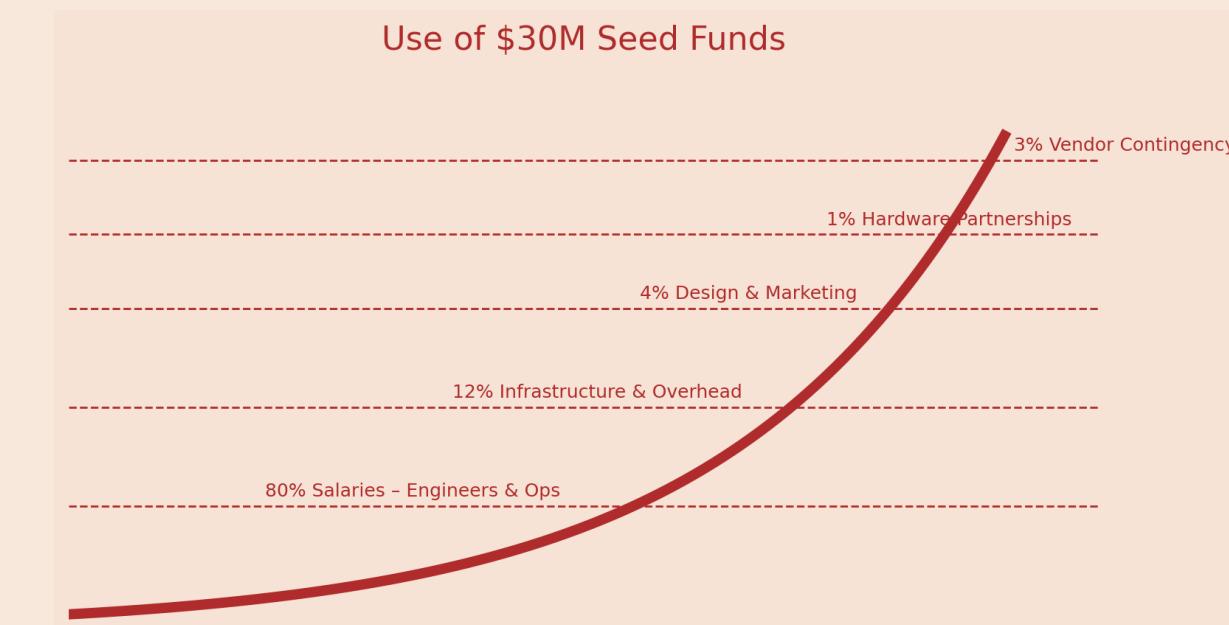


Use Of Funds & Path To Unicorn

48-month Sprint To Market Dominance

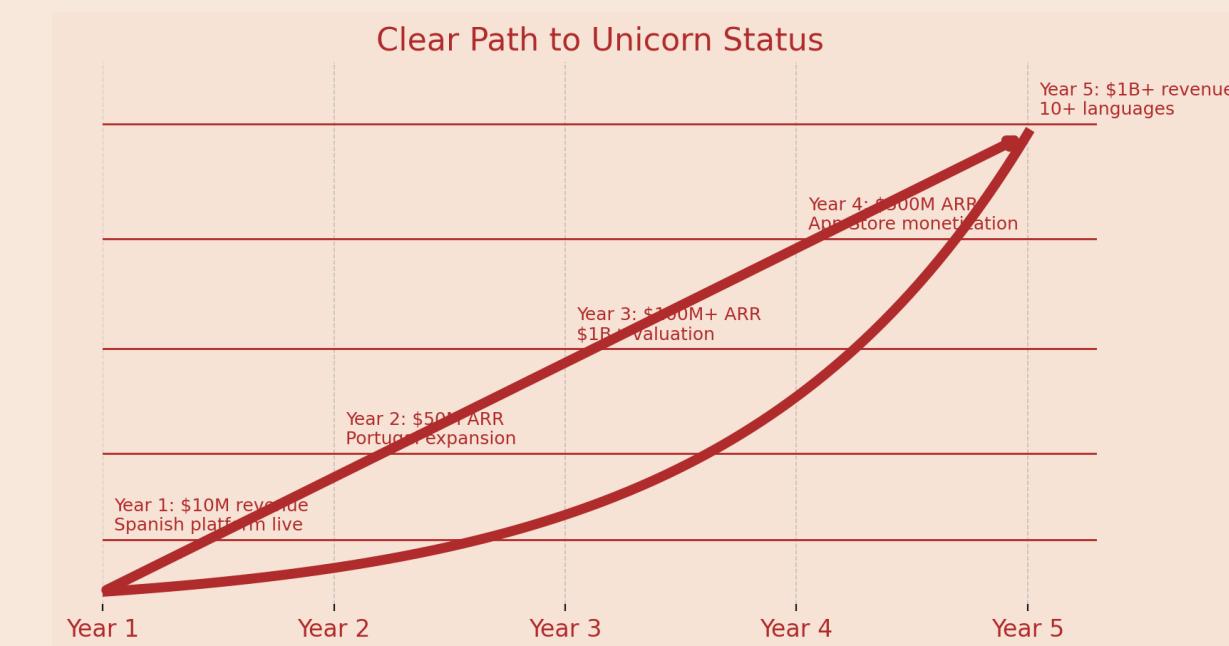
Use of \$30M Seed Funds:

- **80% Salaries:** Core AI engineers, researchers, and regional annotation managers building the DEMOS platform and multilingual datasets
- **12% Overhead:** Infrastructure, legal, and administrative operations to support distributed Edge AI development.
- **4% Design & Marketing:** UX/UI development for all language-specific DEMOS platforms (e.g., *Hermigas*, *Fermigas*, etc.), plus marketing to recruit annotators and drive government/enterprise adoption
- **1% Physical AI Partnership Development:** Hardware integration, drone testing, and edge device partnerships
- **3% Miscellaneous & Vendor Payments:** Contingency for scaling costs, early-stage refinements, and payments to third-party annotators and dataset curation vendors



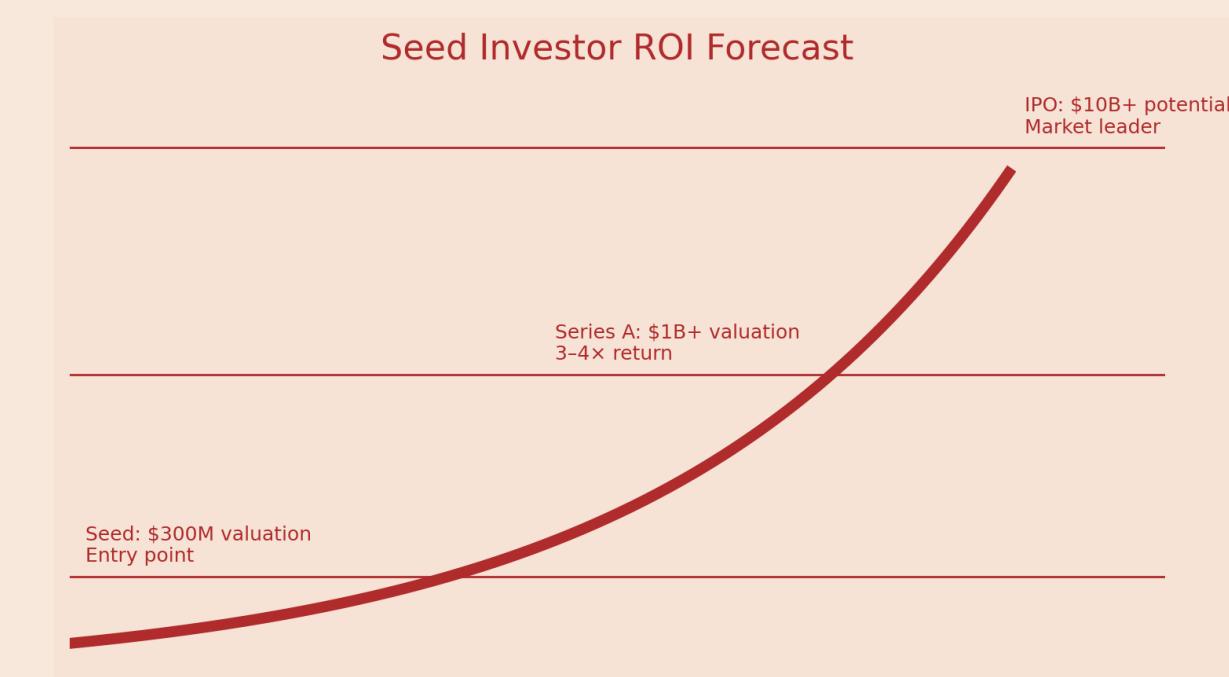
Clear Path to Unicorn Status:

- **Year 1 (2027):** \$10M revenue – Spanish platform live; initial government and enterprise annotation contracts executed
- **Year 2 (2028):** \$16M to \$30M ARR – Portuguese platform launched; expansion into sovereign data services across Europe and Latin America
- **Year 3 (2029):** \$30M to \$53M; Ant Colony Architecture fully deployed with active edge retraining nodes.
- **Year 4 (2030):** \$120M ARR & Billion Dollar Valuation – Omnestra App Store launch drives ecosystem revenue across Europe and Middle East.
- **Year 5 (2031):** \$100M to \$200M ARR & 1B+ Valuation – 10+ language ecosystems active across Europe, Latin America, and the Middle East, positioning DEMOS as the global leader in sovereign multilingual AI



ROI for Seed Investors:

- **Seed Stage:** \$300M entry valuation → **\$1B+** valuation within 48 months (≈3–4× return as DEMOS scales through Phases 0–2) → IPO: \$1B+ potential (Palantir/Scale AI trajectory)
- **Series A → Liquidity Event:** Multi-billion-dollar potential as DEMOS transitions from sovereign data infrastructure to a global multilingual AI ecosystem – following the growth trajectory of Palantir, Snowflake, and other enterprise AI leaders



The Vision

Puerto Rico: The Perfect Launchpad For Hermigas

Why Puerto Rico Becomes The Silicon Valley Of Spanish AI

- **Ideal Testing Ground for Spanish Platform:**
 - 3.2M Spanish speakers for immediate Hermigas deployment
 - Infrastructure challenges (blackouts) prove our offline capabilities
 - U.S. jurisdiction but Latin culture—perfect bridge market
- **Strategic Advantages:**
 - Act 60 tax incentives: 4% corporate tax, 0% capital gains
 - Educated bilingual workforce (highest PhD per capita in LatAm)
 - U.S. legal framework with Latin American cultural ties



Puerto Rico: The Perfect Launchpad For Hermigas

Why Puerto Rico Becomes The Silicon Valley Of Spanish AI

Manufacturing & Jobs Hub:

- Drone assembly facility creating 1000+ jobs
- Annotation center employing hundreds
- AI training hub for Spanish-speaking world

Gateway to 600M Spanish Speakers:

- Direct flights to all major LatAm cities
- Cultural credibility across Spanish-speaking world
- Proven resilience narrative resonates globally

From Adversity to Advantage: Hurricane Maria's blackouts aren't our weakness—they're our proving ground

Join The **AI Revolution**

Be Part Of History: Democratizing Ai For 6.5 Billion People

The Opportunity:

- First sovereign Edge AI platform serving the 6.5B people cloud AI cannot reach
- 10× larger market than English AI – currently valued at 1/10th the price
- Building the Palantir + Scale AI of multilingual, on device AI infrastructure
- Democratizing intelligent systems for every language, nation, and device

Why NOW is the Moment:

- **Lightweight > Heavyweight:** Research like DeepSeek proves smaller, efficient models win – DEMOS commercializes this at global scale
- **Strategic Timing:** Pentagon, EU, and Global South governments urgently need multilingual, disconnected AI
- **Data Moat:** Big Tech depends on English-language datasets – only DEMOS provides culturally native, sovereign data pipelines
- **Resilience Advantage:** Infrastructure failures in Puerto Rico, Ukraine, and Africa highlight the need for offline, edge-native AI

Join The **AI Revolution**

Be Part Of History: Democratizing AI For 6.5 Billion People

The Ask:

- Investors: Join at ground floor of next multi billion dollar company
- Partners: Governments, enterprises, drone manufacturers
- Customers: Be first to deploy truly sovereign AI

What Success Looks Like:

- Teacher in Guatemala earning dignified wages annotating data
- Farmer in Peru using \$5 agricultural agent offline
- U.S. soldier using lightweight translator in Afghanistan
- Every human accessing AI in their language, on their terms



AI For Every Language, Every Human

Language Is The Horizon Of Consciousness Itself

Our Promise: We're not just building technology. We're ensuring every language, every culture, every community can participate in the AI revolution with dignity and sovereignty.

The Choice Before Us: A world where AI is controlled by a few in a few languages, extracting value from the many... OR A world where every community owns their AI destiny through lightweight, sovereign technology.

Closing Statement

"Demos Is Democracy's Shield Building AI Infrastructure for the Next 6.5 Billion Users"



The Appendix: Branding

Appendix: Branding

Why THE ANTS?

Micro Intelligence. Massive Coordination: We deliberately chose the ant as our symbol because it reflects the power of micro and small-parameter models—autonomous, efficient, and collectively intelligent

Just as ants coordinate through shared signals, DEMOS's distributed models communicate and learn through the Ant Colony Architecture:

- Global Identity:** Every regional platform—Hermigas, Formigas, and others—takes its name from the word “ant” in its local language, symbolizing unity across diversity
- Collective Intelligence:** Each “ant” represents human contributors, devices, and models working together to build localized, sovereign AI
- Natural Efficiency:** Ants thrive through low-power collaboration—just like DEMOS’s lightweight edge models
- Universally Understood:** Found on every continent, the ant transcends language and culture—mirroring DEMOS’s mission to build AI that speaks to everyone